

# Quantization

# Contents

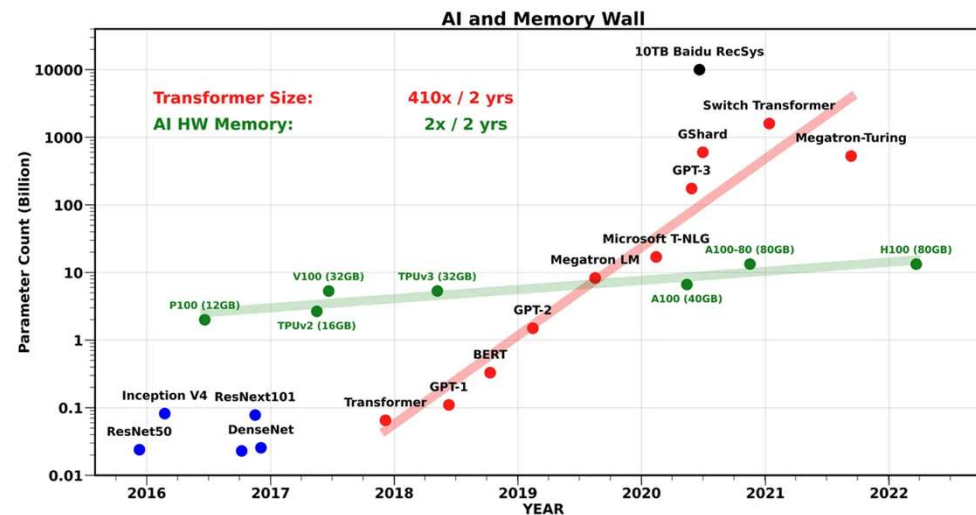
- **Background**
- **Preliminaries**
- **Quantization Basics**
- **Quantization Strategy**
- **Practice**

# Background

- **Why is Quantization Important?**
  - Imbalance in the improvement of hardware resources

# Background

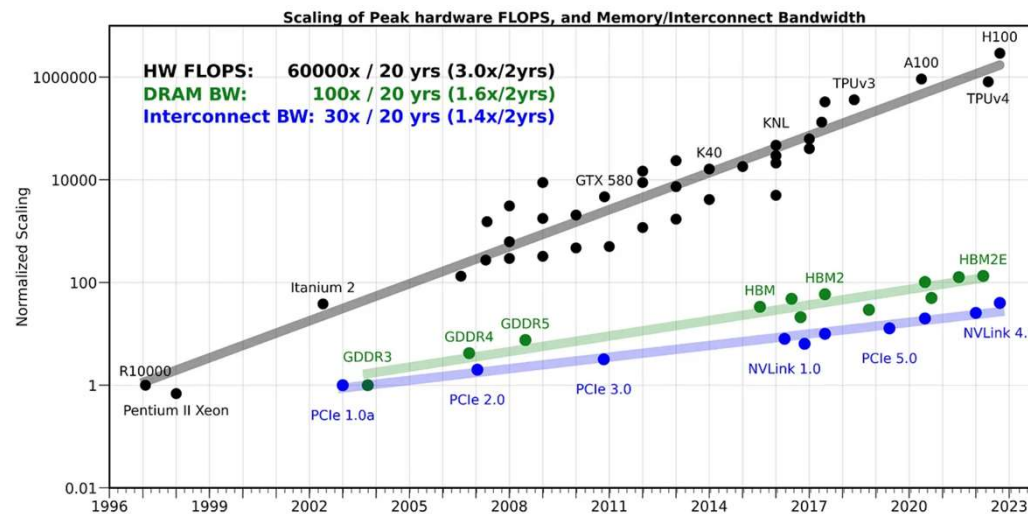
- **Why is Quantization Important?**
  - Imbalance in the improvement of hardware resources



1. **Models are expanding rapidly**, but **hardware resources** (particularly VRAM) are not keeping pace, primarily due to cost constraints.

# Background

- **Why is Quantization Important?**
  - Imbalance in the improvement of hardware resources



2. Hardware **compute power is increasing fast**, while **memory bandwidth** is growing at a **slower rate**.

# Background

- **Why is Quantization Important?**
  - Imbalance in the improvement of hardware resources
  - To maximize throughput/latency, we need to fully utilize the hardware.  
(We want our model to run faster on newer H/W)

# Background

- **Why is Quantization Important?**
  - Imbalance in the improvement of hardware resources
  - To maximize throughput/latency, we need to fully utilize the hardware.  
(We want our model to run faster on newer H/W)
  - Quantization is one of the most practical method to solve this problem.

# Background – Techniques

- **Optimizing the NN architecture**
  - Designing efficient NN model architecture
  - Co-designing NN architecture and hardware together
- **Lightweighting NN model**
  - Knowledge Distillation
  - Pruning
  - Quantization



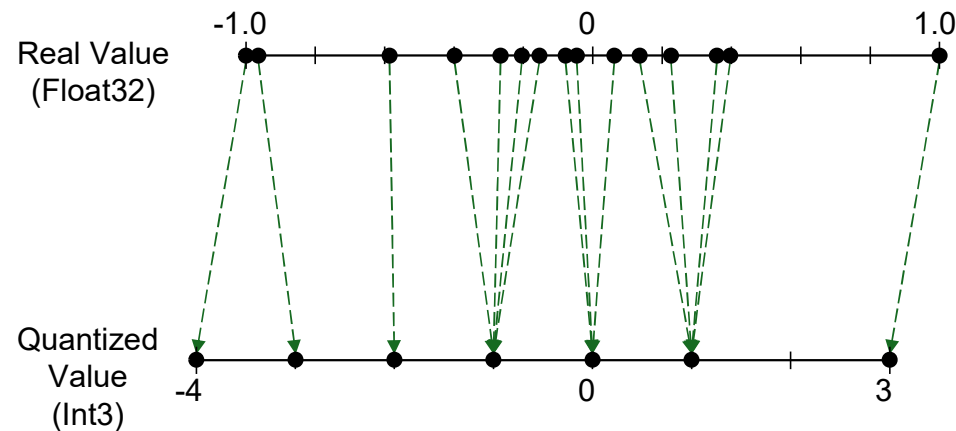
# Background – Techniques

- **Optimizing the NN architecture**
  - Designing efficient NN model architecture
  - Co-designing NN architecture and hardware together
- **Lightweighting NN model**
  - Knowledge Distillation
  - Pruning
  - **Quantization**

# What is Quantization?

Mapping **continuous (Float)** values into smaller set of **discrete (Integer)** values

- Example (3-Bit Uniform Quantization)



# How do we optimize with Quantization?

Quantization inevitably have a **negative impact on accuracy**

How to overcome this problem?

# How do we optimize with Quantization?

Quantization inevitably have a **negative impact on accuracy**

How to overcome this problem?

1. Quantize everything and find way to restore accuracy

# How do we optimize with Quantization?

Quantization inevitably have a **negative impact on accuracy**

How to overcome this problem?

1. Quantize everything and find way to restore accuracy
2. Partially apply quantization to avoid affecting accuracy

# How do we optimize with Quantization?

Quantization inevitably have a **negative impact on accuracy**

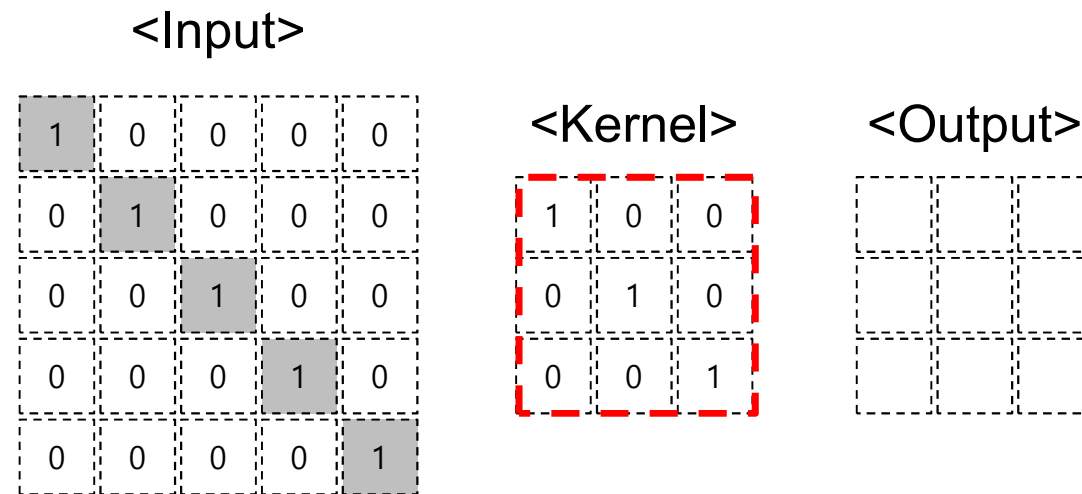
How to overcome this problem?

1. Quantize everything and find way to restore accuracy
2. Partially apply quantization to avoid affecting accuracy
  - Quantizing a few operation **would not change** model output **much**
  - **Find expensive operation** with a long latency **and quantize** them

# How do we optimize with Quantization?

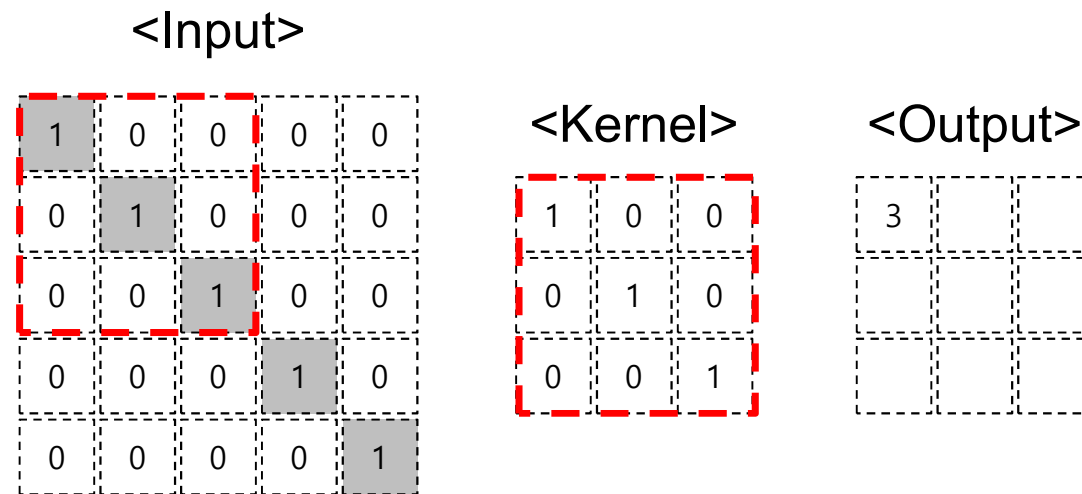
- Examples of **expensive** operations in DL
  - **Convolution operation**
  - **Linear operation**
- Examples of inexpensive operations in DL
  - Normalization Layer
  - Element-wise Operation (Residual add/concatenation)

# Preliminaries – Convolution operation

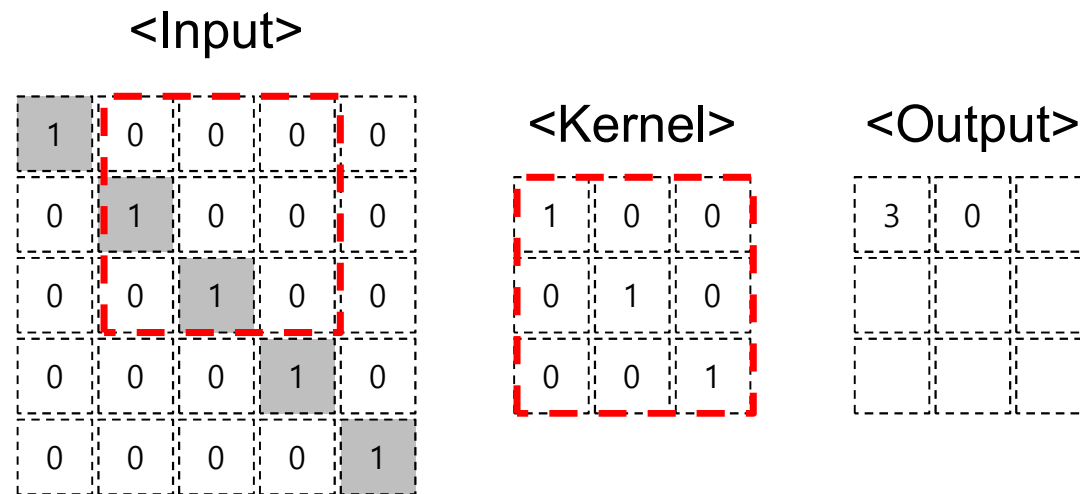




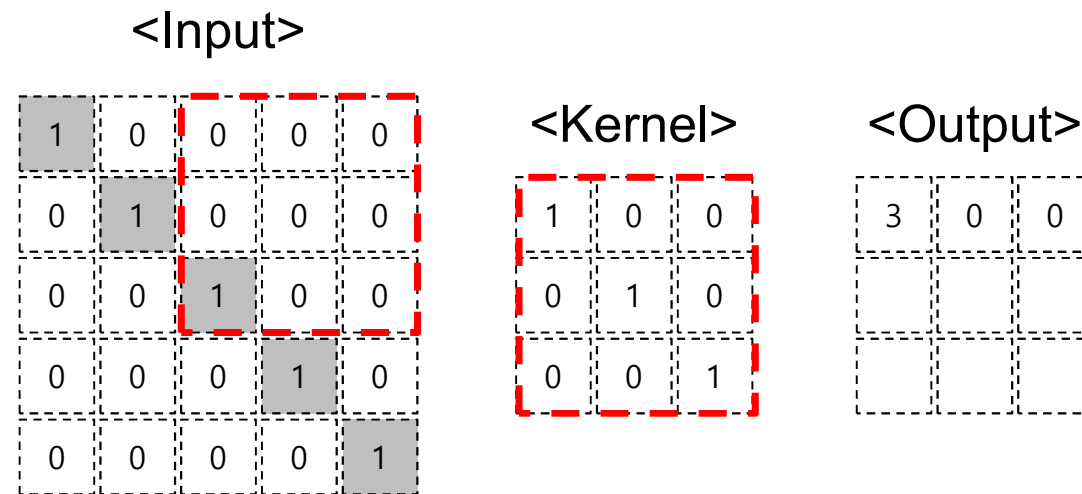
# Preliminaries – Convolution operation



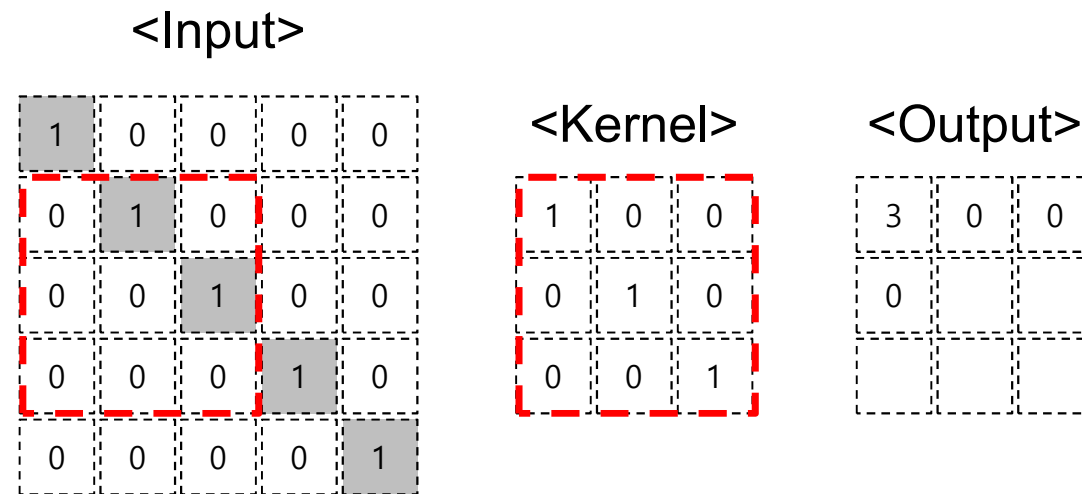
# Preliminaries – Convolution operation



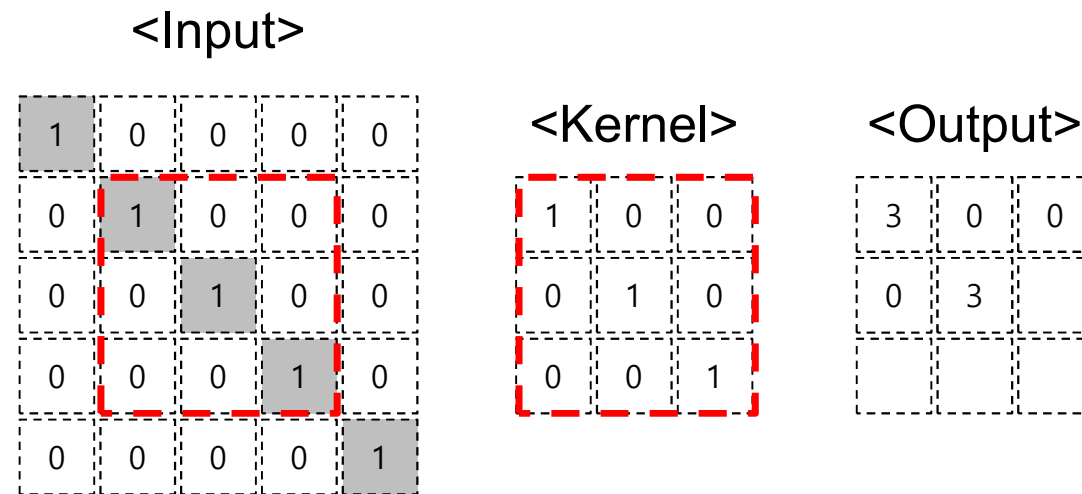
# Preliminaries – Convolution operation



# Preliminaries – Convolution operation



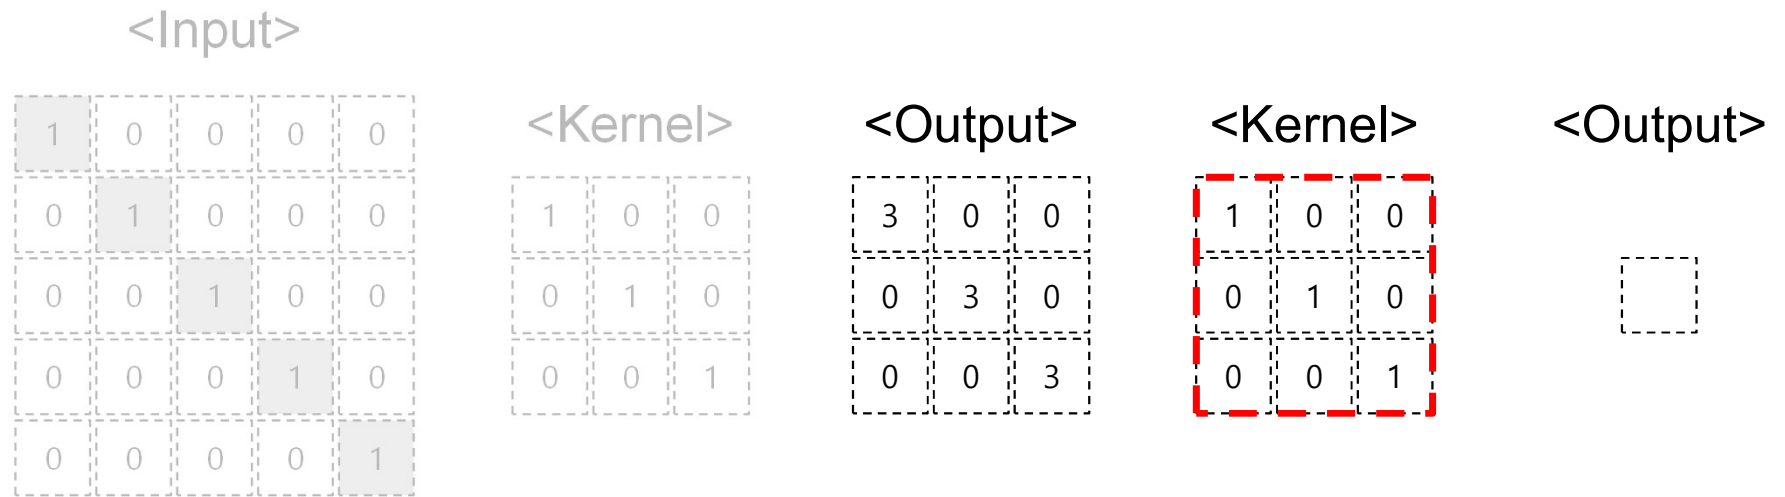
# Preliminaries – Convolution operation



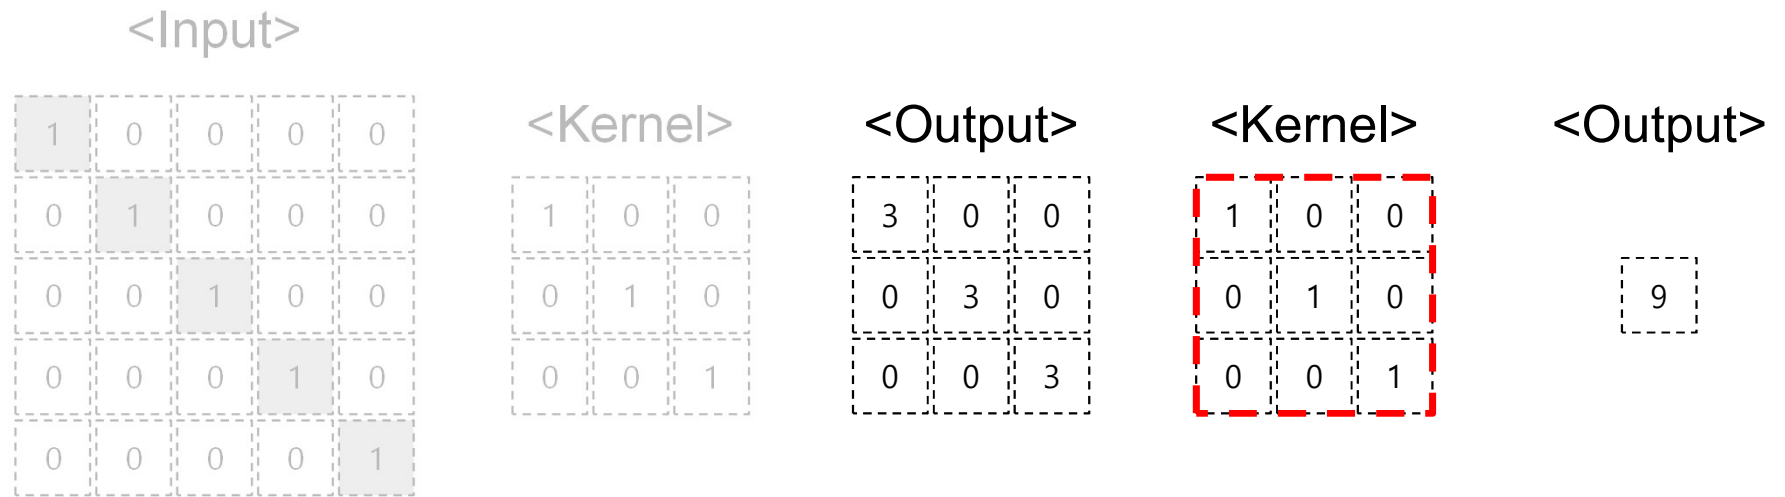
# Preliminaries – Convolution operation

<Input>					<Kernel>			<Output>		
1	0	0	0	0	1	0	0	3	0	0
0	1	0	0	0	0	1	0	0	3	0
0	0	1	0	0	0	0	1	0	0	3
0	0	0	1	0						
0	0	0	0	1						

# Preliminaries – Convolution operation

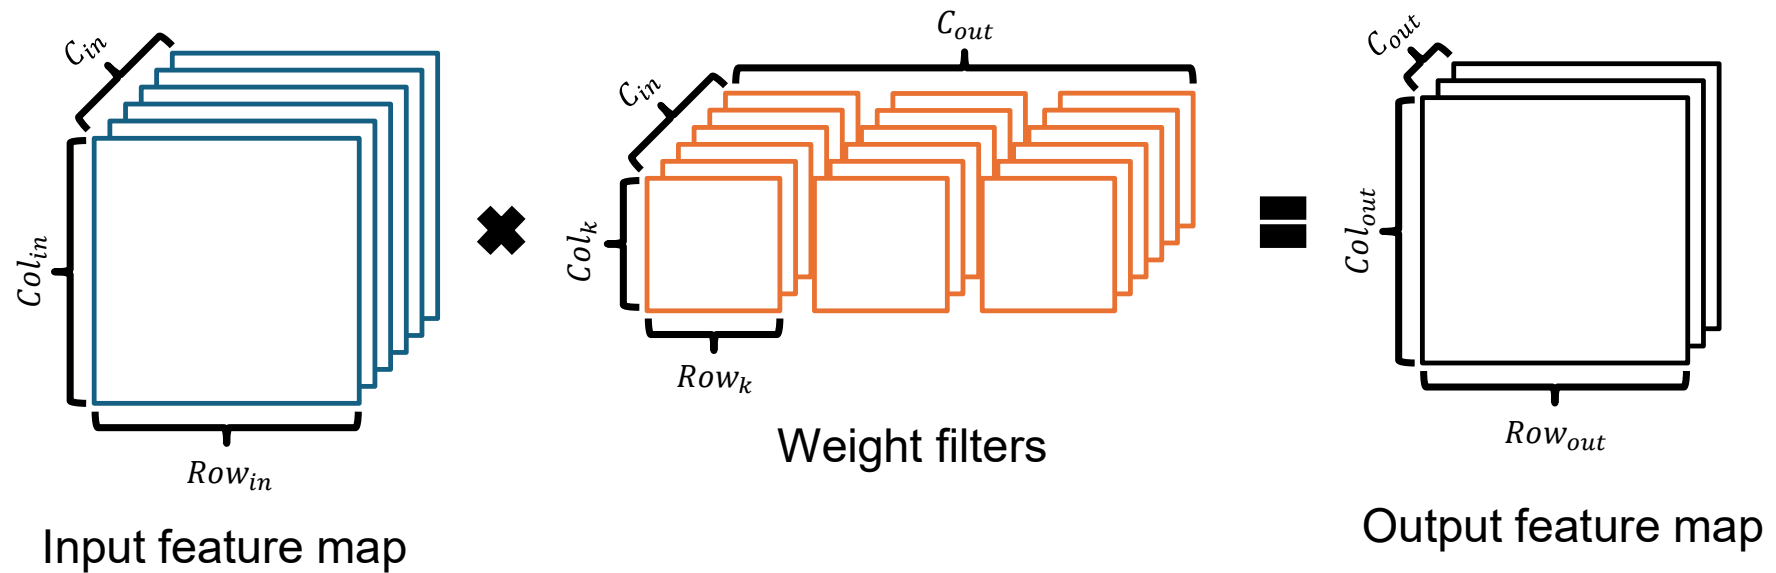


# Preliminaries – Convolution operation

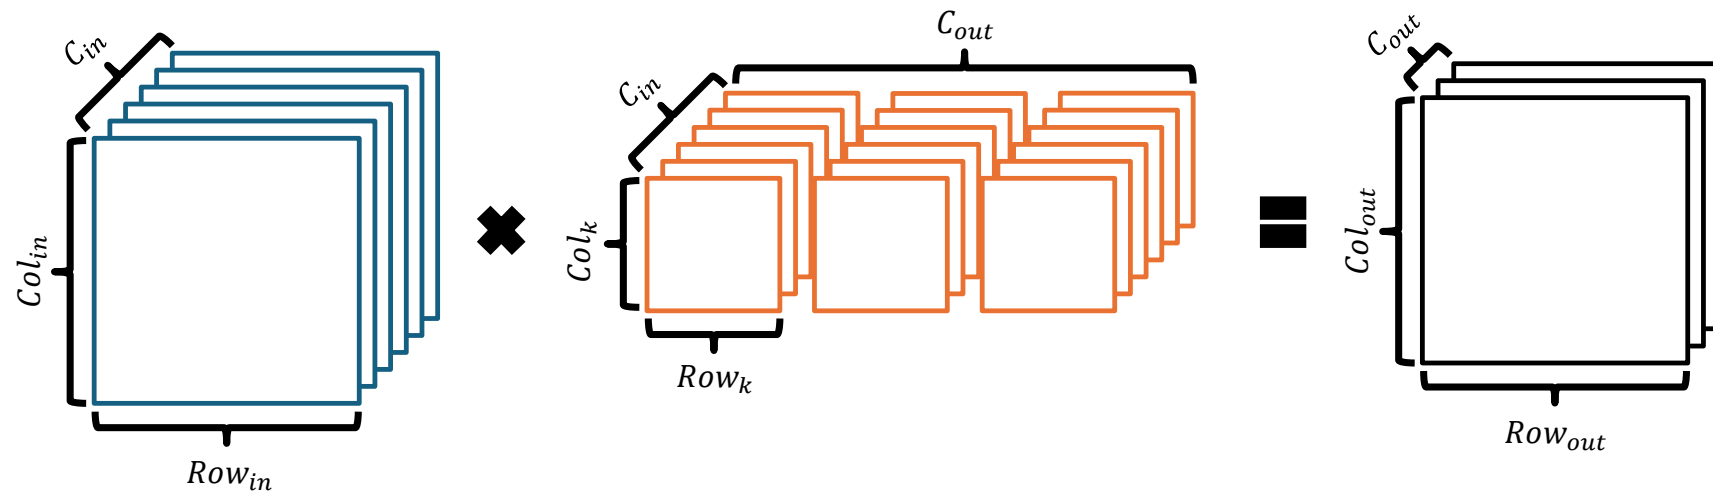




# Preliminaries – Convolution operation

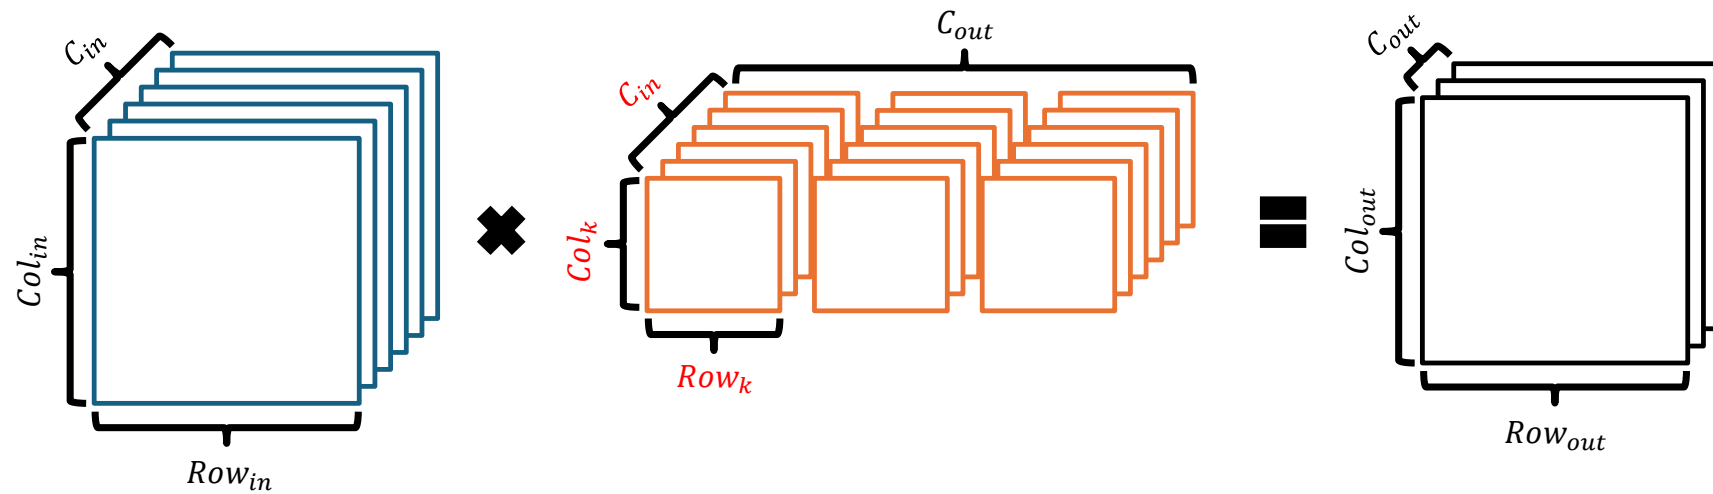


# Preliminaries – Convolution operation



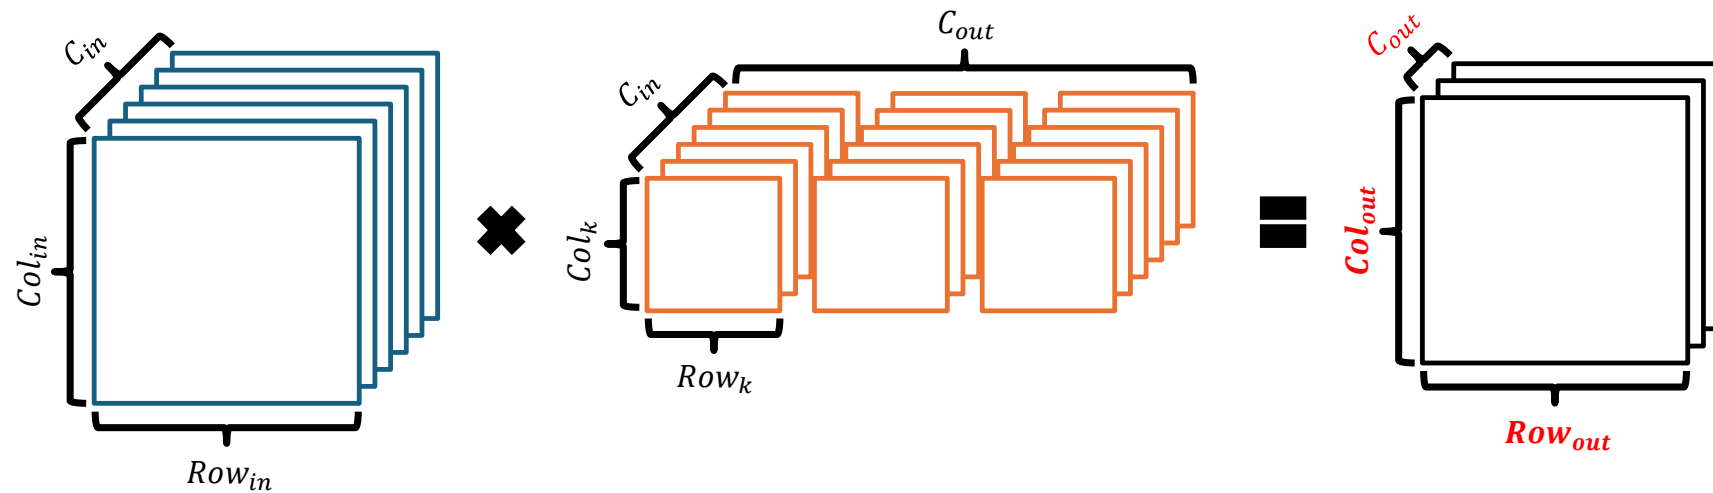
$$FLOPs = Col_k \times Row_k \times C_{in} \times Col_{out} \times Row_{out} \times C_{out}$$

# Preliminaries – Convolution operation



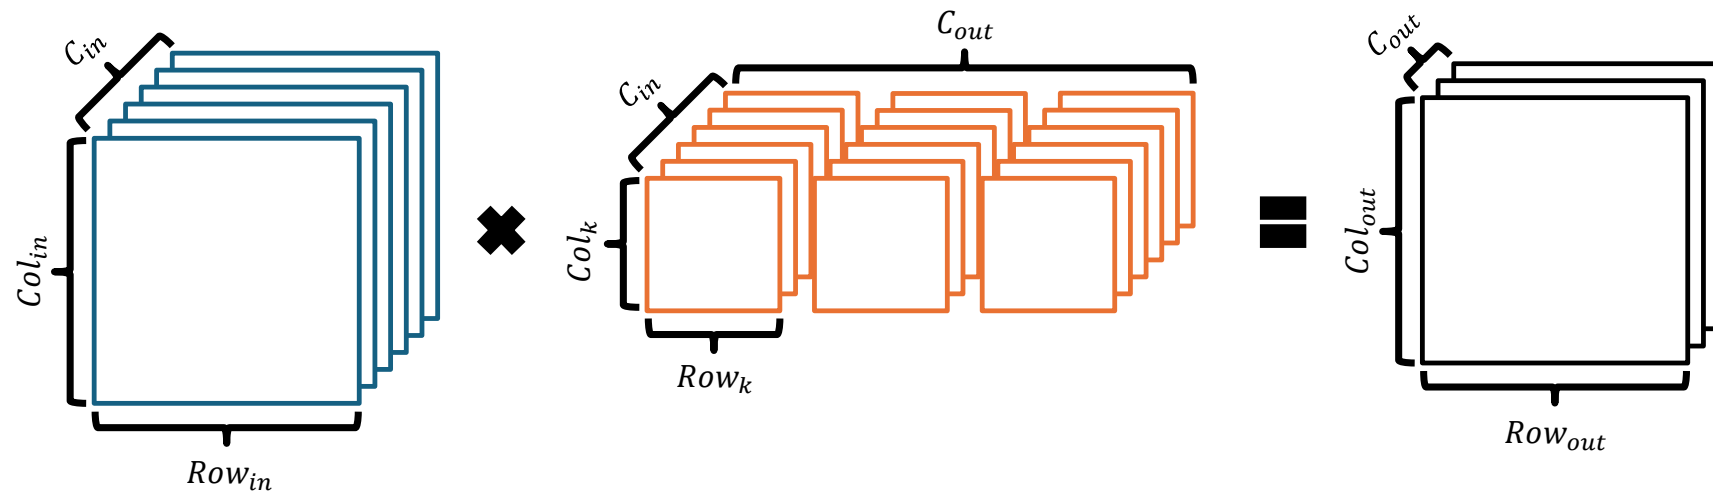
$$FLOPs = Col_k \times Row_k \times C_{in} \times Col_{out} \times Row_{out} \times C_{out}$$

# Preliminaries – Convolution operation



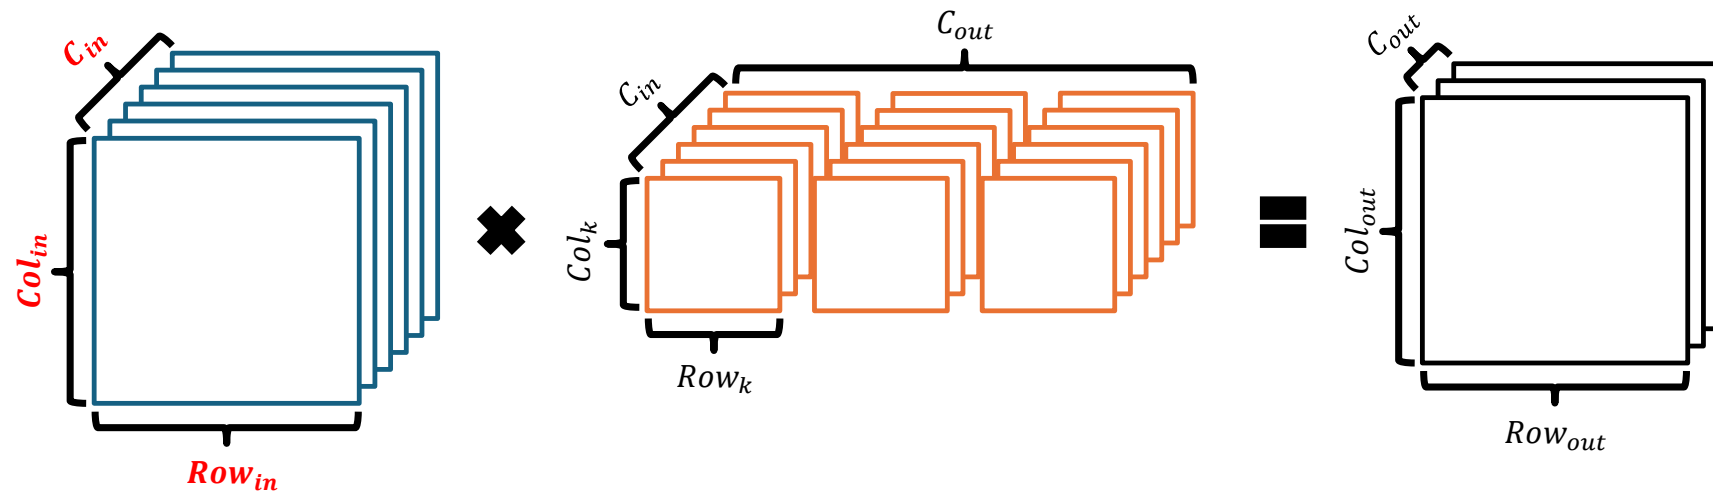
$$FLOPs = Col_k \times Row_k \times C_{in} \times Col_{out} \times Row_{out} \times C_{out}$$

# Preliminaries – Convolution operation



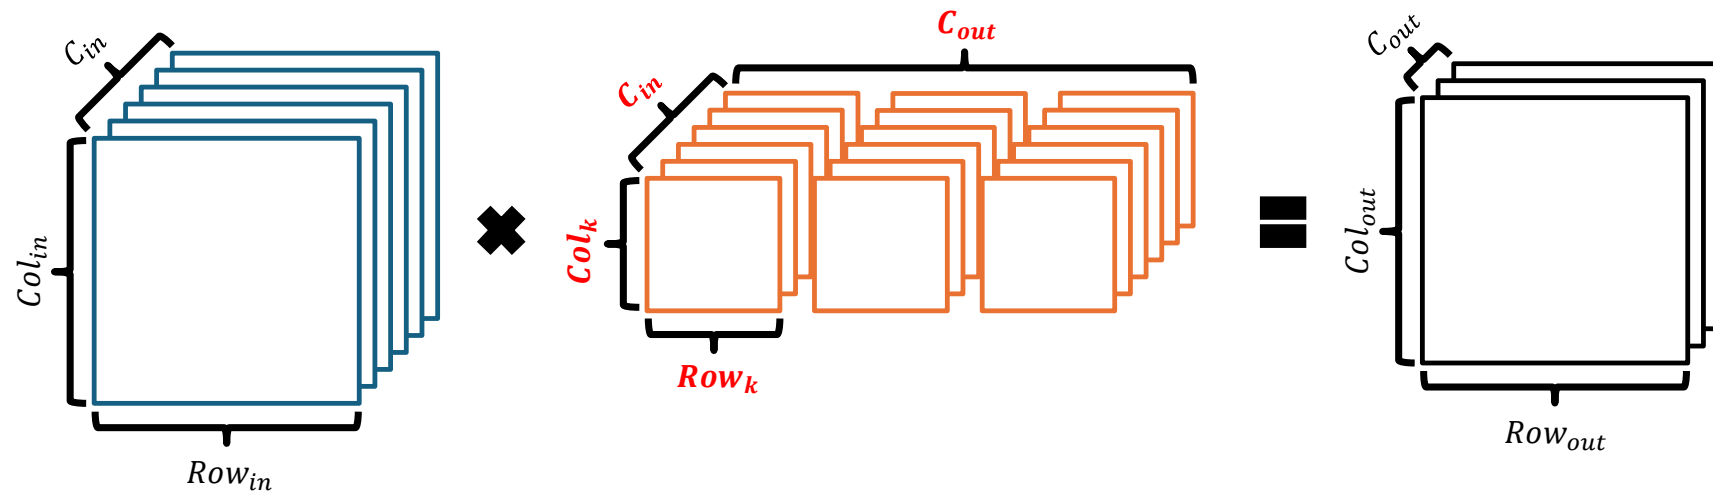
$$MEM = [Col_{in} \times Row_{in} \times C_{in}] + [Col_k \times Row_k \times C_{in} \times C_{out}]$$

# Preliminaries – Convolution operation



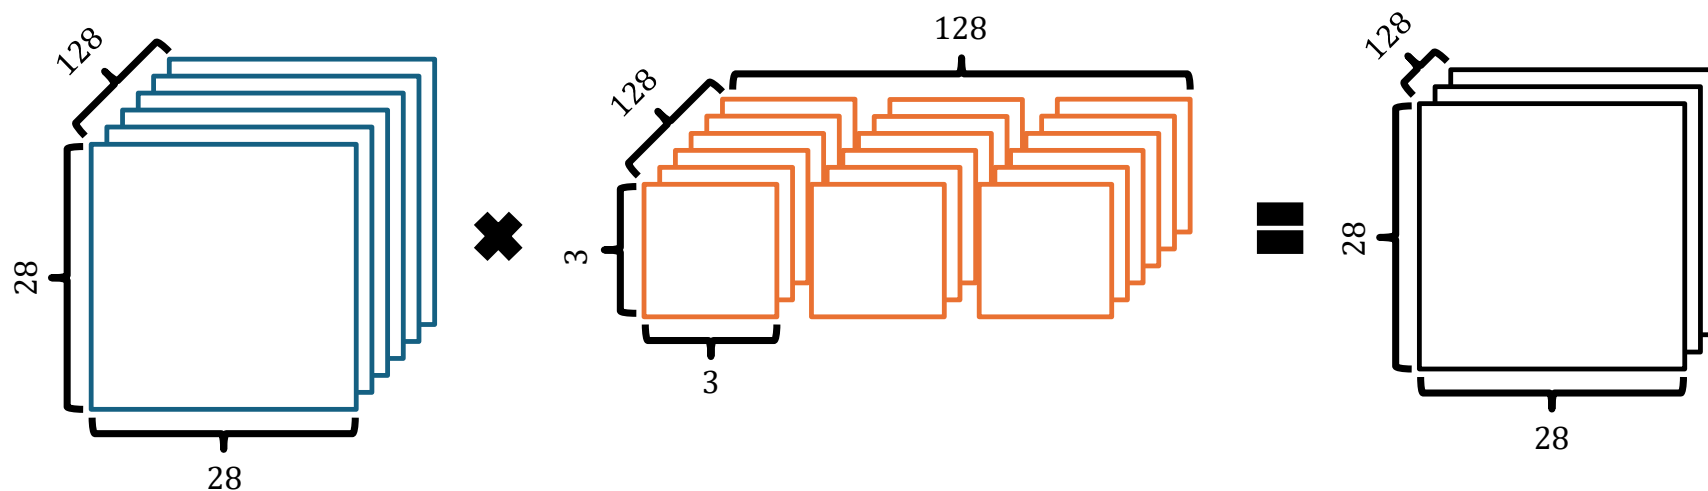
$$MEM = [Col_{in} \times Row_{in} \times C_{in}] + [Col_k \times Row_k \times C_{in} \times C_{out}]$$

# Preliminaries – Convolution operation



$$MEM = [Col_{in} \times Row_{in} \times C_{in}] + [Col_k \times Row_k \times C_{in} \times C_{out}]$$

# Preliminaries – Convolution operation



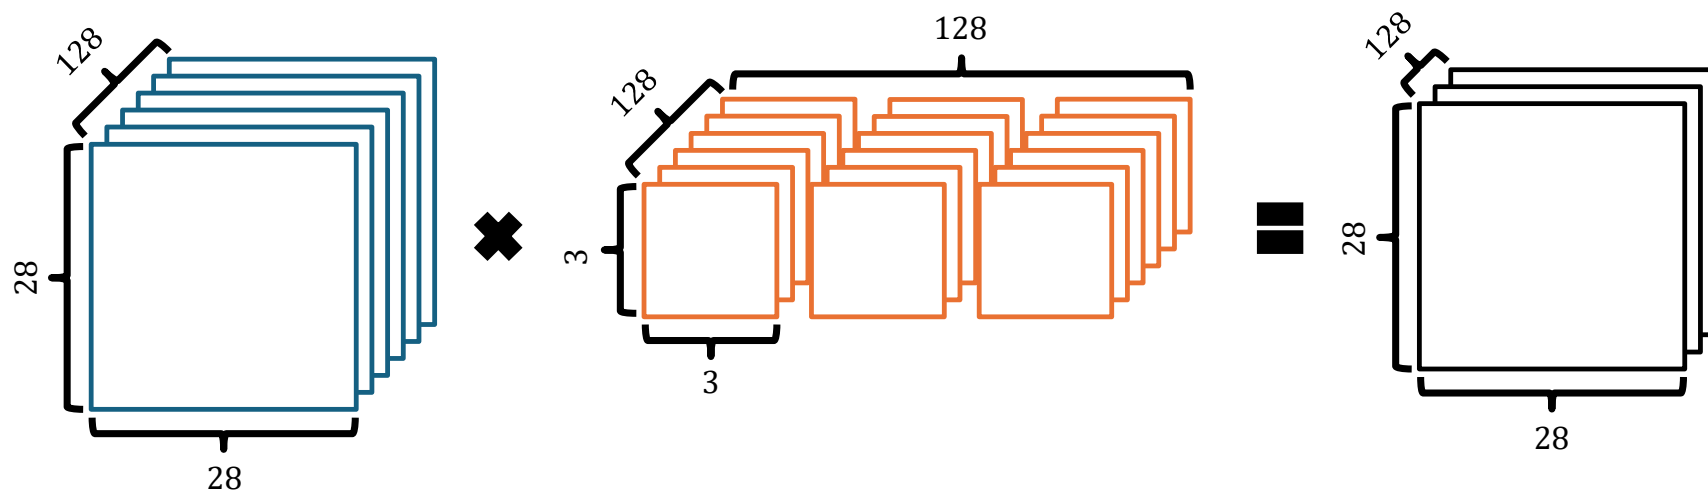
Example with ResNet50 (middle layer)

$$FLOPs = 3 \times 3 \times 28 \times 28 \times 128 \times 128 = 115M$$

$$MEM = [28 \times 28 \times 128] + [3 \times 3 \times 128 \times 128] = 0.10M + 0.15M$$



# Preliminaries – Convolution operation

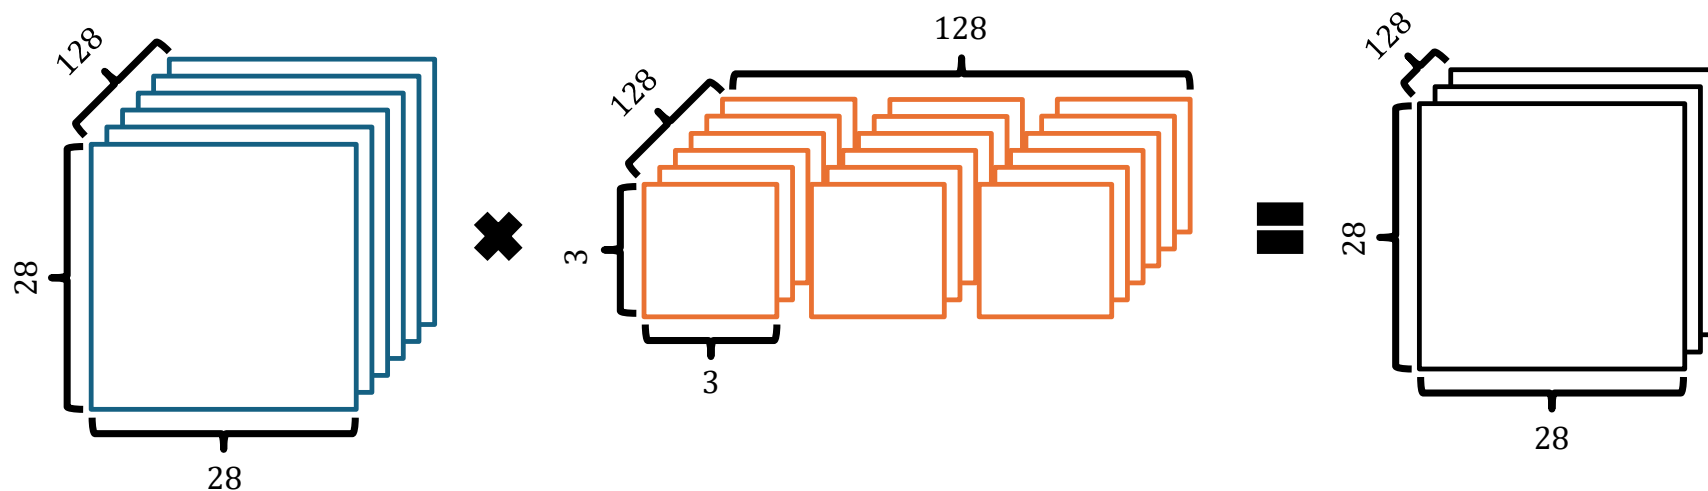


Example with ResNet50 (middle layer)

$$FLOPs = 115M$$

$$MEM = 0.10M + 0.15M$$

# Preliminaries – Convolution operation



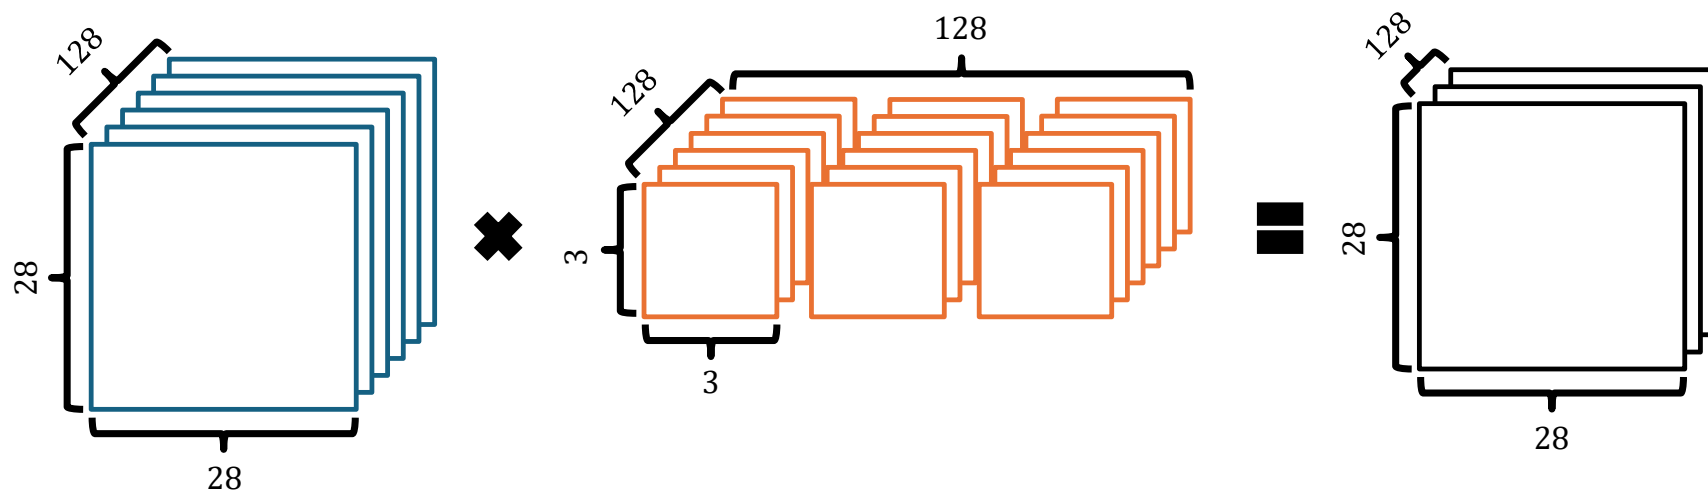
Example with ResNet50 (middle layer)

$$FLOPs = 115M$$

Suppose we are using **FP16**..

$$MEM = 0.10M + 0.15M = \mathbf{0.5MB}$$

# Preliminaries – Convolution operation



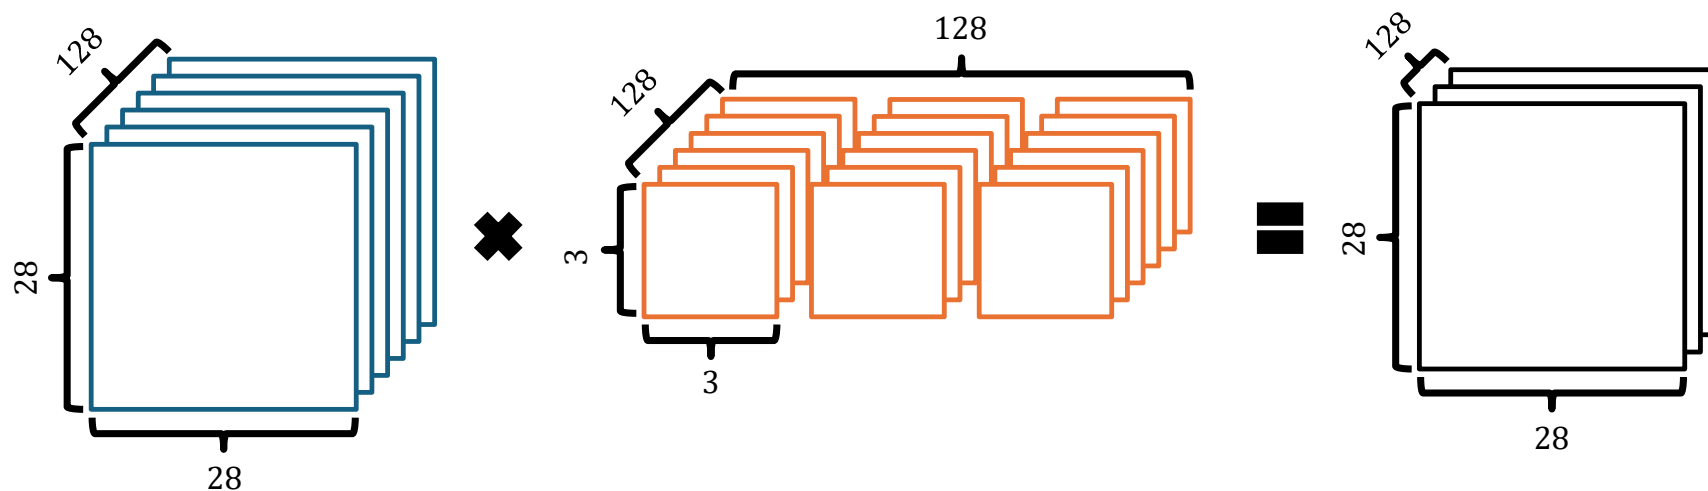
Example with ResNet50 (middle layer)

$FLOPs = 115M$

$MEM = 0.10M + 0.15M = 0.5MB$

So roughly assuming, our hardware need capability to process **115M** computation while reading **0.5MB**

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

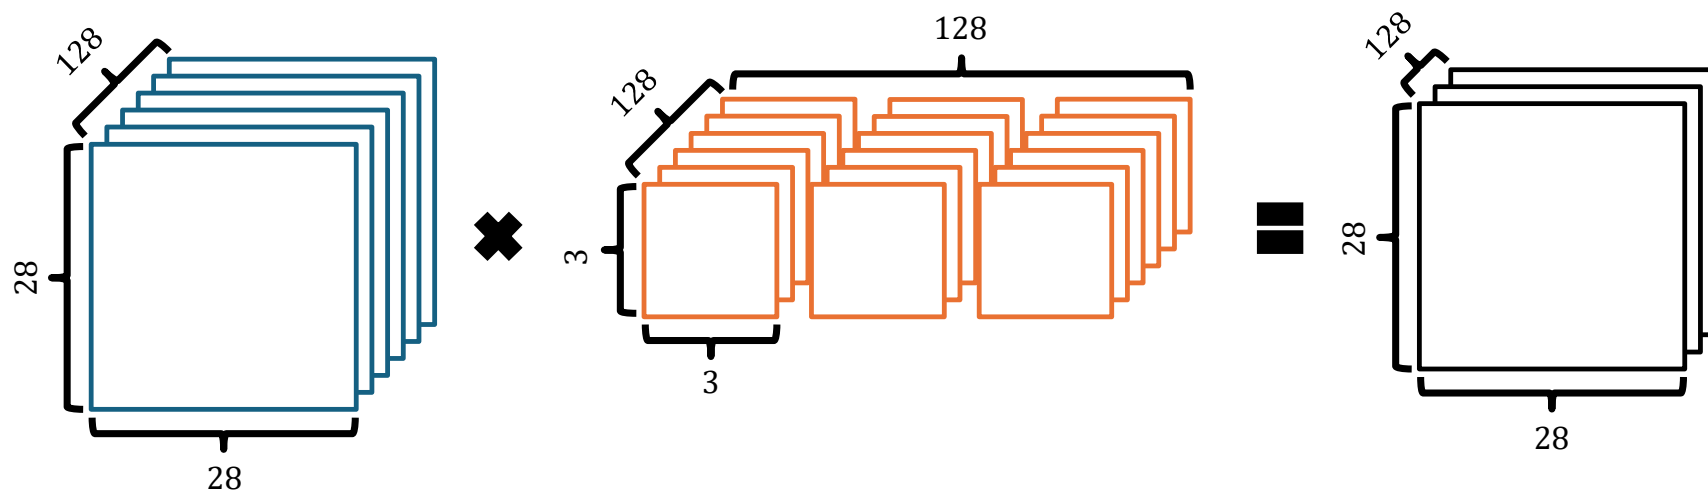
$$FLOPs = 115M$$

$$MEM = 0.10M + 0.15M = 0.5MB$$

So roughly assuming, our hardware need capability to process **115M** computation while reading **0.5MB**

This is called **Arithmetic intensity** : **FLOPs / Bytes**

# Preliminaries – Convolution operation



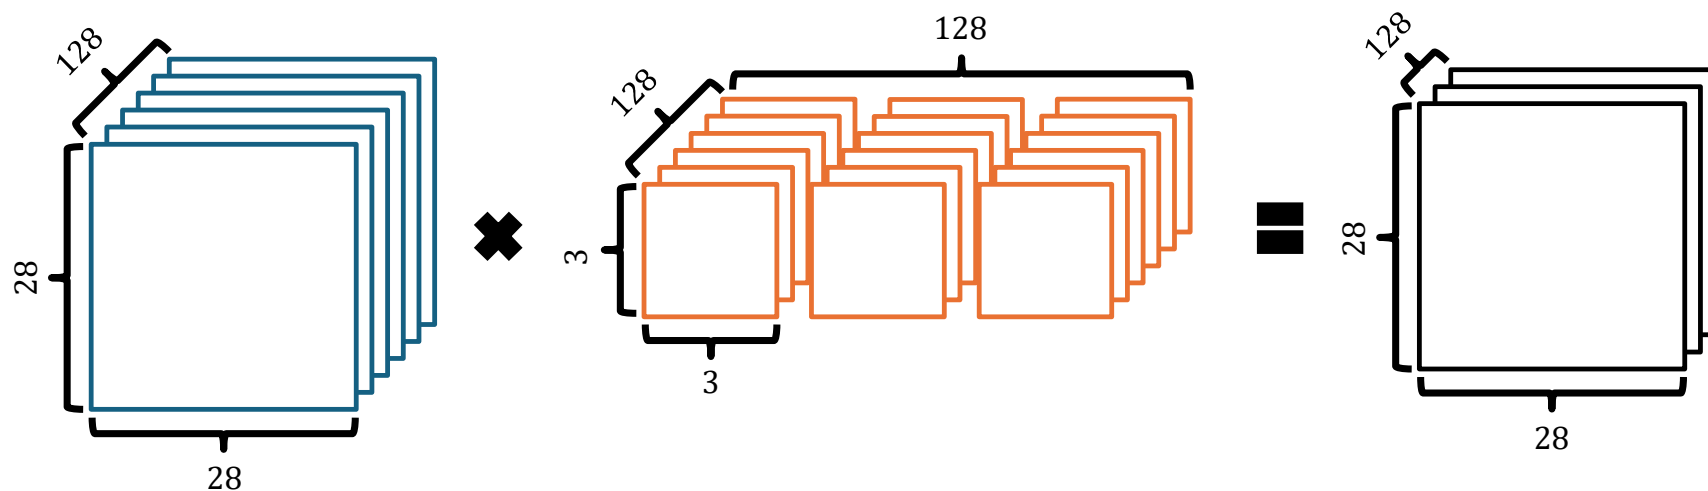
Example with ResNet50 (middle layer)

$$FLOPs = 115M$$

$$MEM = 0.10M + 0.15M = 0.5MB$$

Arithmetic intensity of A100 with FP16 :  
**208 FLOPs / Bytes**

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

$$FLOPs = 115M$$

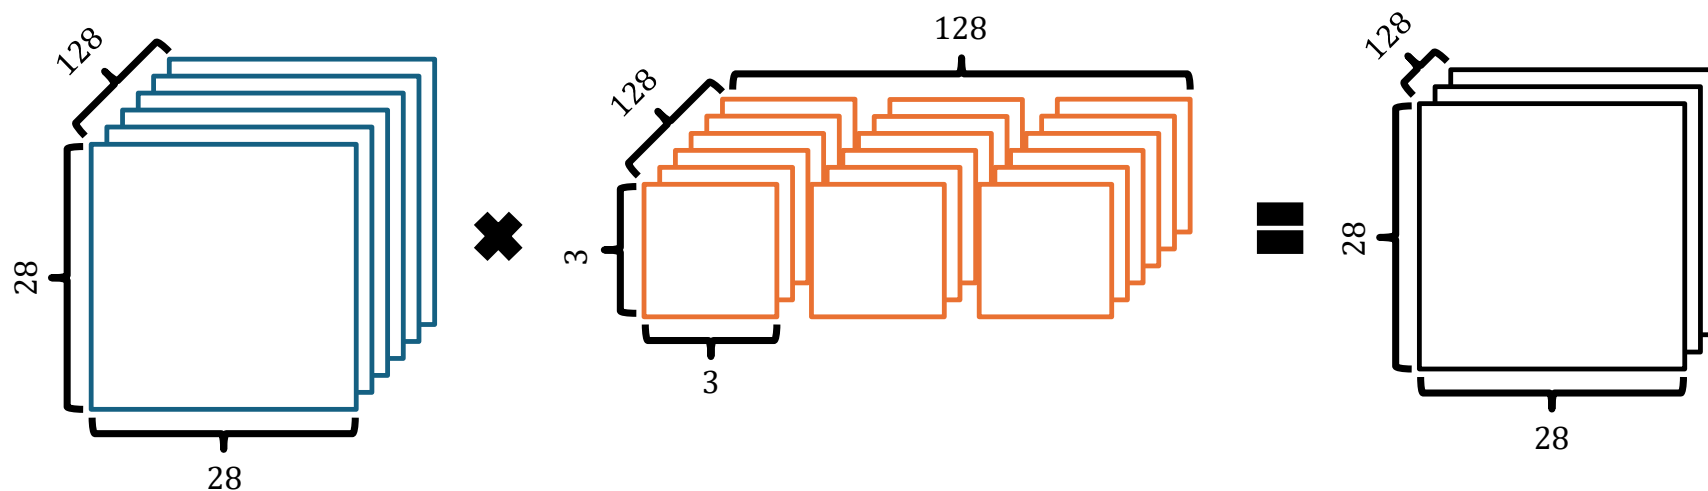
$$MEM = 0.10M + 0.15M = 0.5MB$$

Arithmetic intensity of A100 with FP16 :

**208 FLOPs / Bytes**

Can process at most 208 Operation every 1 Byte Read.

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

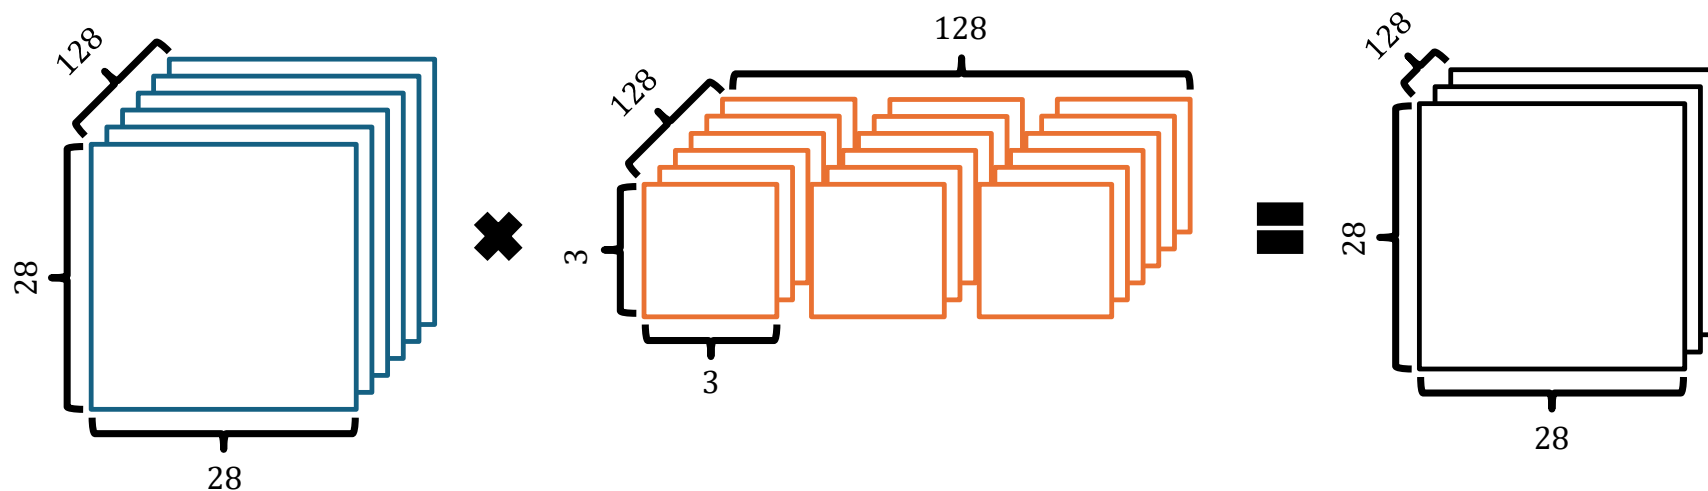
$$FLOPs = 115M$$

$$MEM = 0.10M + 0.15M = 0.5MB$$

Arithmetic intensity of A100 with FP16 :  
**208 FLOPs / Bytes**

Arithmetic intensity of example:  
**230 FLOPs / Bytes**

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

$$FLOPs = 115M$$

$$MEM = 0.10M + 0.15M = 0.5MB$$

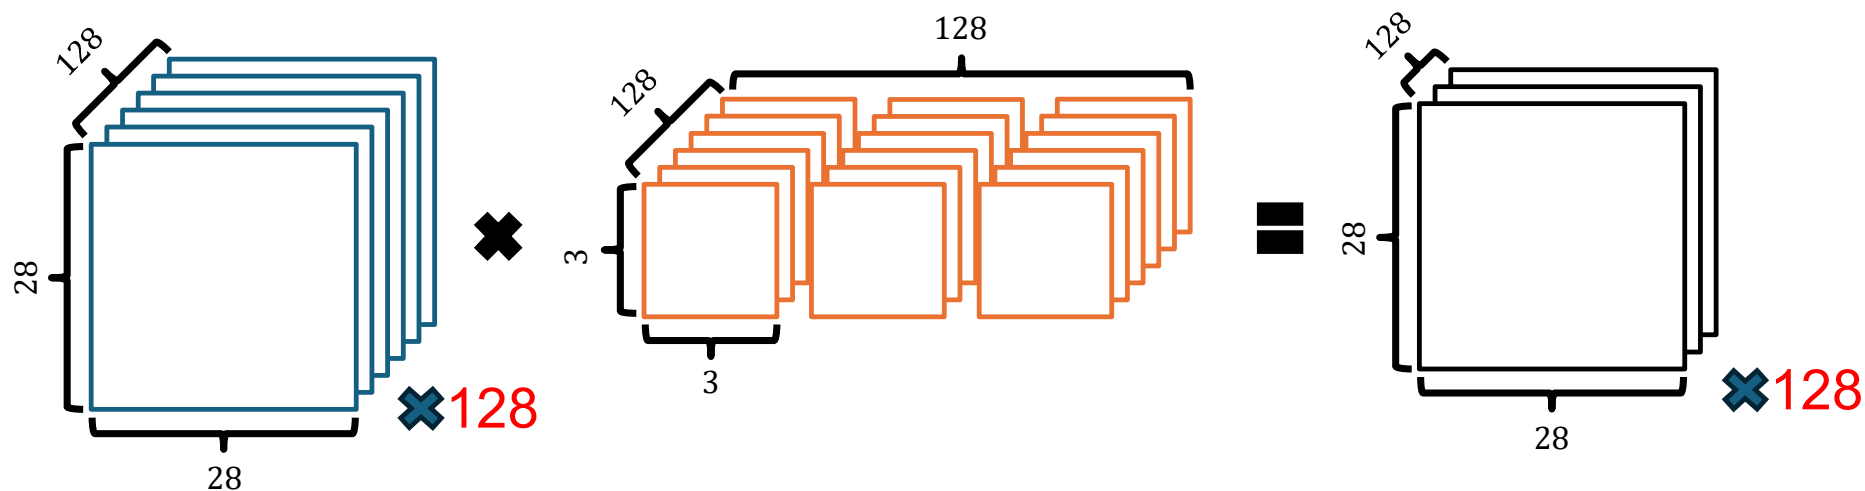
Arithmetic intensity of A100 with FP16 :  
**208 FLOPs / Bytes**

Arithmetic intensity of example:  
**230 FLOPs / Bytes**

**Going to suffer Compute bound!** ➡



# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

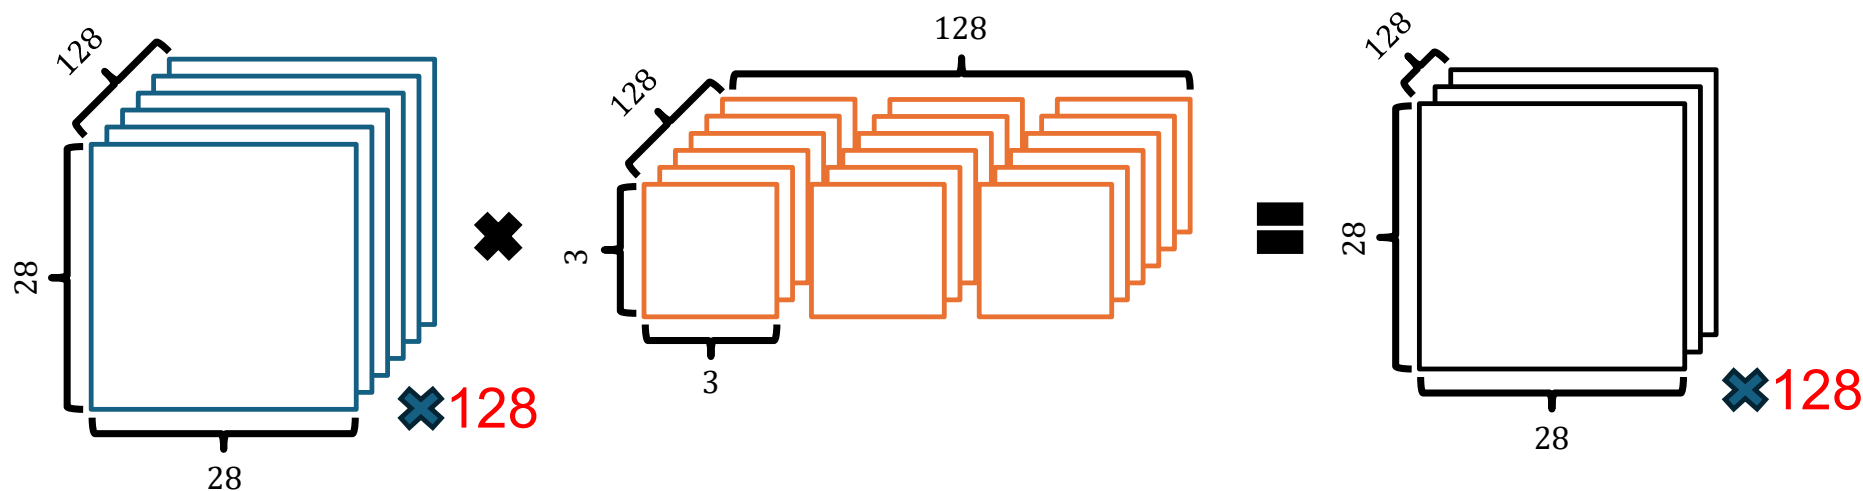
$$FLOPs = 115M \times 128 = 14.8B$$

$$MEM = 12.8M + 0.15M = 26.0MB$$

Arithmetic intensity of A100 with FP16 :  
**208 FLOPs / Bytes**

Arithmetic intensity with **128** batch:  
**569 FLOPs / Bytes**

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

$$FLOPs = 115M \times 128 = 14.8B$$

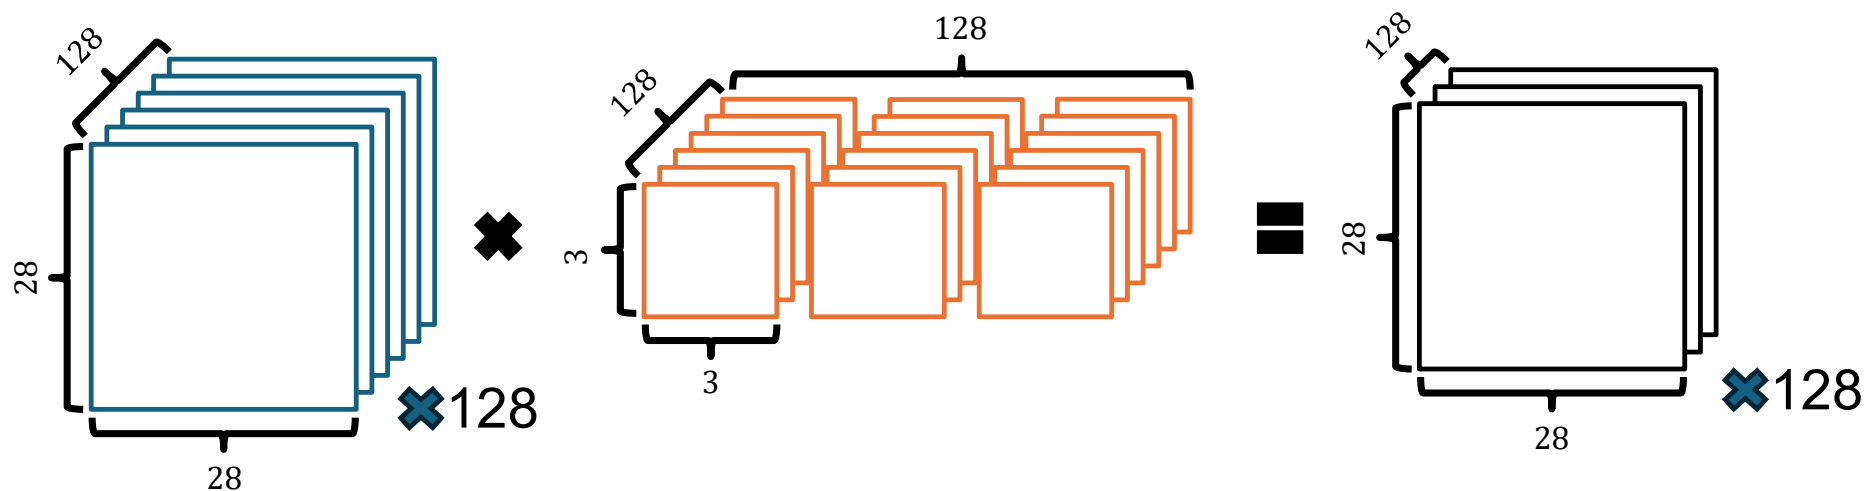
$$MEM = 12.8M + 0.15M = 26.0MB$$

Arithmetic intensity of A100 with FP16 :  
**208 FLOPs / Bytes**

Arithmetic intensity with 128 batch:

**Very Compute bound!** ➡ **569 FLOPs / Bytes**

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

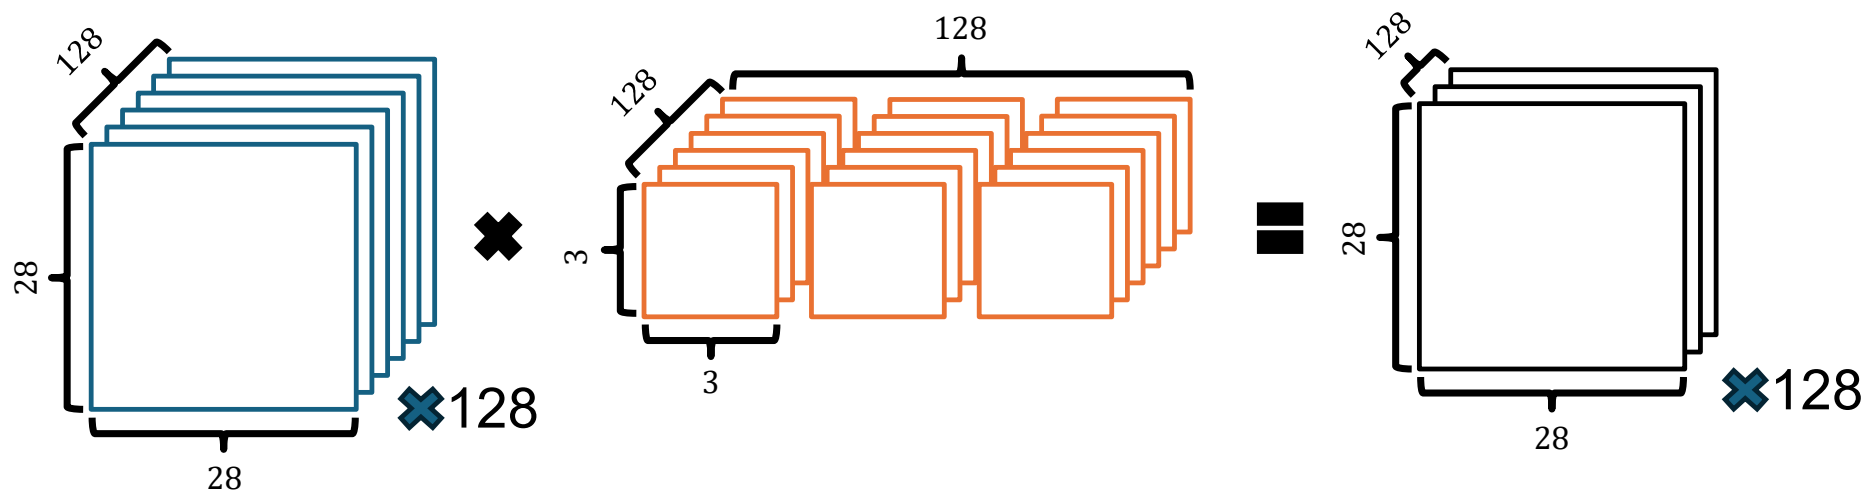
$$FLOPs = 115M \times 128 = 14.8B$$

$$MEM = 12.8M + 0.15M = 26.0MB$$

Arithmetic intensity of A100 with **INT8** :  
**832 OPs / Bytes**

Arithmetic intensity with 128 batch:  
**569 FLOPs / Bytes**

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

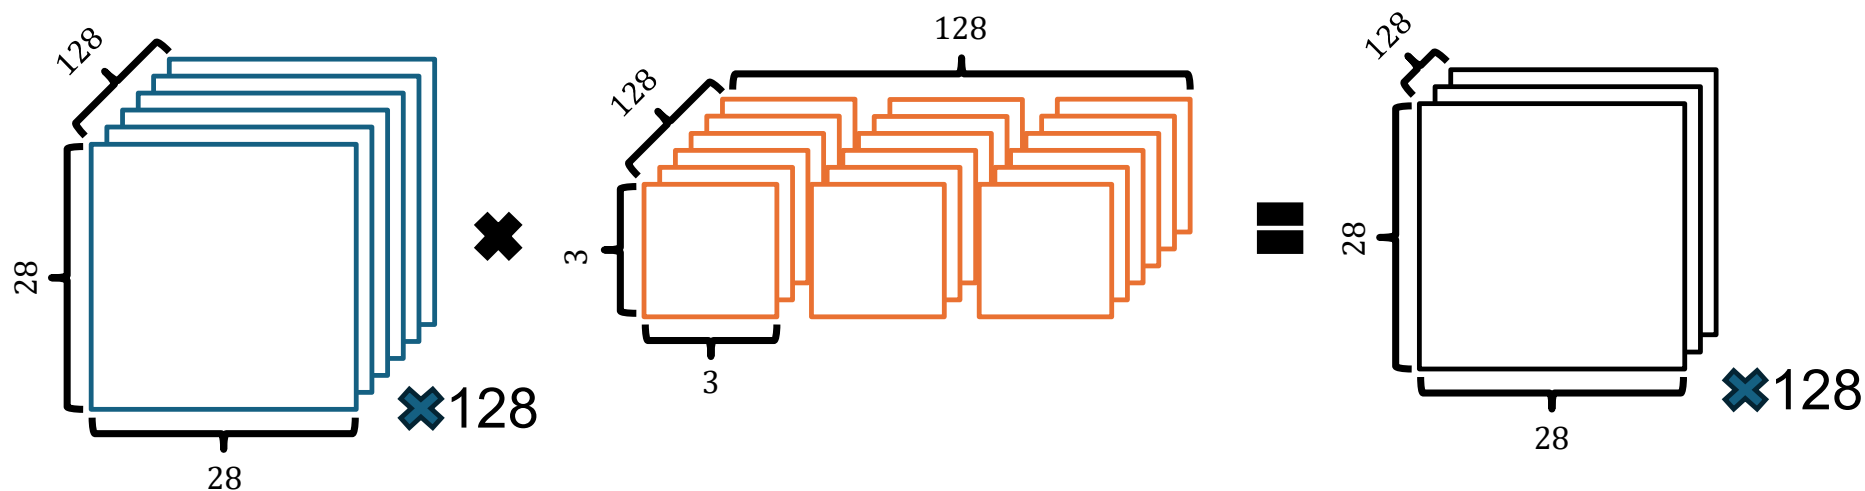
$$FLOPs = 115M \times 128 = 14.8B$$

$$MEM = 12.8M + 0.15M = 13.0MB$$

Arithmetic intensity of A100 with **INT8** :  
**832 OPs / Bytes**

Arithmetic intensity with 128 batch:  
**1139 OPs / Bytes**

# Preliminaries – Convolution operation



Example with ResNet50 (middle layer)

$$FLOPs = 115M \times 128 = 14.8B$$

$$MEM = 12.8M + 0.15M = 13.0MB$$

Arithmetic intensity of A100 with **INT8** :  
**832 OPs / Bytes**

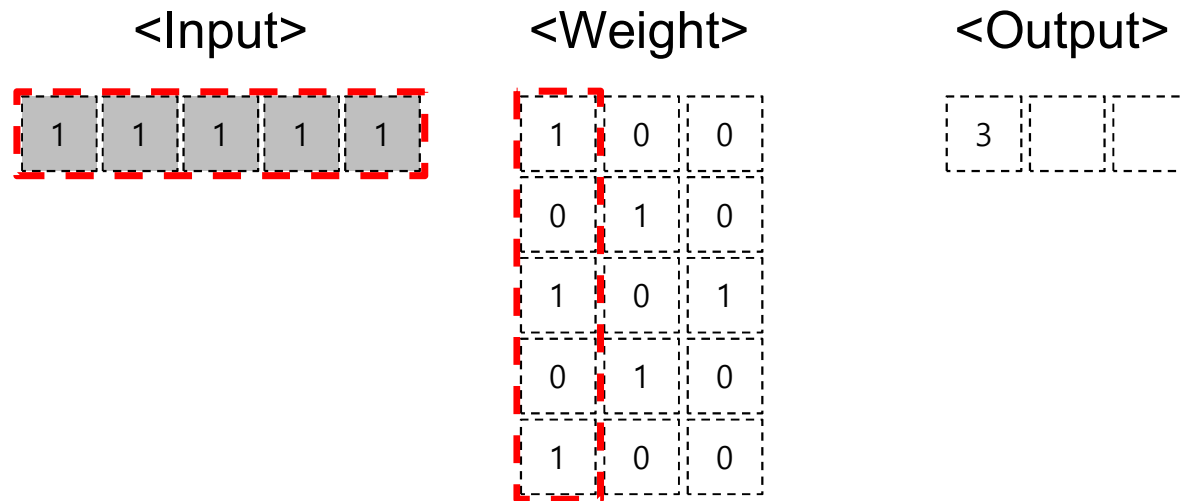
Arithmetic intensity with 128 batch:  
**1139 OPs / Bytes**

**We have chance to accelerate the layer with INT8 quantized computation!**

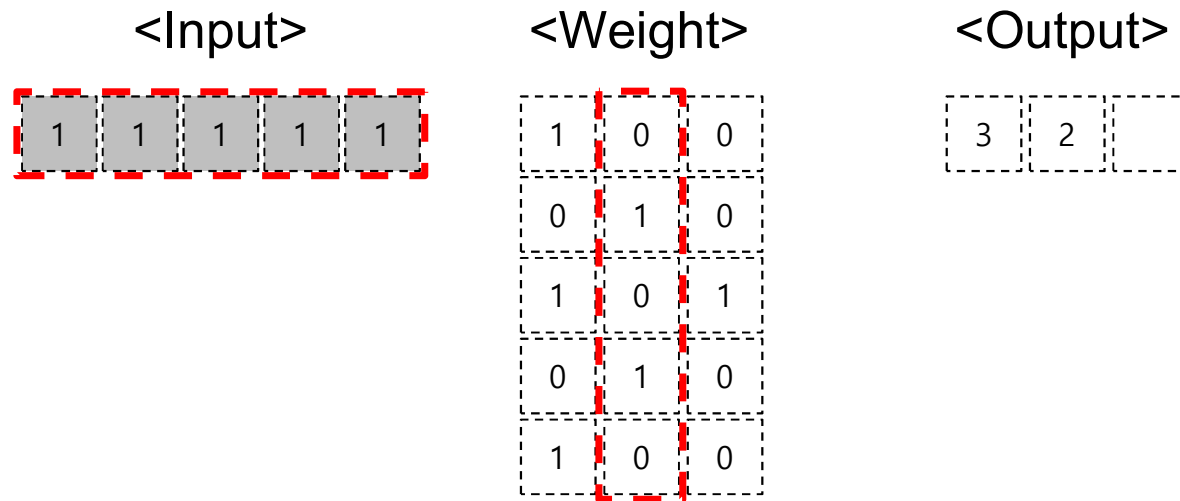
# Preliminaries – Linear operation

<Input>					<Weight>			<Output>		
1	1	1	1	1	1	0	0			
					0	1	0			
					1	0	1			
					0	1	0			
					1	0	0			

# Preliminaries – Linear operation

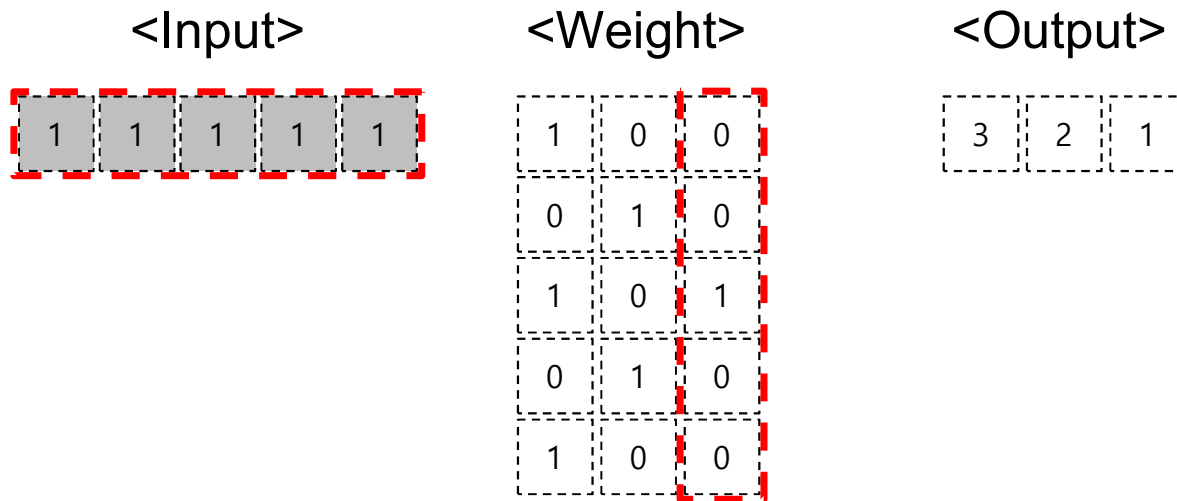


# Preliminaries – Linear operation

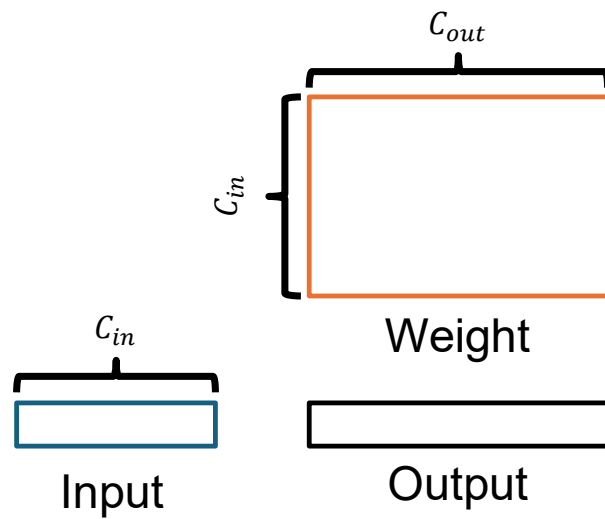




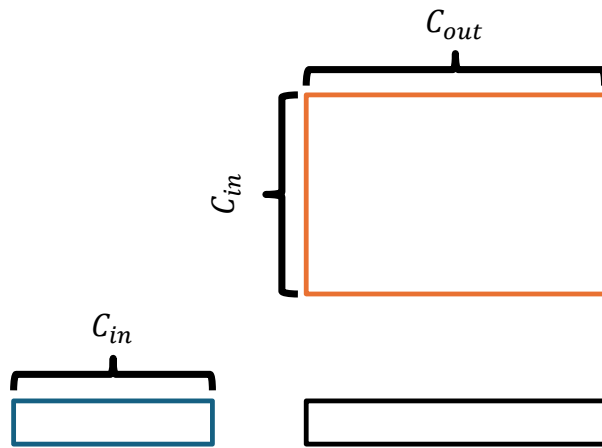
# Preliminaries – Linear operation



# Preliminaries – Linear operation

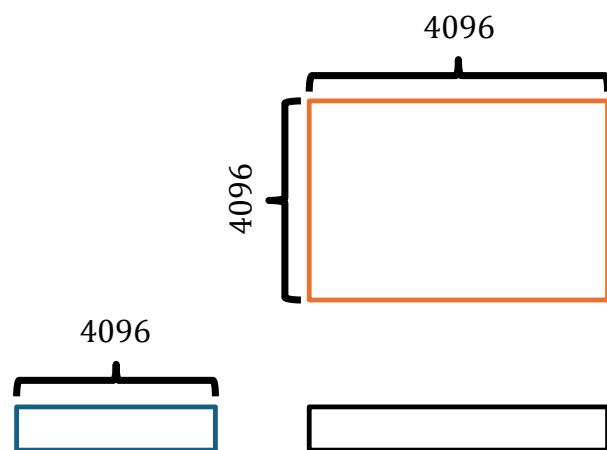


# Preliminaries – Linear operation



$$FLOPs = (2 * C_{in} - 1) \times C_{out}$$
$$MEM = C_{in} + [C_{in} \times C_{out}]$$

# Preliminaries – Linear operation

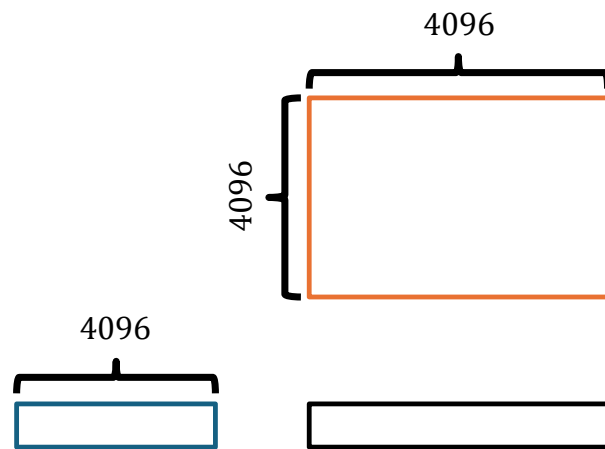


Example with AlexNet

$$FLOPs = (2 * 4096 - 1) \times 4096$$

$$MEM = 4096 + [4096 \times 4096]$$

# Preliminaries – Linear operation

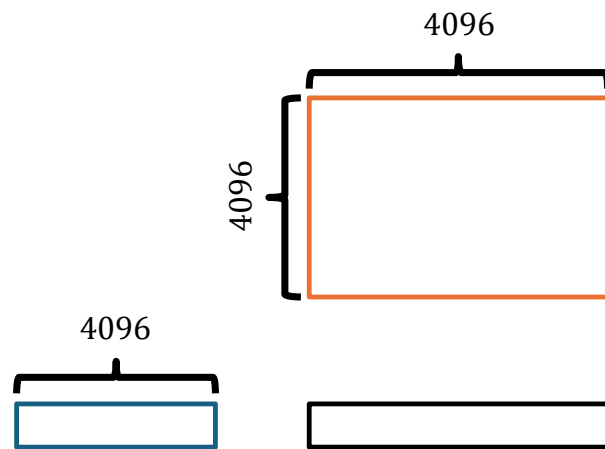


Example with AlexNet

$$FLOPs = 33.55M$$

$$MEM = 4K + 16.7M = 16.78M$$

# Preliminaries – Linear operation

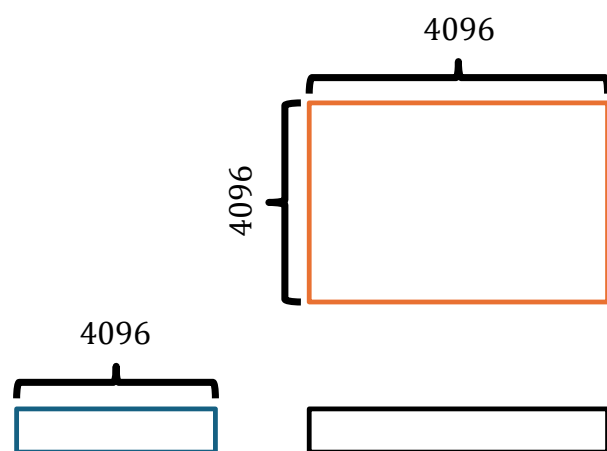


Example with AlexNet

$$FLOPs = 33.55M$$

$$MEM = 4K + 16.7M = 33.56 \text{ **FP16 MB**}$$

# Preliminaries – Linear operation



Example with AlexNet

$$FLOPs = 33.55M$$

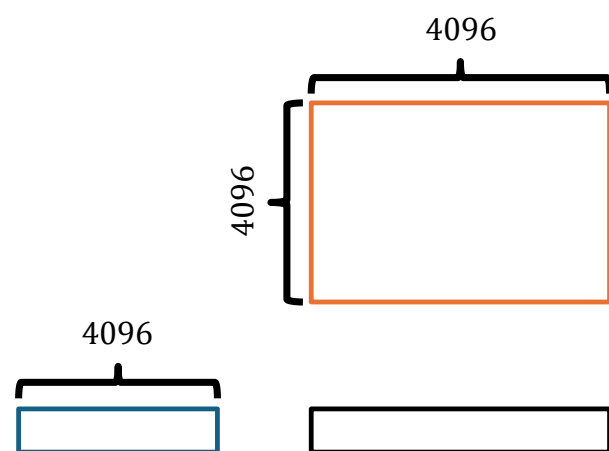
$$MEM = 4K + 16.7M = 33.56MB$$

Arithmetic intensity of example:

**1.0 FLOPs / Bytes**

Can process 1 Operation every 1 Byte Read.

# Preliminaries – Linear operation



Example with AlexNet

$$FLOPs = 33.55M$$

$$MEM = 4K + 16.7M = 33.56MB$$

Arithmetic intensity of example:

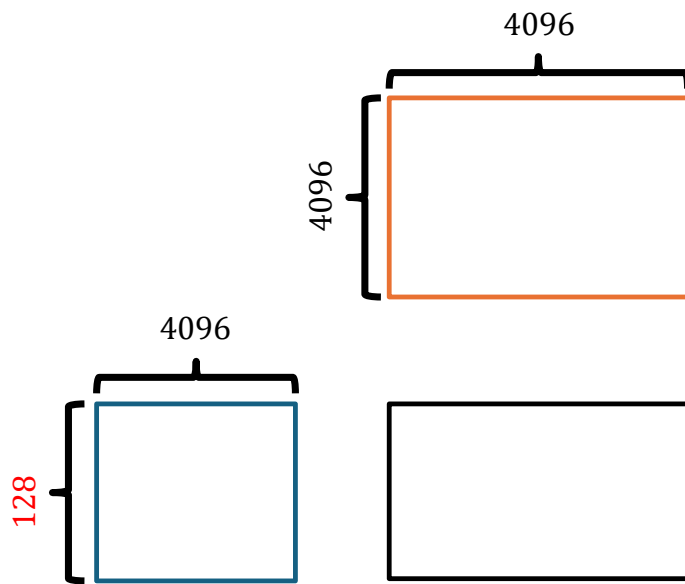
**1.0 FLOPs / Bytes**



**Going to strongly suffer memory bound!**



# Preliminaries – Linear operation



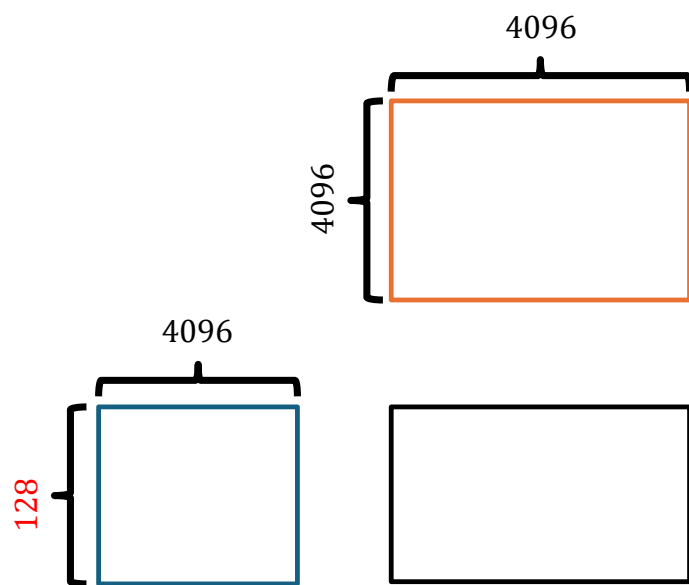
Example with AlexNet

$$FLOPs = 33.55M \times 128 = 4.3B$$

$$MEM = 0.5M + 16.7M = 34.60MB$$

Arithmetic intensity with 128 batch :  
**124.1 FLOPs / Bytes**

# Preliminaries – Linear operation



Example with AlexNet

$$FLOPs = 33.55M \times 128 = 4.3B$$

$$MEM = 0.5M + 16.7M = 34.60MB$$

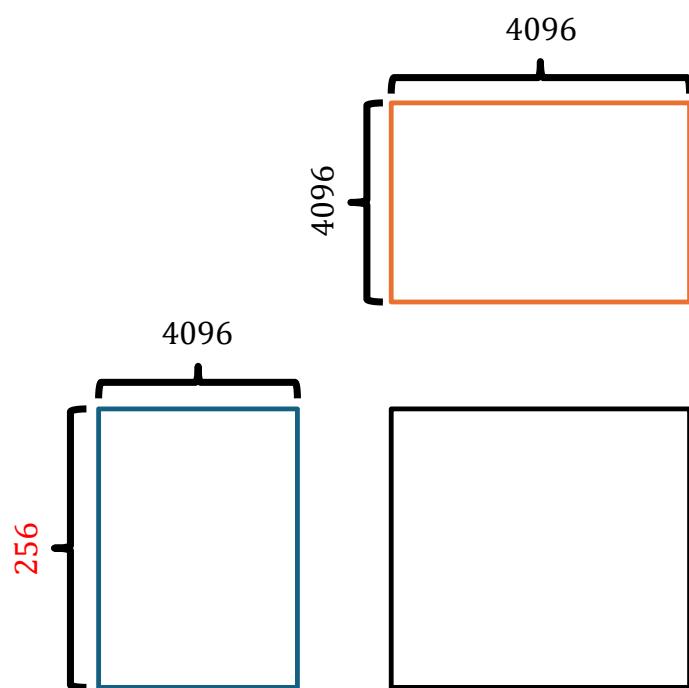
Arithmetic intensity with 128 batch :

**124.1 FLOPs / Bytes**



**Still not enough!**

# Preliminaries – Linear operation



Example with AlexNet

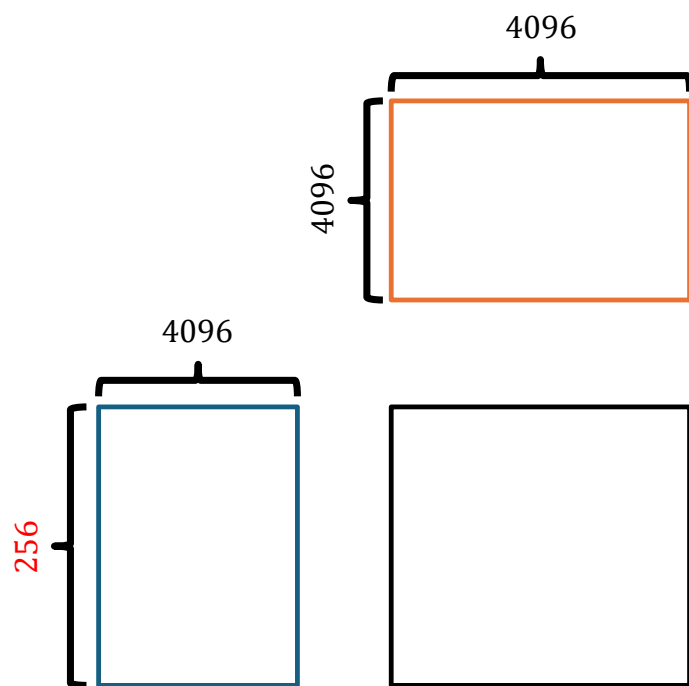
$$FLOPs = 33.55M \times 256 = 8.6B$$

$$MEM = 1.0M + 16.7M = 35.65MB$$

Arithmetic intensity with 256 batch :

**240.91 FLOPs / Bytes**

# Preliminaries – Linear operation



Example with AlexNet

$$FLOPs = 33.55M \times 256 = 8.6B$$

$$MEM = 1.0M + 16.7M = 35.65MB$$

Arithmetic intensity with 256 batch :

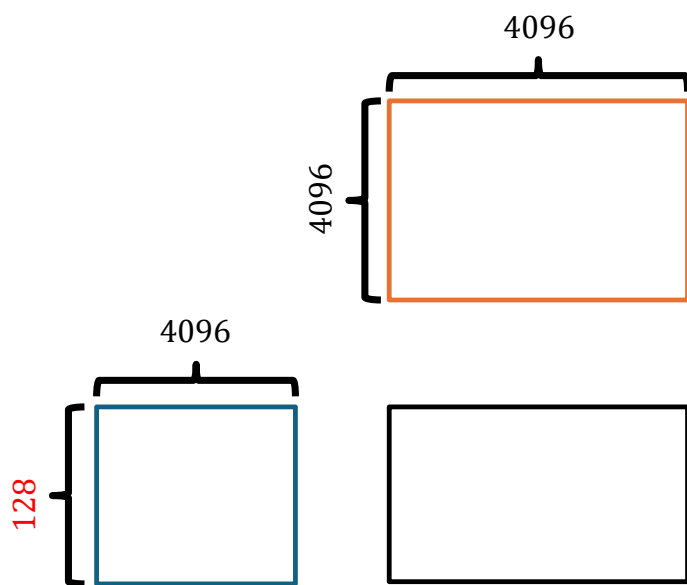
**240.91 FLOPs / Bytes**



**Now compute bound!**

**But batch size might be too big for inference.**

# Preliminaries – Linear operation



Example with AlexNet

If we quantize weight to INT8 with 128 batch

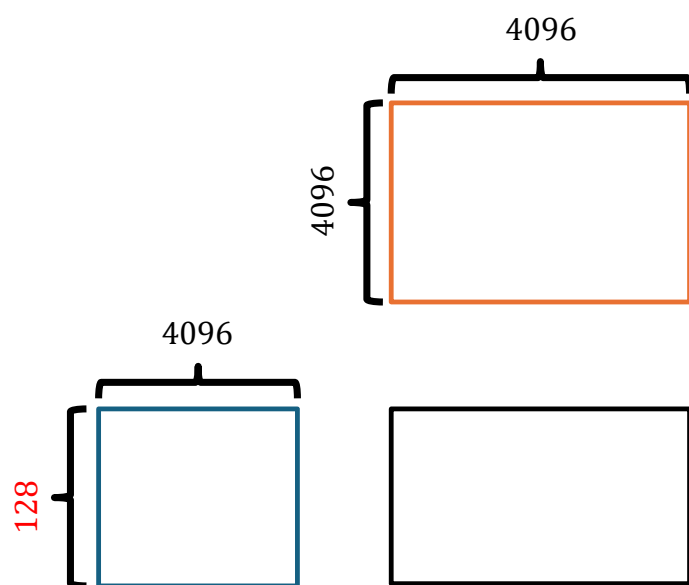
$$FLOPs = 33.55M \times 128 = 4.3B$$

$$MEM = 0.5M + 16.7M = 34.60MB$$

Arithmetic intensity with 128 batch :

**124.1 FLOPs / Bytes**

# Preliminaries – Linear operation



Example with AlexNet

If we quantize weight to INT8 with 128 batch

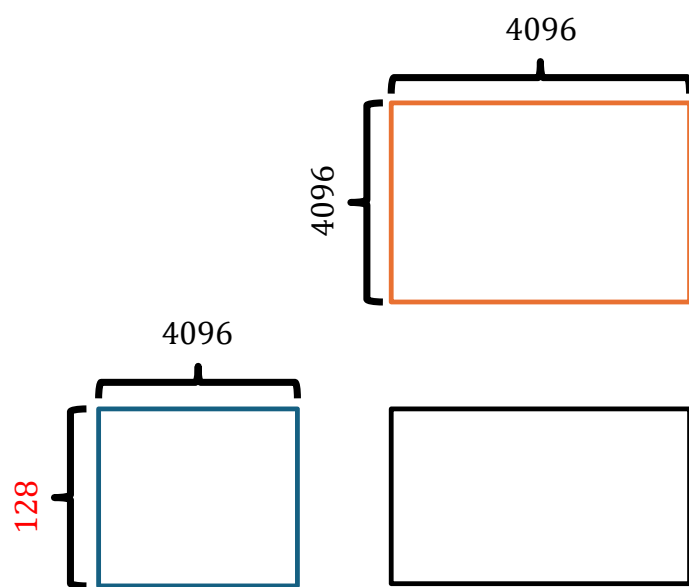
$$FLOPs = 33.55M \times 128 = 4.3B$$

$$MEM = 0.5M + 16.7M = 17.8MB$$

Arithmetic intensity with 128 batch :

**240.91 FLOPs / Bytes**

# Preliminaries – Linear operation



Example with AlexNet

If we quantize weight to INT8 with 128 batch

$$FLOPs = 33.55M \times 128 = 4.3B$$

$$MEM = 0.5M + 16.7M = 17.8MB$$

Arithmetic intensity with 128 batch :

**240.91 FLOPs / Bytes**



Lowering memory pressure can increase arithmetic intensity,  
thus accelerate the layer!

# Preliminaries

- Major data types used in Deep Learning
  - **Integer**
  - **Floating Point**



# Preliminaries – Data Type

- Major data types used in Deep Learning

- **Integer**

- Unsigned Integer:  $[0, 2^n - 1]$

0	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

$2^5 + 2^4 + 2^0 = 49$

- Signed Integer:  $[-2^{n-1}, 2^{n-1} - 1]$

Sign Integer

1	0	1	1	0	0	0	1
---	---	---	---	---	---	---	---

$-2^7 + 2^5 + 2^4 + 2^0 = -79$

# Preliminaries – Data Type

- Major data types used in Deep Learning
  - **Floating Point**
    - **Numerical Form:**  $(-1)^s M 2^E$ 
      - **s (Sign bit)** determines whether number is negative or positive
      - **M (Significand)** normally a fractional value in range [1.0,2.0)
      - **E (Exponent)** weights value by power of 2

# Preliminaries – Data Type

- Major data types used in Deep Learning
  - **Floating Point**
    - **Numerical Form:**  $(-1)^s M 2^E$ 
      - **s (Sign bit)** determines whether number is negative or positive
      - **M (Significand)** normally a fractional value in range  $[1.0, 2.0)$
      - **E (Exponent)** weights value by power of 2
    - **Encoding**
      - **s** is sign bit **s**
      - **exp** field encodes **E** (but is not equal to **E**)
      - **frac** field encodes **M** (but is not equal to **M**)



# Preliminaries – Data Type

- Major data types used in Deep Learning
  - **Floating Point**
    - **Single Precision:** 32 bits  
 $\approx 7$  decimal digits, max/min:  $10^{\pm 38}$



# Preliminaries – Data Type

$$- V = (-1)^s M 2^E$$

- Major data types used in Deep Learning
  - **Floating Point**
    - **Exponent** coded as a biased value:  $E = \text{exp} - \text{Bias}$ 
      - **exp**: unsigned value of **exp** field
      - $\text{Bias} = 2^{k-1} - 1$ , where  $k$  is number of exponent bits
        - Single precision: 127 (**exp**: 1...254, E: -126...127)



# Preliminaries – Data Type

$$- V = (-1)^S M 2^E$$

- Major data types used in Deep Learning
  - **Floating Point**
    - **Significand** coded with implied leading 1:  $M = 1.xxx...x_2$ 
      - $xxx...x$ : bits of **frac** field
      - Get extra leading bit for “free”



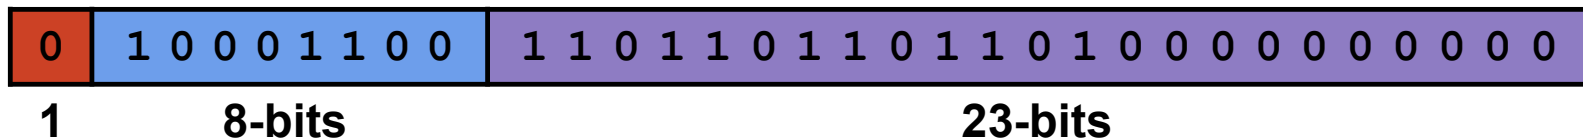
# Preliminaries – Data Type

- $V = (-1)^S M 2^E$
- $E = \text{exp} - \text{Bias}$

- Major data types used in Deep Learning

- **Floating Point** – Example (float32)

- **Value** =  $15213.0 = 1.1101101101101_2 \times 2^{13}$
- **Significand**
  - $M = 1.\underline{1101101101101}_2$
  - $\text{frac} = \underline{1101101101101}0000000000_2$
- **Exponent**
  - $E = 13$
  - $\text{Bias} = 127$
  - $\text{exp} = 140 = 10001100_2$
- **Result**



# Preliminaries – Data Type

- Major data types used in Deep Learning
  - **Floating Point** – Example

- IEEE 754 Single Precision 32-bit Float (IEEE FP32)



- Google Brain Float (BF16)



- IEEE 754 Half Precision 16-bit Float (IEEE FP16)



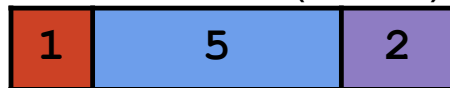


# Preliminaries – Data Type

- Major data types used in Deep Learning
  - **Floating Point** – Example
    - IEEE 754 Half Precision 16-bit Float (IEEE FP16)



- NVIDIA FP8 (E5M2)



- NVIDIA FP8 (E4M3)



# Break

# Quantization Basics

- Fundamentals
- Scheme
- Type

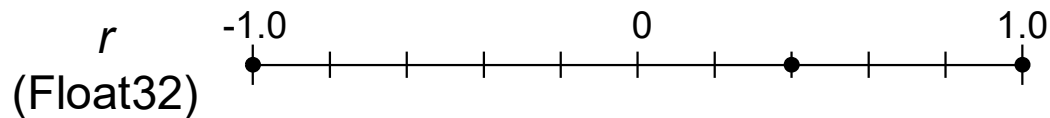
# Quantization Basics – Fundamentals

- **Example** (3-Bit Uniform Quantization)

# Quantization Basics – Fundamentals

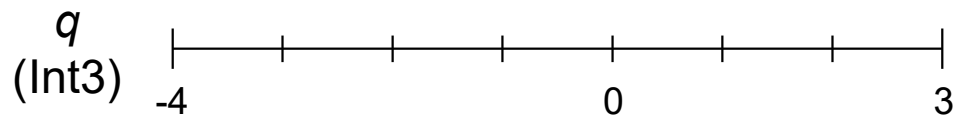
## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



$$• \quad S = (Max_r - Min_r) / (2^n - 1)$$

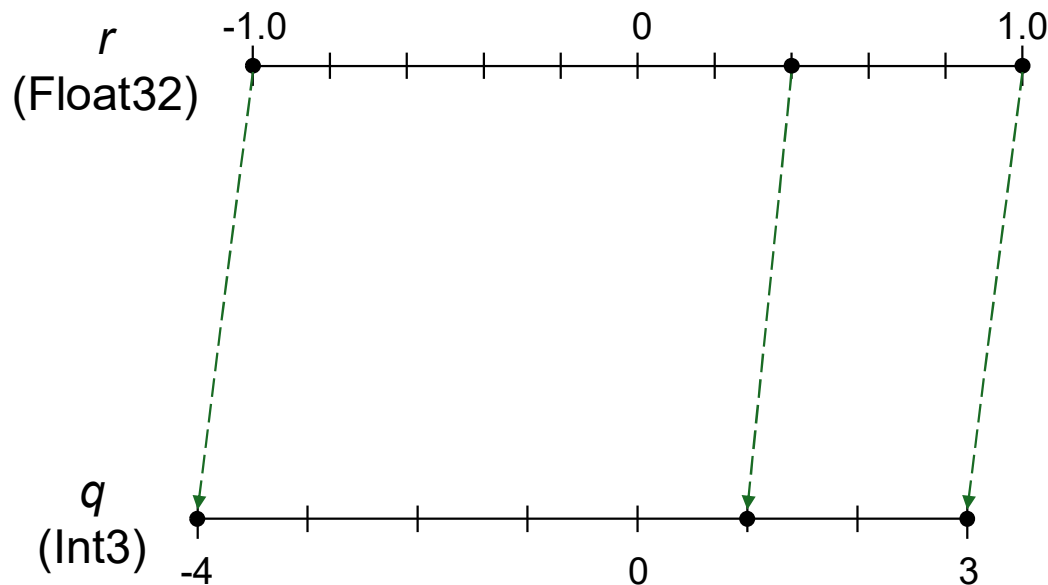
$$• \quad q = [r / S]$$



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



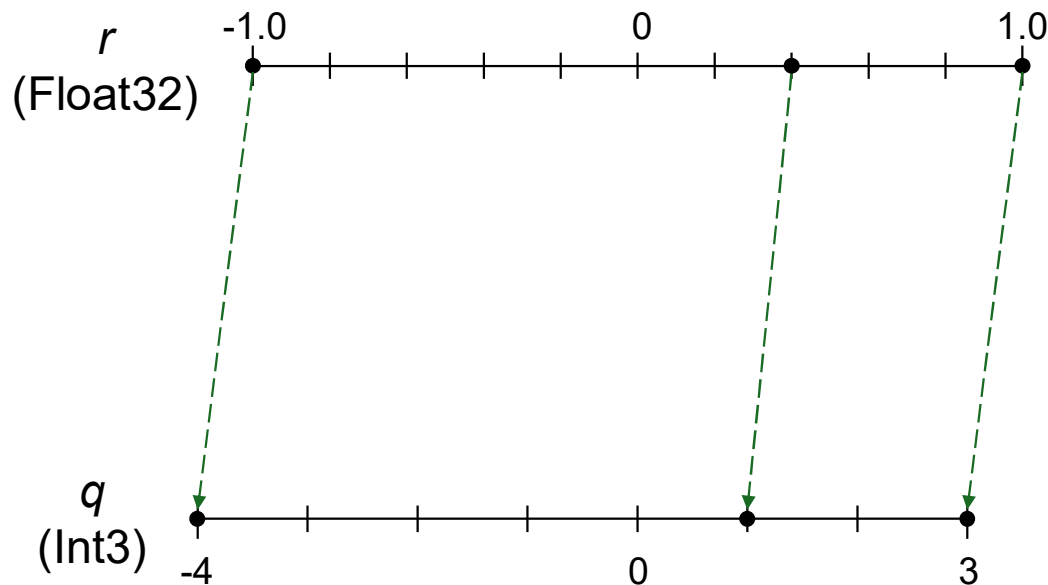
$$• S = (Max_r - Min_r) / (2^n - 1)$$

$$• q = [r / S]$$

# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



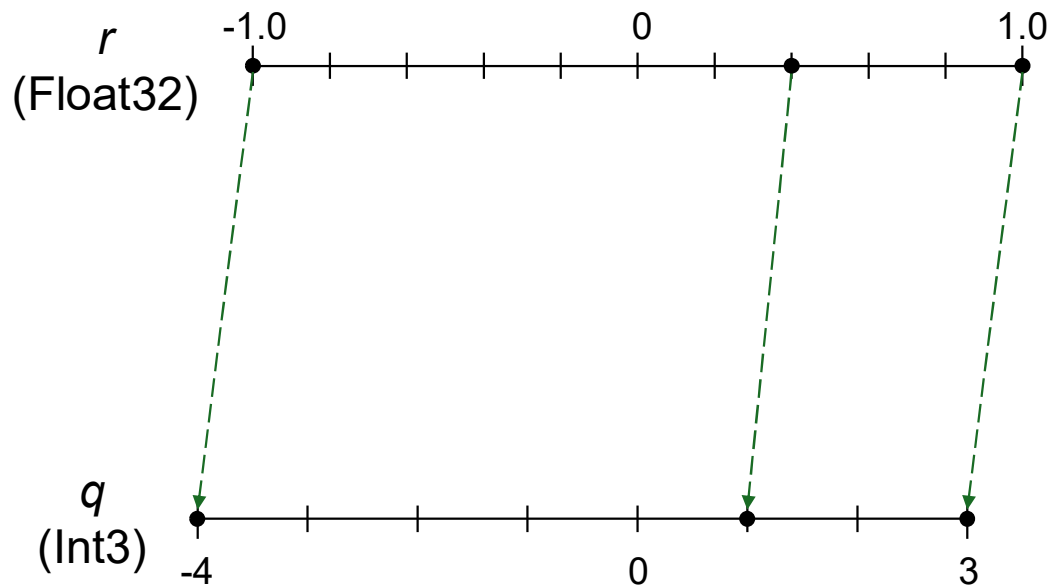
$$• S = (1.0 - (-1.0)) / (2^3 - 1)$$

$$• q = [r / S]$$

# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



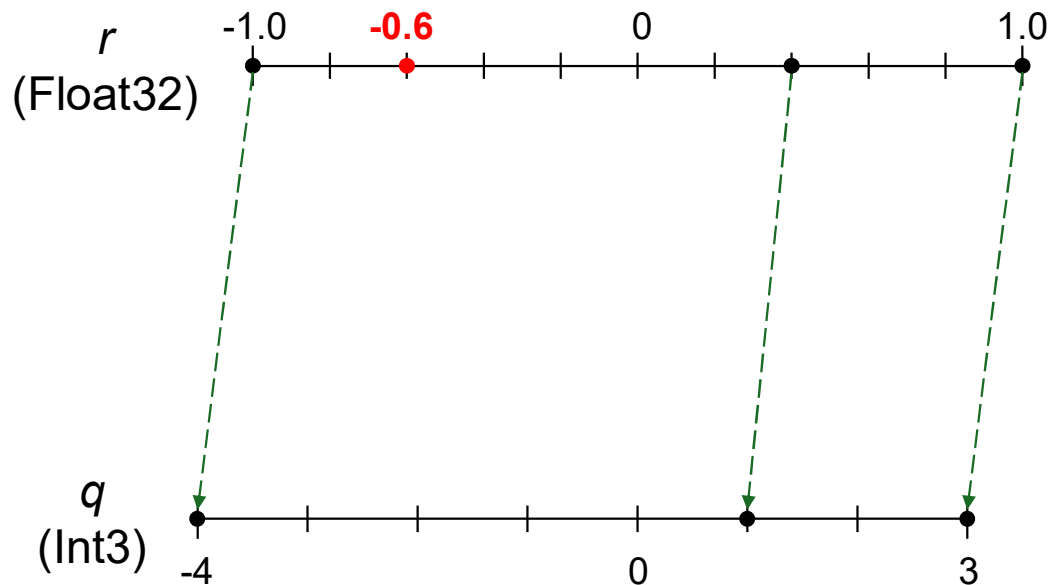
- $S = (1.0 - (-1.0)) / (2^3 - 1)$   
 $= 0.28571$
- $q = [r / S]$



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor

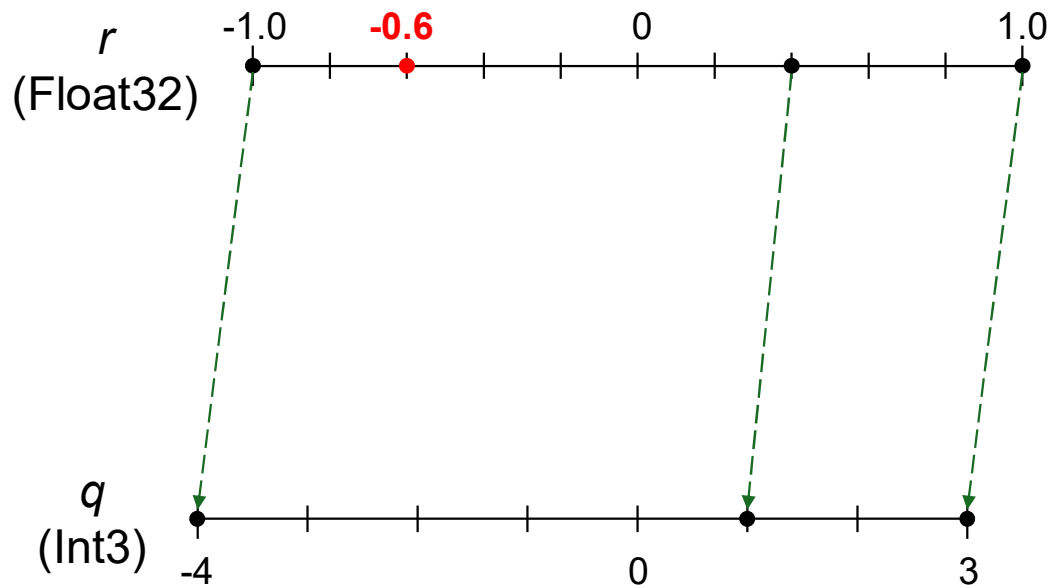


- $S = (1.0 - (-1.0)) / (2^3 - 1)$   
 $= 0.28571$
- $q = [r / S]$

# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor

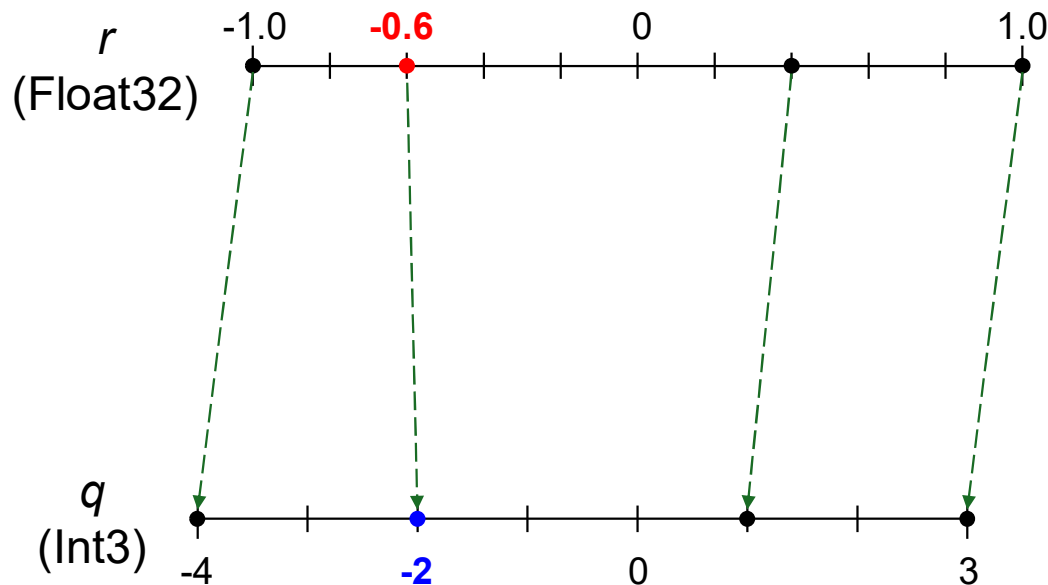


- $S = (1.0 - (-1.0)) / (2^3 - 1)$   
 $= 0.28571$
- $q = \lfloor -0.6 / 0.28571 \rfloor$

# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor

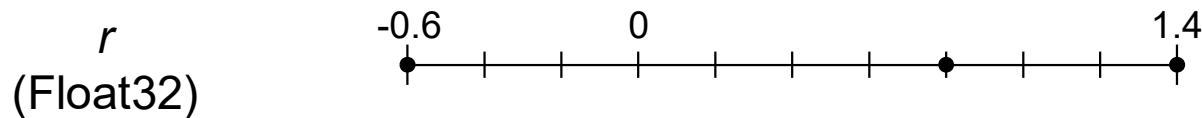


- $S = (1.0 - (-1.0)) / (2^3 - 1)$   
 $= 0.28571$
- $q = \lfloor -0.6 / 0.28571 \rfloor$   
 $= -2$

# Quantization Basics – Fundamentals

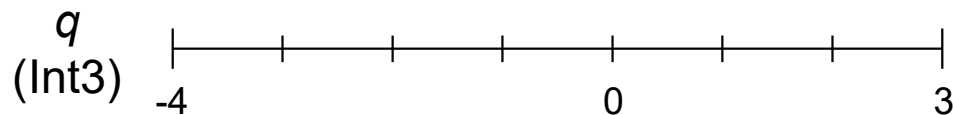
## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[\ ]$  : round
- $S$  : scaling factor



$$\bullet \quad S = (Max_r - Min_r) / (2^n - 1)$$

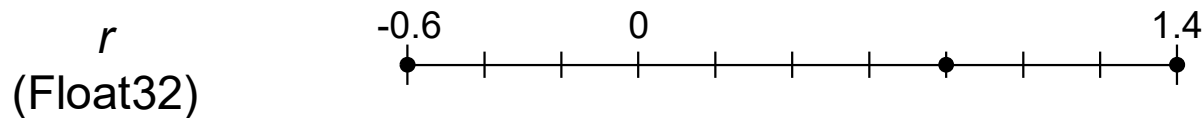
$$\bullet \quad q = [r / S]$$



# Quantization Basics – Fundamentals

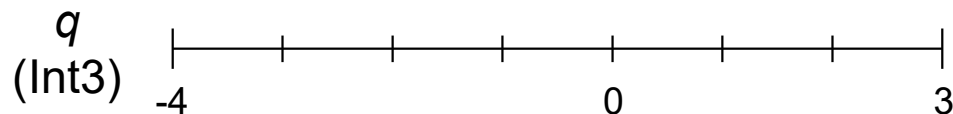
## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



$$• \quad S = (1.4 - (-0.6)) / (2^n - 1)$$

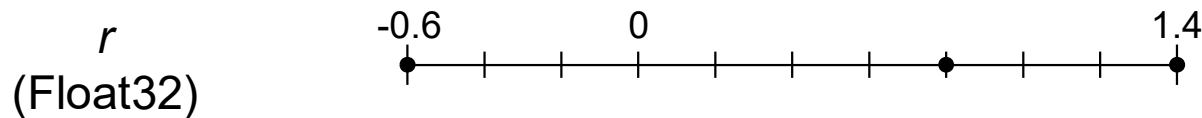
$$• \quad q = [r / S]$$



# Quantization Basics – Fundamentals

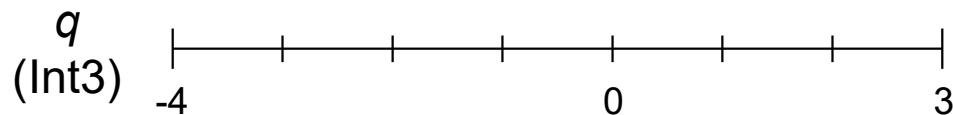
## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



$$\begin{aligned} \bullet \quad S &= (1.4 - (-0.6)) / (2^n - 1) \\ &= 0.28571 \end{aligned}$$

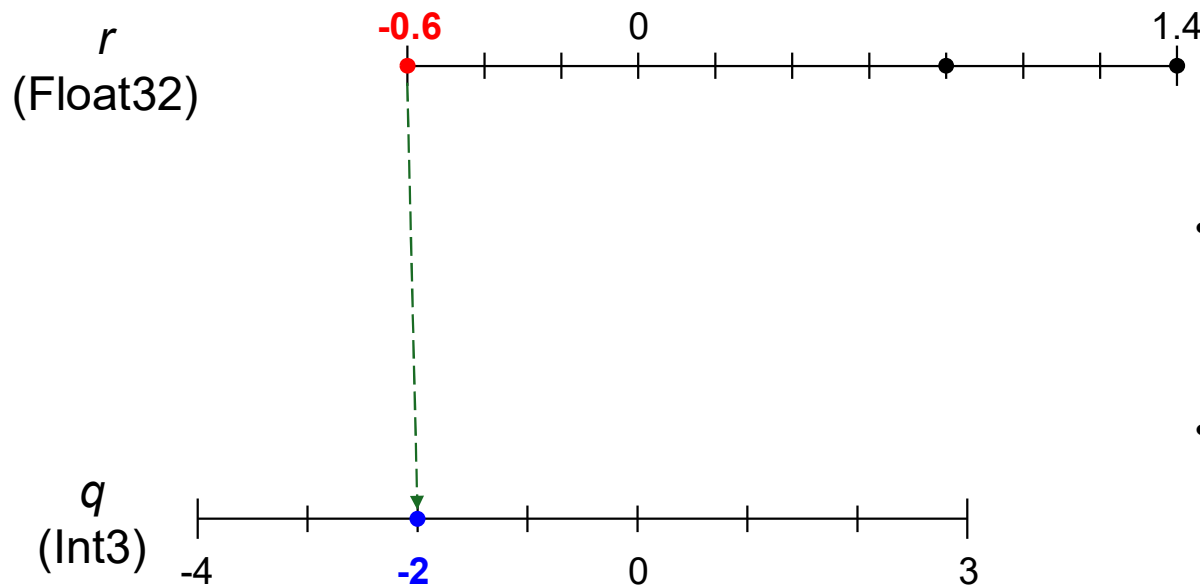
$$\bullet \quad q = [r / S]$$



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

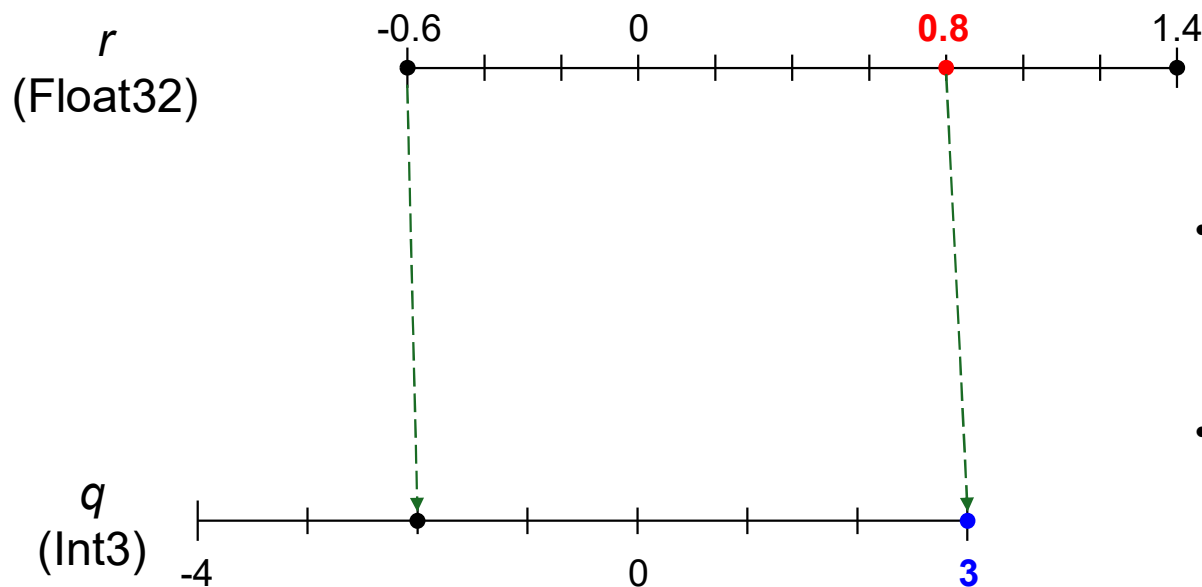
- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



- $S = (1.4 - (-0.6)) / (2^n - 1)$   
 $= 0.28571$
- $q = \lfloor -0.6 / 0.28571 \rfloor$   
 $= -2$

# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)



- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor

$$\begin{aligned} \bullet \quad S &= (1.4 - (-0.6)) / (2^n - 1) \\ &= 0.28571 \end{aligned}$$

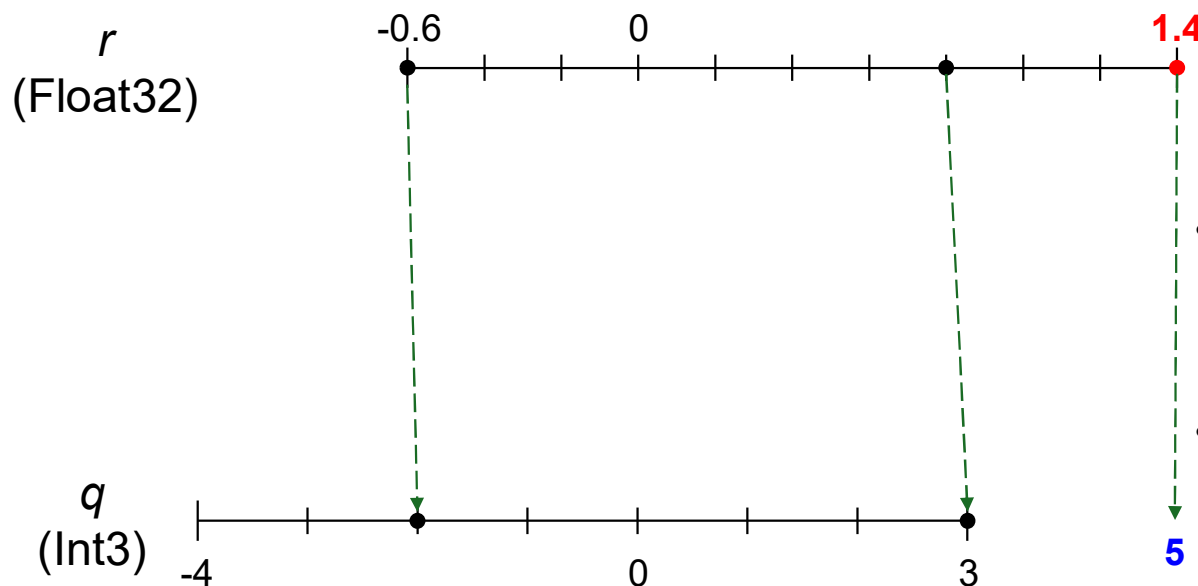
$$\begin{aligned} \bullet \quad q &= [0.8 / 0.28571] \\ &= 3 \end{aligned}$$



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



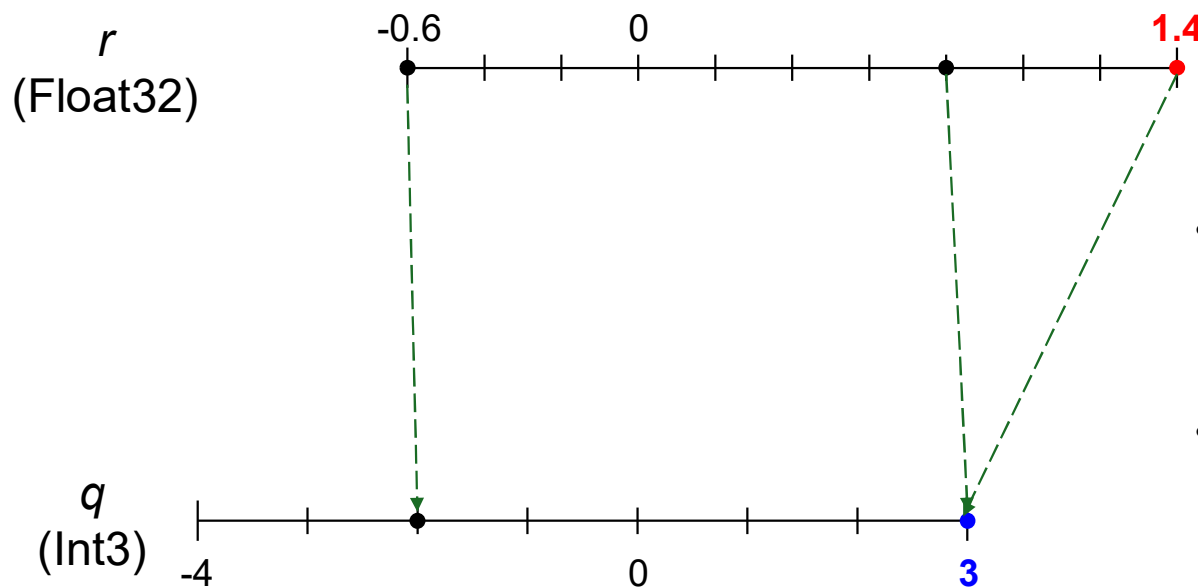
$$\begin{aligned} \bullet \quad S &= (1.4 - (-0.6)) / (2^n - 1) \\ &= 0.28571 \end{aligned}$$

$$\begin{aligned} \bullet \quad q &= \lfloor 1.4 / 0.28571 \rfloor \\ &= 5 \end{aligned}$$

# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor

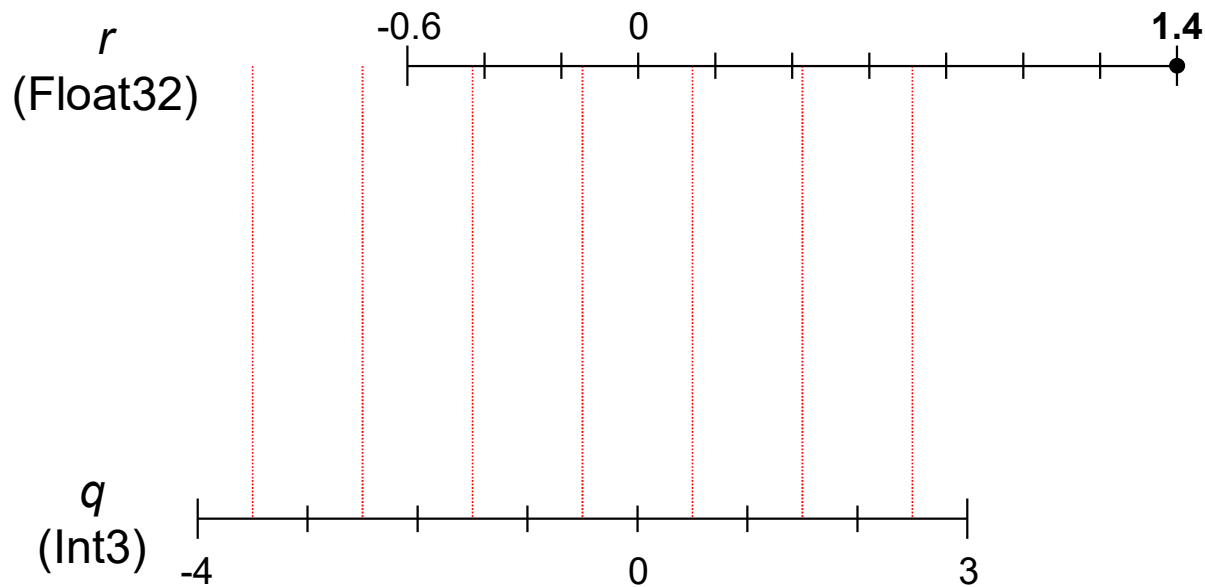


- $S = (1.4 - (-0.6)) / (2^n - 1)$   
 $= 0.28571$
- $q = \text{clamp}(\lfloor 1.4 / 0.28571 \rfloor)$   
 $= 3$

# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

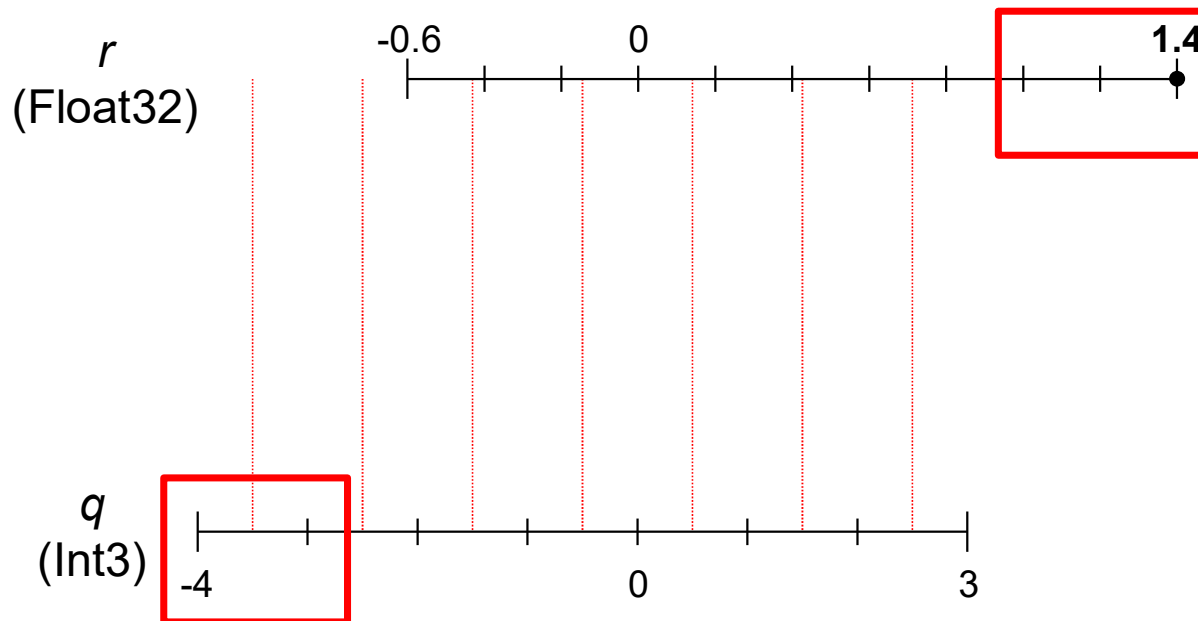
- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

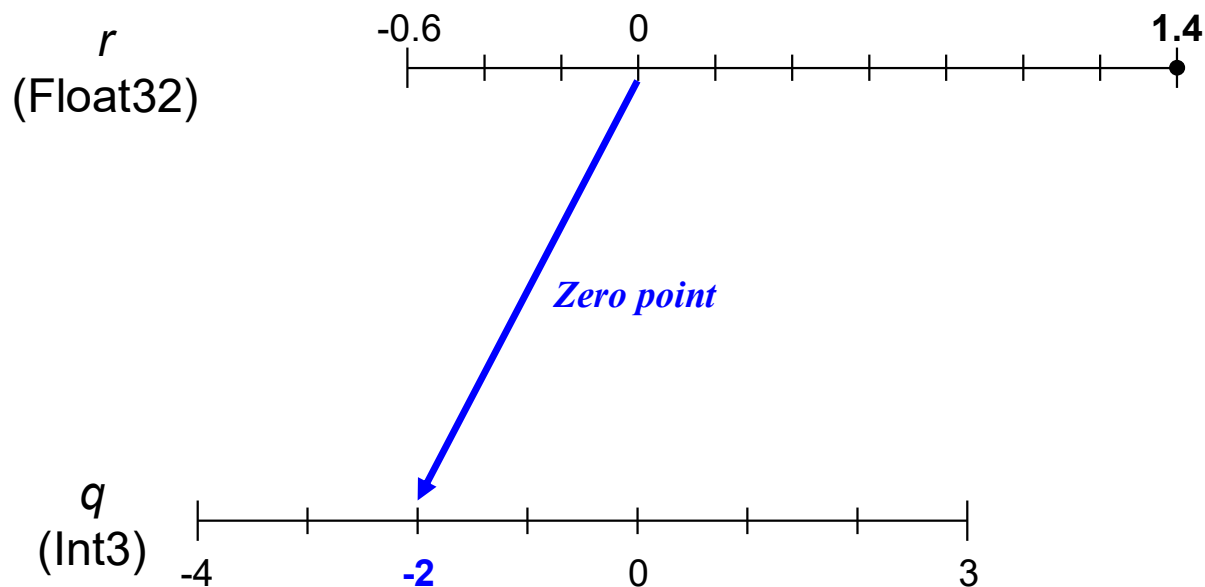
- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

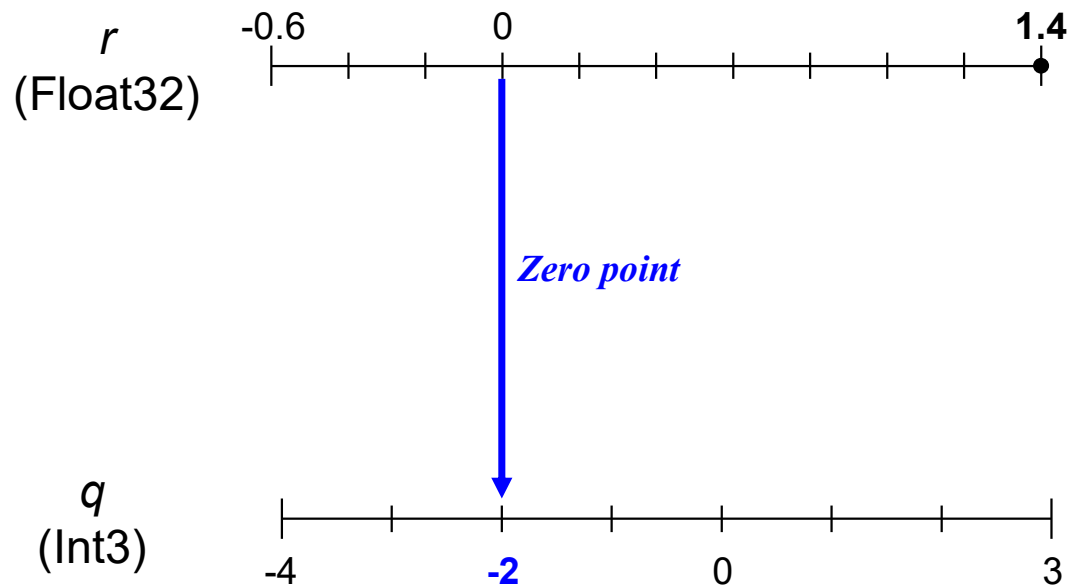
- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)

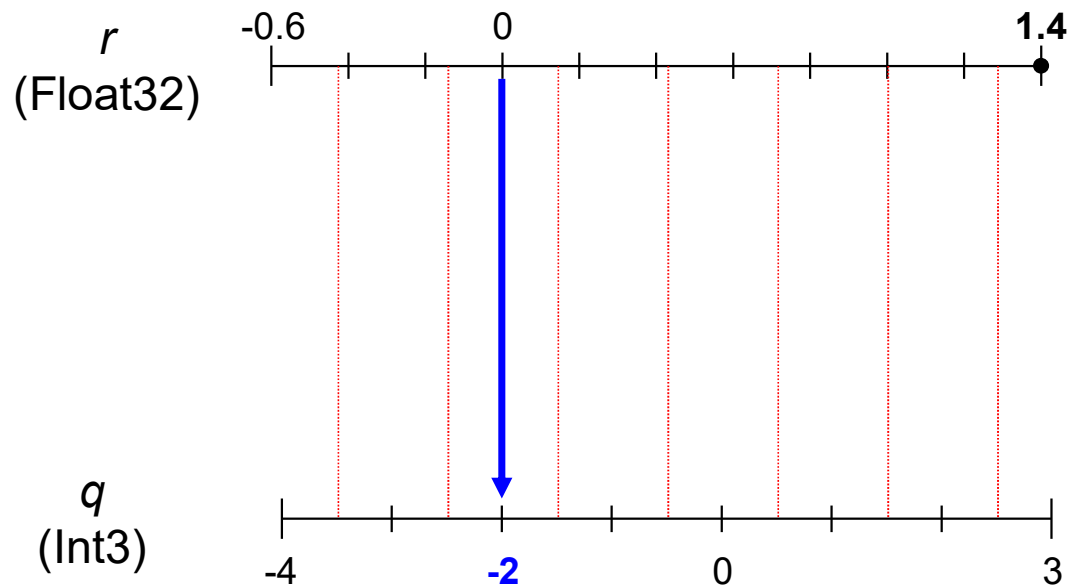
- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



# Quantization Basics – Fundamentals

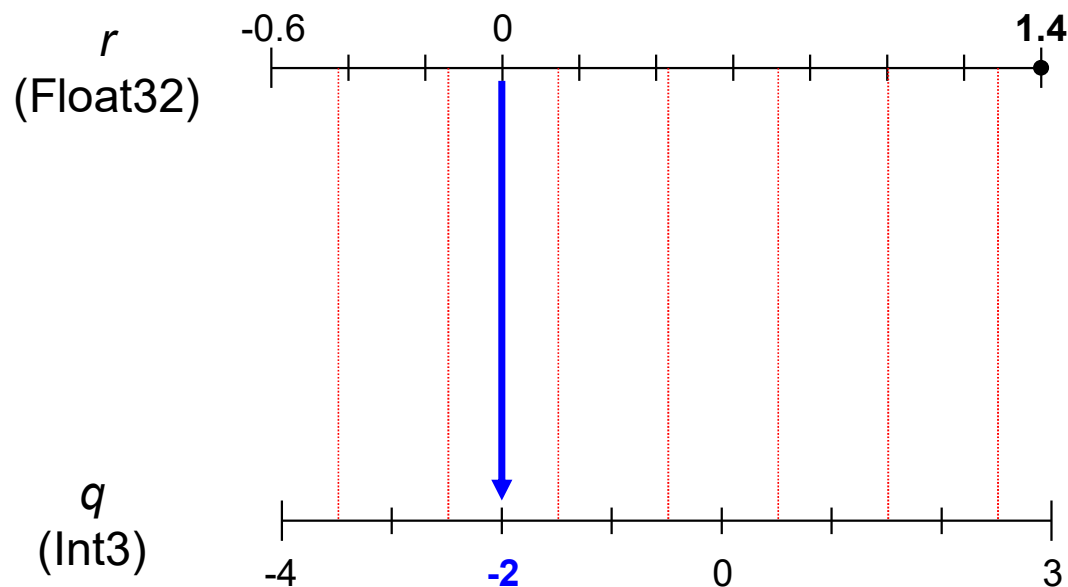
## • **Example** (3-Bit Uniform Quantization)

- $r$  : real value
- $q$  : quantized value
- $[]$  : round
- $S$  : scaling factor



# Quantization Basics – Fundamentals

## • **Example** (3-Bit Uniform Quantization)



- $r$  : real value
- $q$  : quantized value
- $[\ ]$  : round
- $S$  : scaling factor
- $Z$  : zero-point

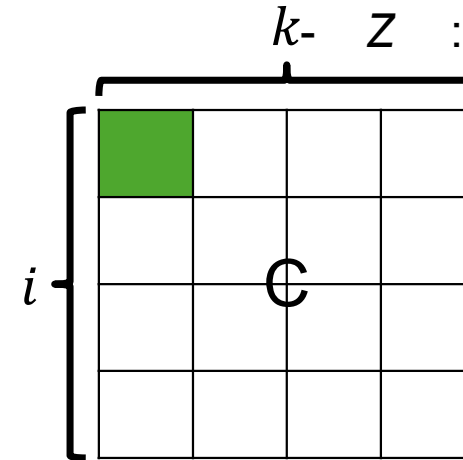
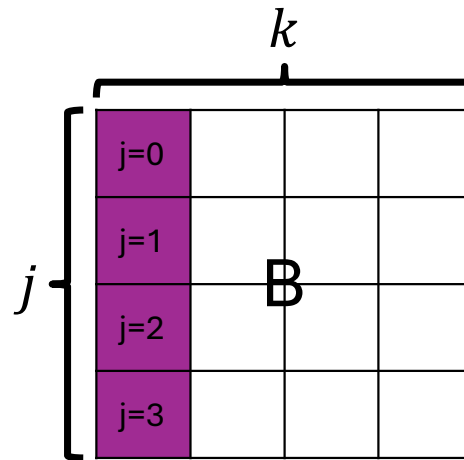
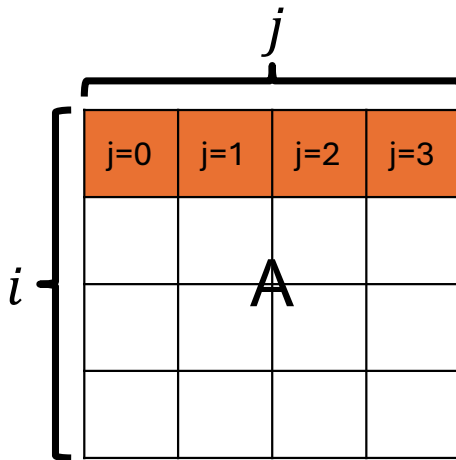
- $S = (Max_r - Min_r) / (2^n - 1)$
- $Z = Min_q - (Min_r / S)$
- $q = [r / S] + Z$



# Quantization Basics – Fundamentals

## Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point

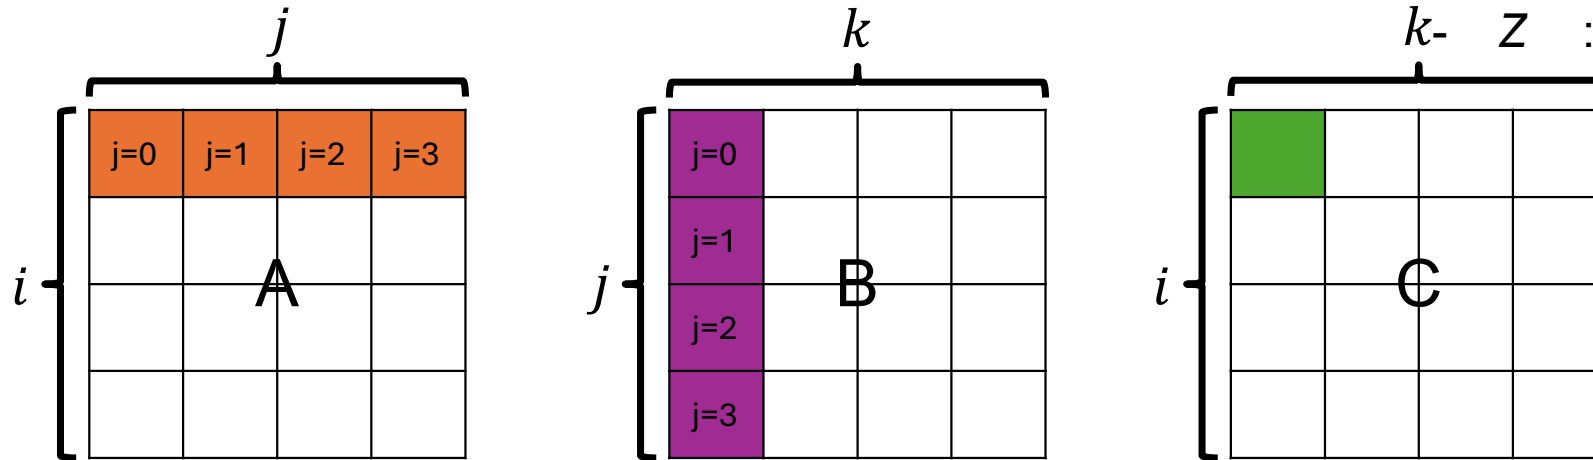


$$r_C^{(i,k)} = \sum_j r_A^{(i,j)} \times r_B^{(j,k)}$$

# Quantization Basics – Fundamentals

## Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point



$$S_C \left( q_C^{(i,k)} - Z_C \right) = \sum_j S_A \left( q_A^{(i,j)} - Z_A \right) S_B \left( q_B^{(j,k)} - Z_B \right)$$

# Quantization Basics – Fundamentals

## • Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point

$$S_C \left( q_C^{(i,k)} - Z_C \right) = \sum_j S_A \left( q_A^{(i,j)} - Z_A \right) S_B \left( q_B^{(j,k)} - Z_B \right)$$

# Quantization Basics – Fundamentals

## • Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point

$$S_C \left( q_C^{(i,k)} - Z_C \right) = \sum_j S_A \left( q_A^{(i,j)} - Z_A \right) S_B \left( q_B^{(j,k)} - Z_B \right)$$



$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \sum_j \left( q_A^{(i,j)} - Z_A \right) \left( q_B^{(j,k)} - Z_B \right)$$

# Quantization Basics – Fundamentals

## • Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point

$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \sum_j \left( q_A^{(i,j)} - Z_A \right) \left( q_B^{(j,k)} - Z_B \right)$$

# Quantization Basics – Fundamentals

## • Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point

$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \sum_j \left( q_A^{(i,j)} - Z_A \right) \left( q_B^{(j,k)} - Z_B \right)$$



$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \left[ N Z_A Z_B - Z_B \sum q_A^{(i,j)} - Z_A \sum q_B^{(j,k)} + \sum q_A^{(i,j)} q_B^{(j,k)} \right]$$

# Quantization Basics – Fundamentals

## • Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point

$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \sum_j \left( q_A^{(i,j)} - Z_A \right) \left( q_B^{(j,k)} - Z_B \right)$$



$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \left[ N Z_A Z_B - Z_B \sum q_A^{(i,j)} - Z_A \sum q_B^{(j,k)} + \sum q_A^{(i,j)} q_B^{(j,k)} \right]$$

Integer Matmul Operation

# Quantization Basics – Fundamentals

## • Quantized Matrix Multiplication

- $r$  : real value
- $q$  : quantized value
- $S$  : scaling factor
- $Z$  : zero point

$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \sum_j \left( q_A^{(i,j)} - Z_A \right) \left( q_B^{(j,k)} - Z_B \right)$$



$$q_C^{(i,k)} = Z_C + \frac{S_A S_B}{S_C} \left[ \underline{N Z_A Z_B - Z_B \sum q_A^{(i,j)} - Z_A \sum q_B^{(j,k)}} + \sum q_A^{(i,j)} q_B^{(j,k)} \right]$$

- If we don't use **Zero Point**, we can get rid of a lot of computations



# Quantization Basics – Fundamentals

- **Asymmetric Quantization**

- Use Zero Point
- Relatively Slow

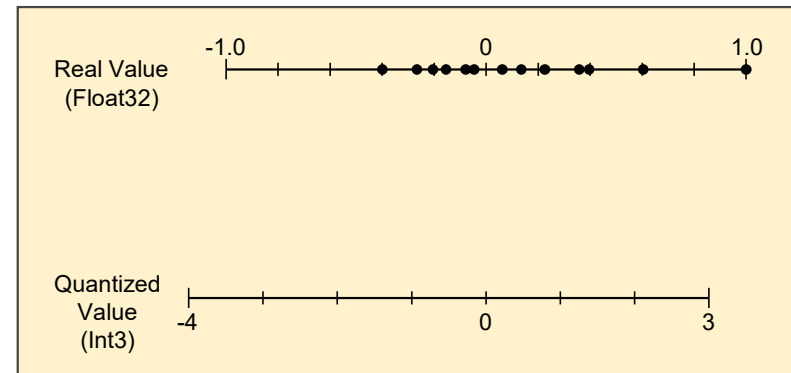
- **Symmetric Quantization**

- Do not use Zero Point
- Relatively Fast

# Quantization Basics – Fundamentals

- **Asymmetric Quantization**

- Use Zero Point
- Relatively Slow



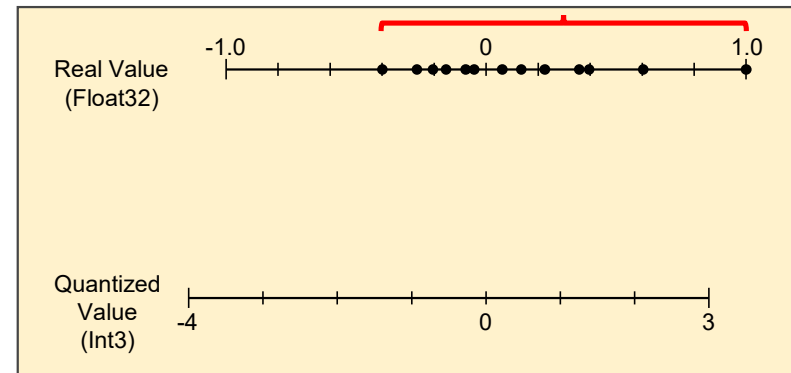
- **Symmetric Quantization**

- Do not use Zero Point
- Relatively Fast

# Quantization Basics – Fundamentals

- **Asymmetric Quantization**

- Use Zero Point
- Relatively Slow



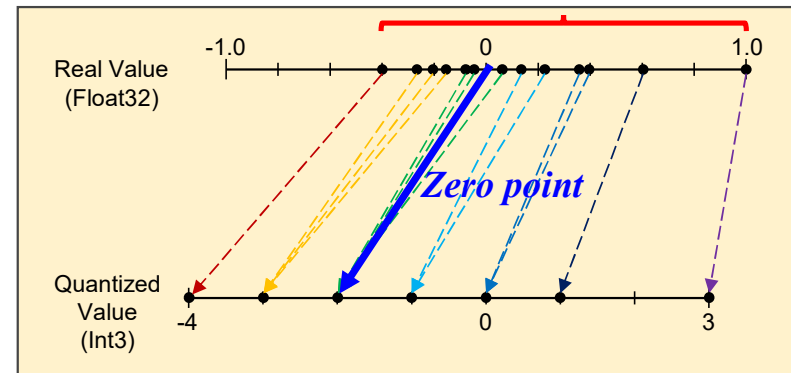
- **Symmetric Quantization**

- Do not use Zero Point
- Relatively Fast

# Quantization Basics – Fundamentals

- **Asymmetric Quantization**

- Use Zero Point
- Relatively Slow

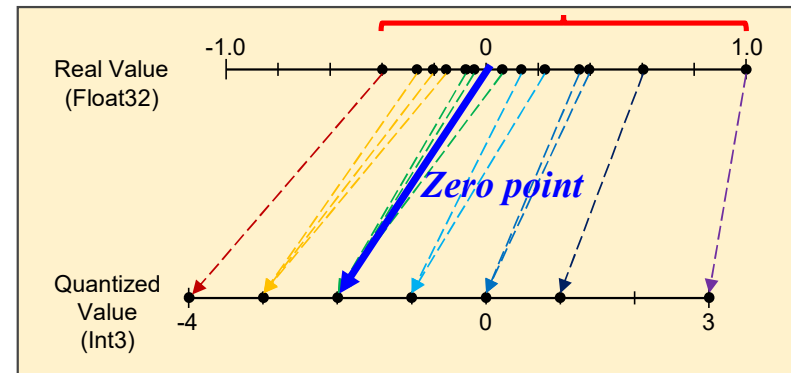


- **Symmetric Quantization**

- Do not use Zero Point
- Relatively Fast

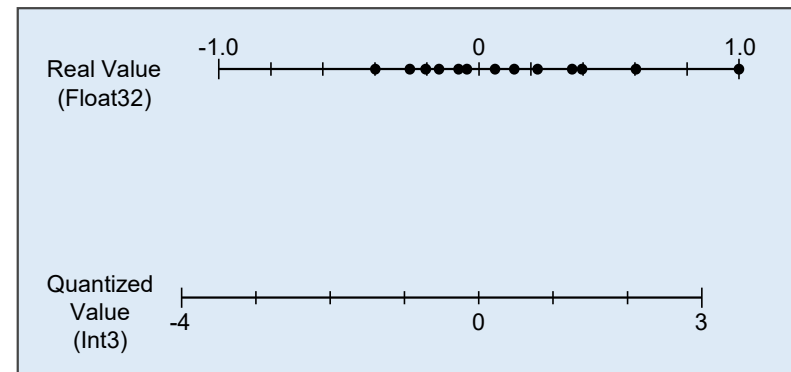
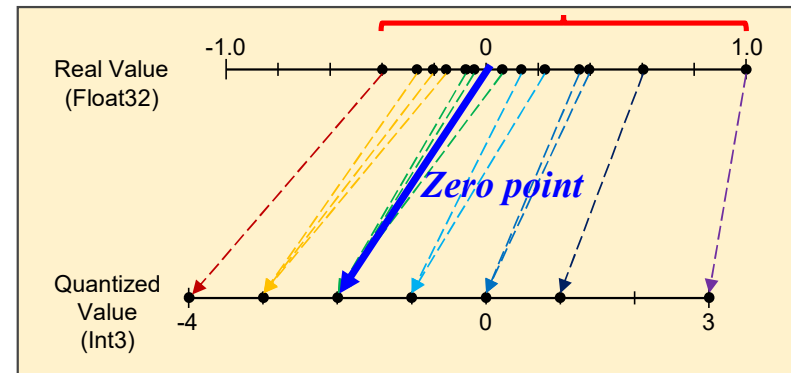
# Quantization Basics – Fundamentals

- **Asymmetric Quantization**
  - Use Zero Point
  - Relatively Slow
  - Can utilize any possible range
- **Symmetric Quantization**
  - Do not use Zero Point
  - Relatively Fast



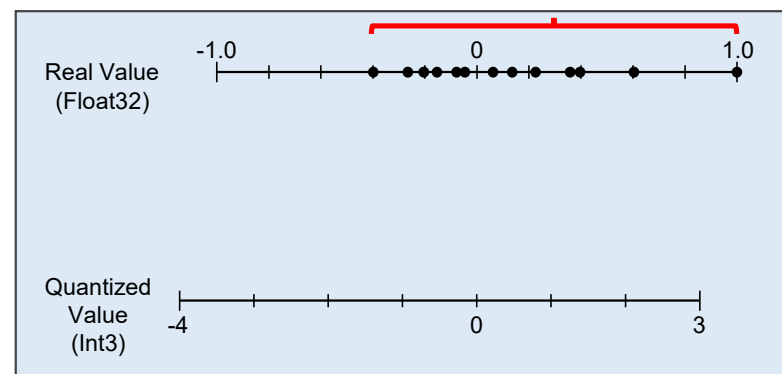
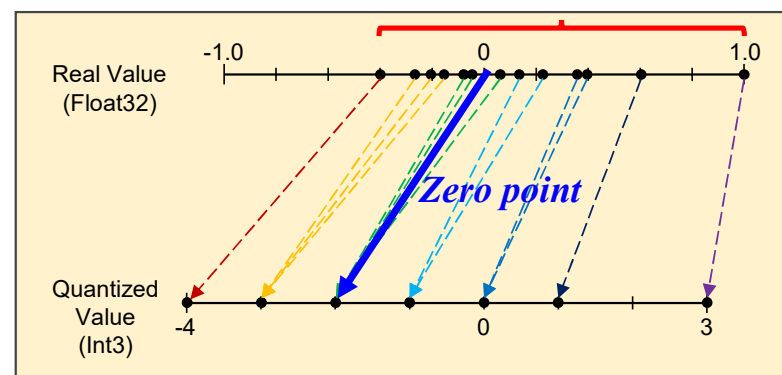
# Quantization Basics – Fundamentals

- **Asymmetric Quantization**
  - Use Zero Point
  - Relatively Slow
  - Can utilize any possible range
- **Symmetric Quantization**
  - Do not use Zero Point
  - Relatively Fast



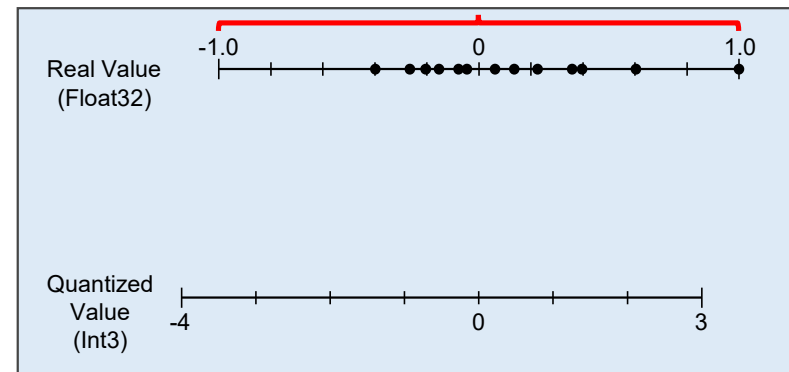
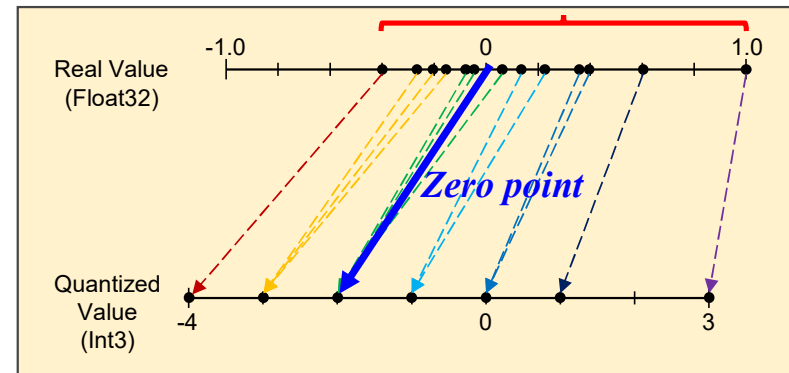
# Quantization Basics – Fundamentals

- **Asymmetric Quantization**
  - Use Zero Point
  - Relatively Slow
  - Can utilize any possible range
- **Symmetric Quantization**
  - Do not use Zero Point
  - Relatively Fast



# Quantization Basics – Fundamentals

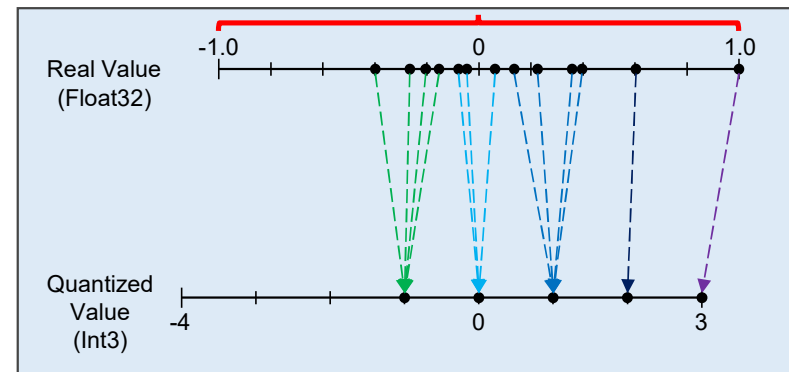
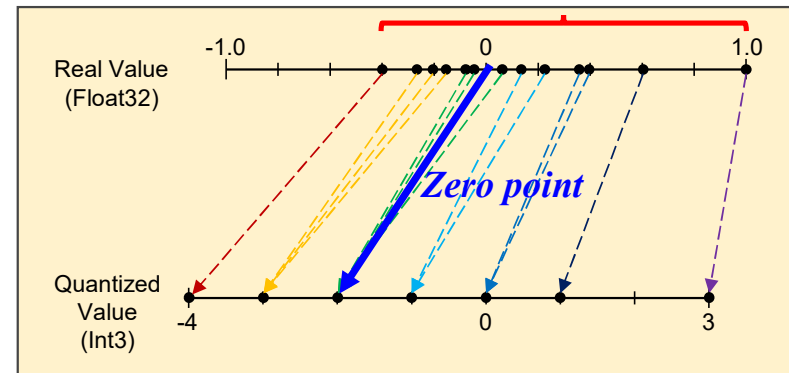
- **Asymmetric Quantization**
  - Use Zero Point
  - Relatively Slow
  - Can utilize any possible range
- **Symmetric Quantization**
  - Do not use Zero Point
  - Relatively Fast





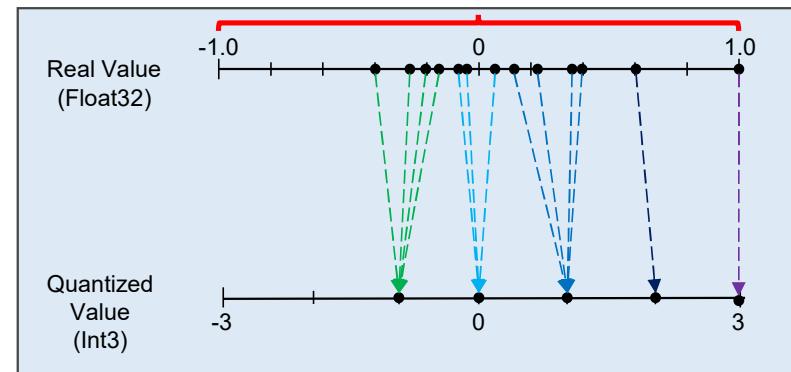
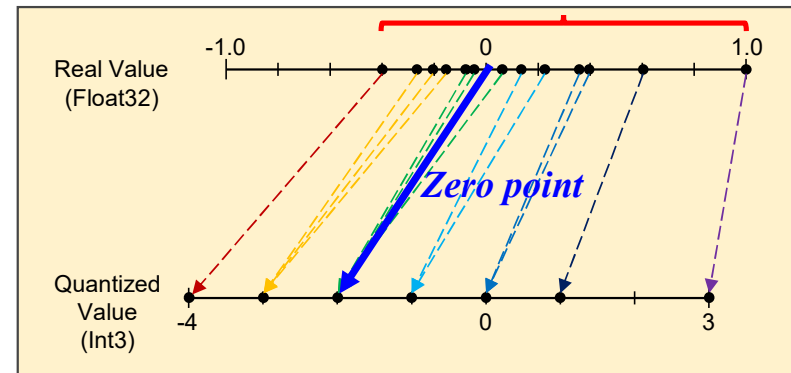
# Quantization Basics – Fundamentals

- **Asymmetric Quantization**
  - Use Zero Point
  - Relatively Slow
  - Can utilize any possible range
- **Symmetric Quantization**
  - Do not use Zero Point
  - Relatively Fast



# Quantization Basics – Fundamentals

- **Asymmetric Quantization**
  - Use Zero Point
  - Relatively Slow
  - Can utilize any possible range
- **Symmetric Quantization**
  - Do not use Zero Point
  - Relatively Fast

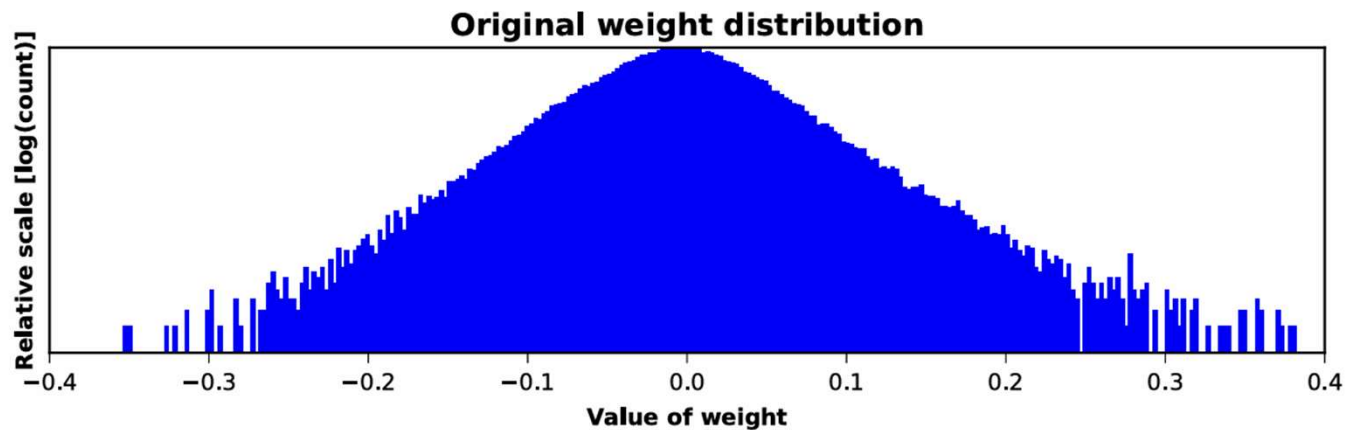


# Quantization Basics – Scheme

- **Uniform Quantization**
- **Non-uniform Quantization**

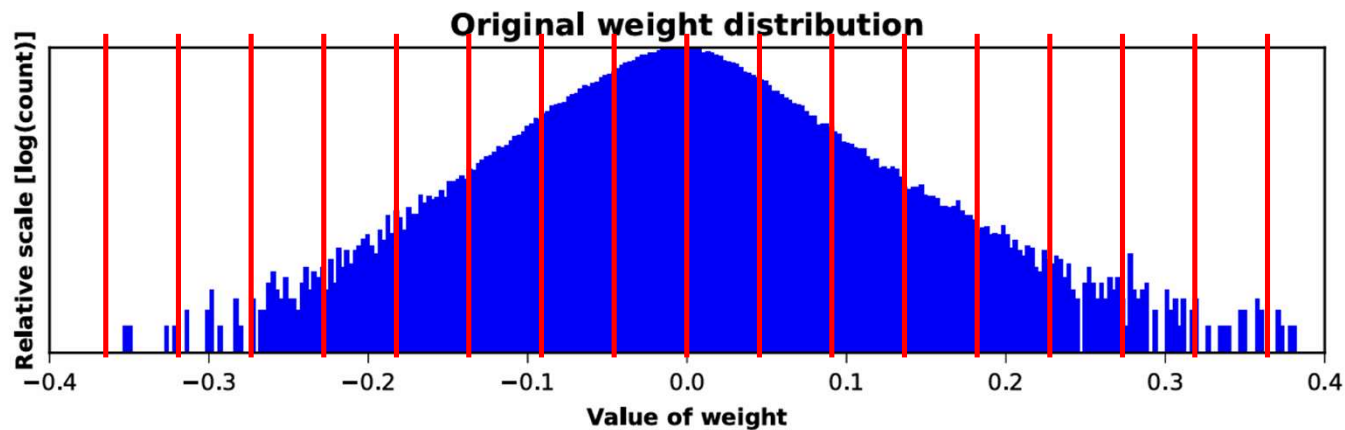
# Quantization Basics – Scheme

- **Uniform Quantization**



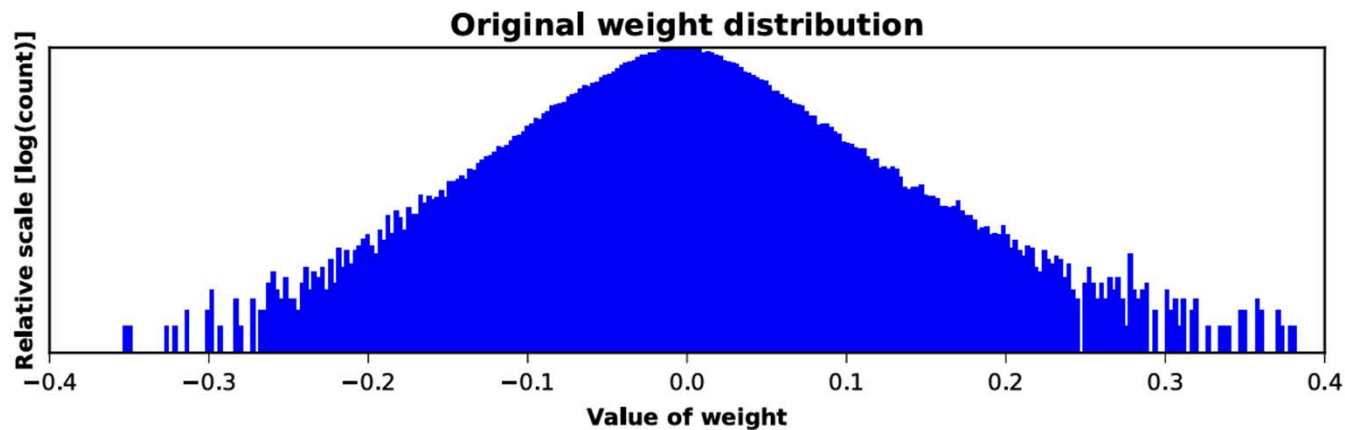
# Quantization Basics – Scheme

- **Uniform Quantization**



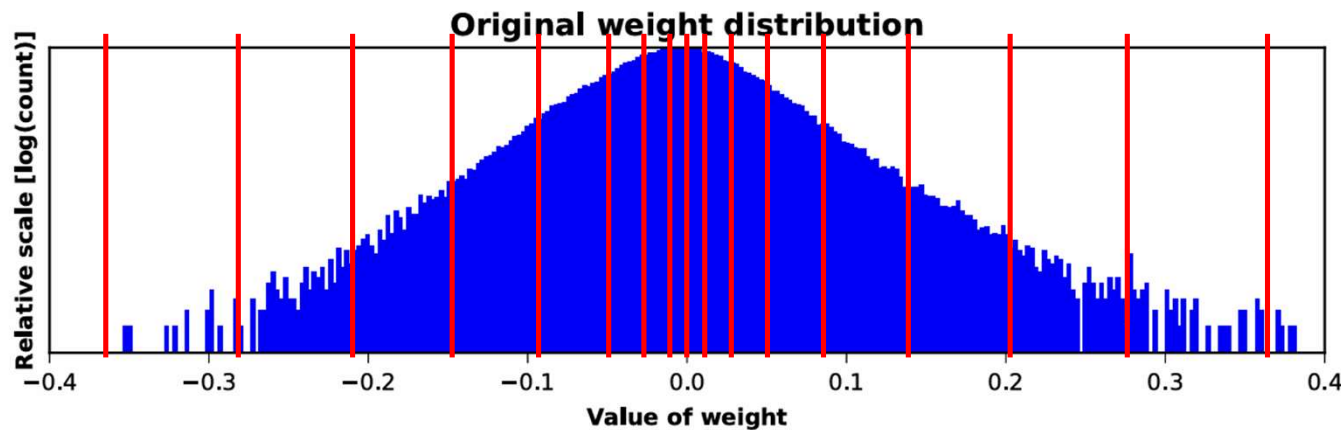
# Quantization Basics – Scheme

- **Non-uniform Quantization**



# Quantization Basics – Scheme

- **Non-uniform Quantization**

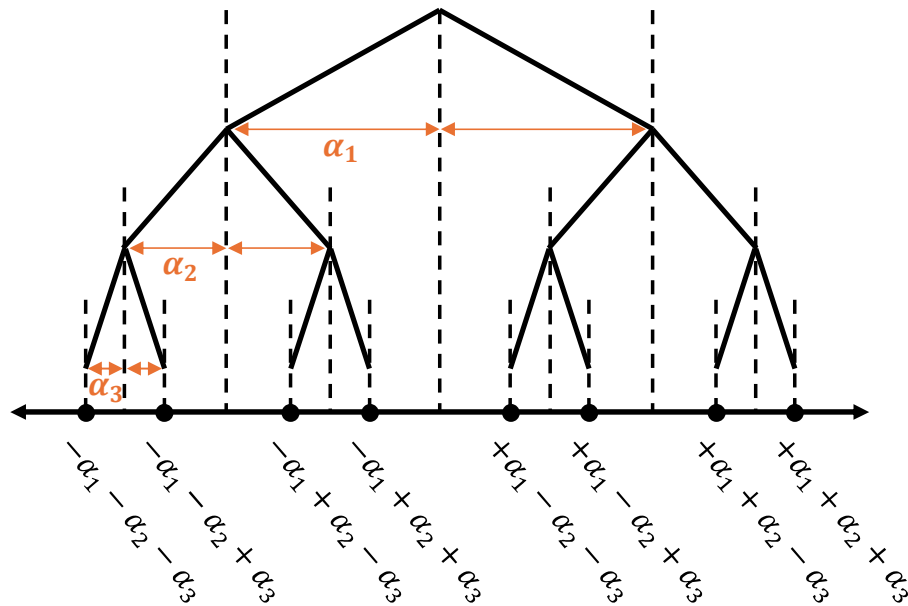


**Logarithmic Quantization (a.k.a Weighted Quantization)**

# Quantization Basics – Scheme

- **Non-uniform Quantization**

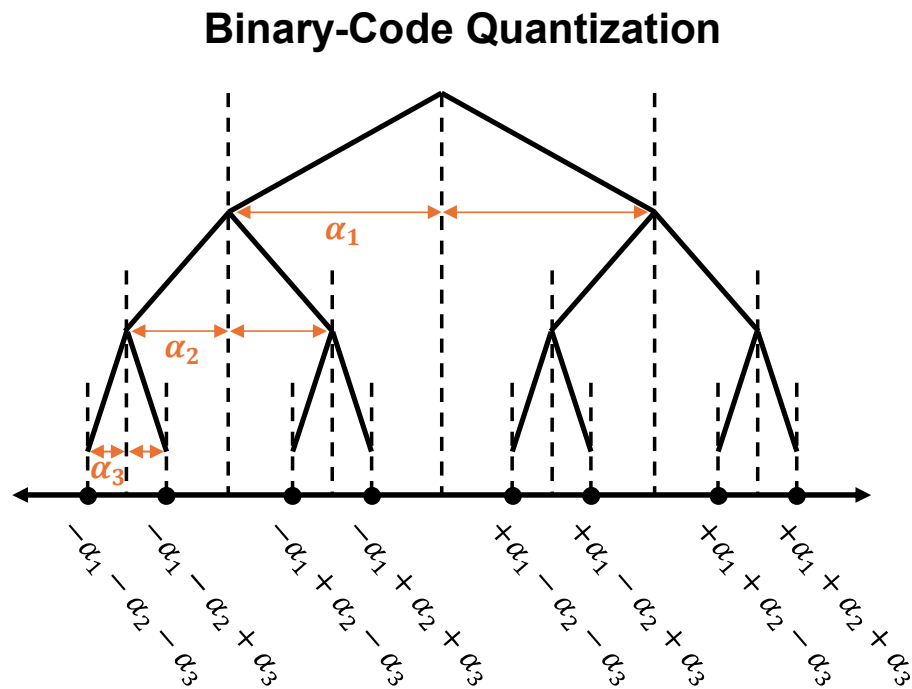
Binary-Code Quantization



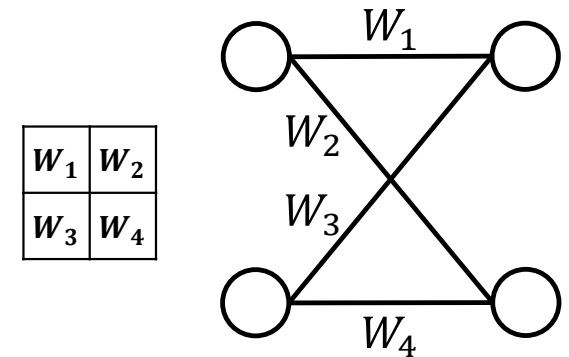


# Quantization Basics – Scheme

- Non-uniform Quantization

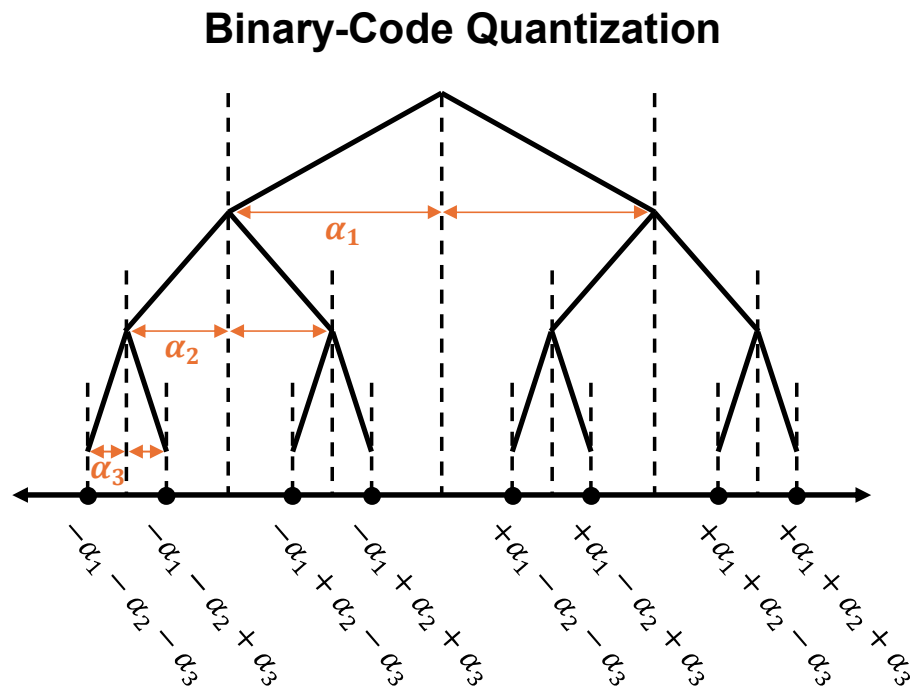


Example of quantizing 2x2 weight



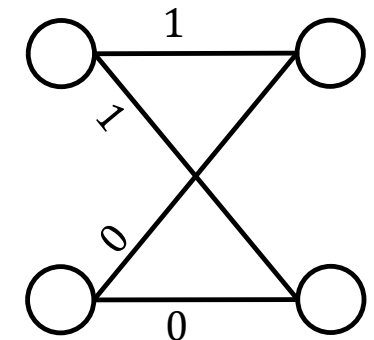
# Quantization Basics – Scheme

- Non-uniform Quantization**



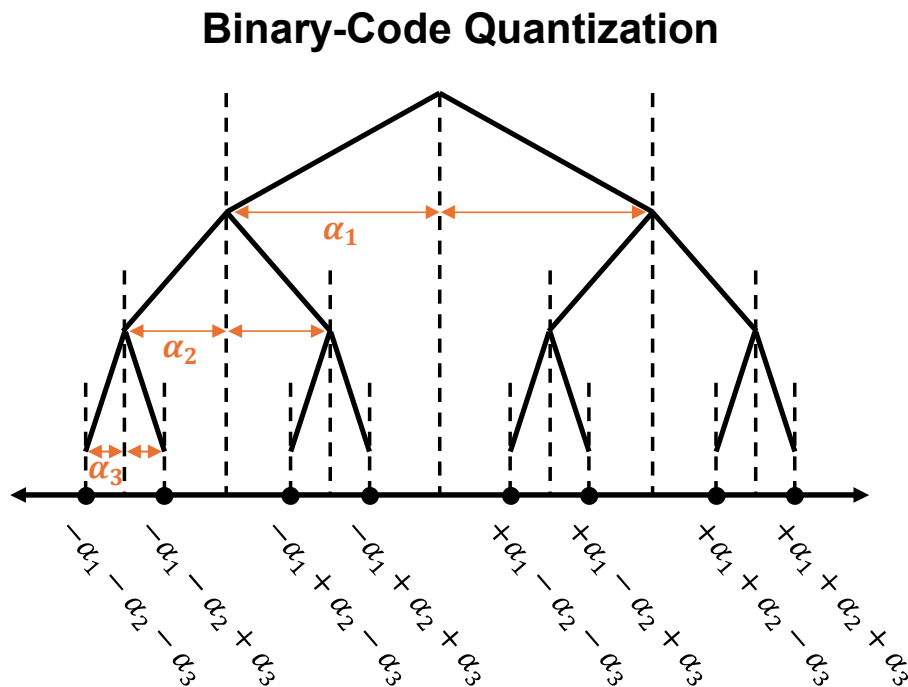
**Example of quantizing 2x2 weight**  
1-bit quantization

$$\begin{bmatrix} \alpha_1 \\ 3.1 \end{bmatrix} \otimes \begin{bmatrix} 1 & 1 \\ -1 & -1 \end{bmatrix} = \begin{bmatrix} W_1 & W_2 \\ W_3 & W_4 \end{bmatrix}$$



# Quantization Basics – Scheme

- Non-uniform Quantization**

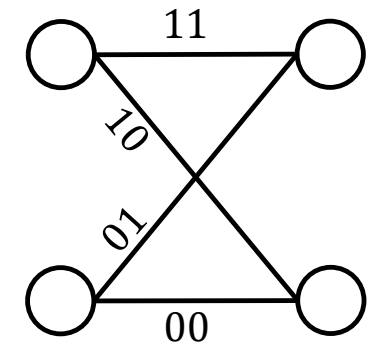


Example of quantizing 2x2 weight  
2-bit quantization

$$\begin{array}{c}
 \alpha_1 \\
 \boxed{3.1}
 \end{array}
 \otimes
 \begin{array}{c}
 B_1 \\
 \begin{array}{|c|c|}
 \hline
 1 & 1 \\
 \hline
 -1 & -1 \\
 \hline
 \end{array}
 \end{array}
 =
 \begin{array}{|c|c|}
 \hline
 W_1 & W_2 \\
 \hline
 W_3 & W_4 \\
 \hline
 \end{array}$$

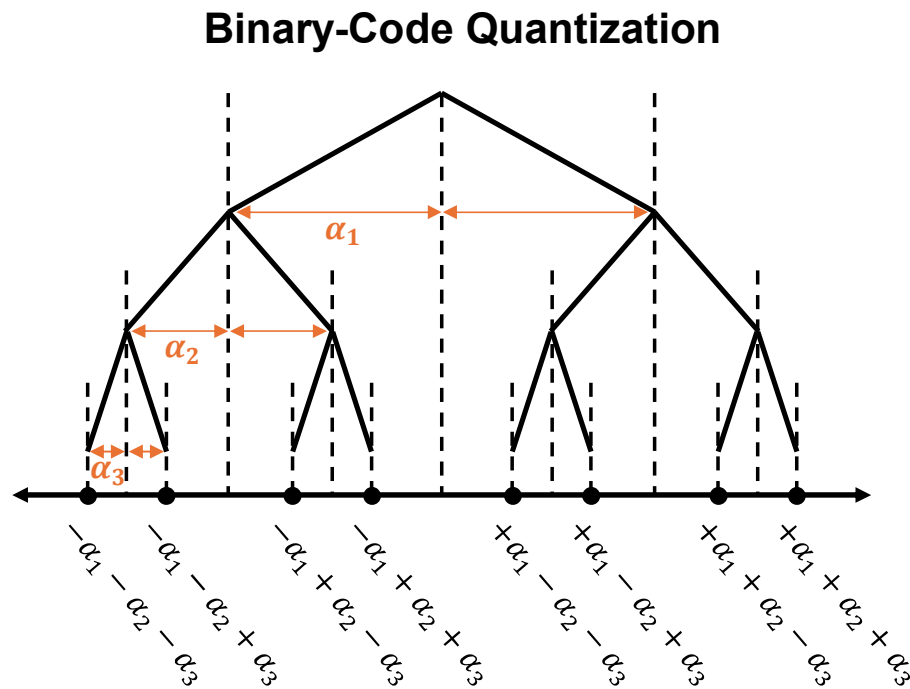
$$+$$

$$\begin{array}{c}
 \alpha_2 \\
 \boxed{1.6}
 \end{array}
 \otimes
 \begin{array}{c}
 B_2 \\
 \begin{array}{|c|c|}
 \hline
 1 & -1 \\
 \hline
 1 & -1 \\
 \hline
 \end{array}
 \end{array}$$



# Quantization Basics – Scheme

- Non-uniform Quantization**



Example of quantizing 2x2 weight  
n-bit quantization

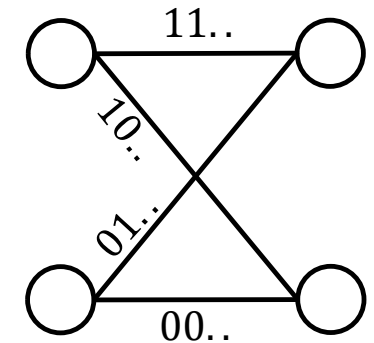
$$\begin{array}{c}
 \alpha_1 \\
 \boxed{3.1}
 \end{array}
 \otimes
 \begin{array}{c}
 B_1 \\
 \begin{array}{|c|c|} \hline 1 & 1 \\ \hline -1 & -1 \\ \hline \end{array}
 \end{array}
 =
 \begin{array}{|c|c|} \hline W_1 & W_2 \\ \hline W_3 & W_4 \\ \hline \end{array}$$

$$+$$

$$\begin{array}{c}
 \alpha_2 \\
 \boxed{1.6}
 \end{array}
 \otimes
 \begin{array}{c}
 B_2 \\
 \begin{array}{|c|c|} \hline 1 & -1 \\ \hline 1 & -1 \\ \hline \end{array}
 \end{array}$$

$$+$$

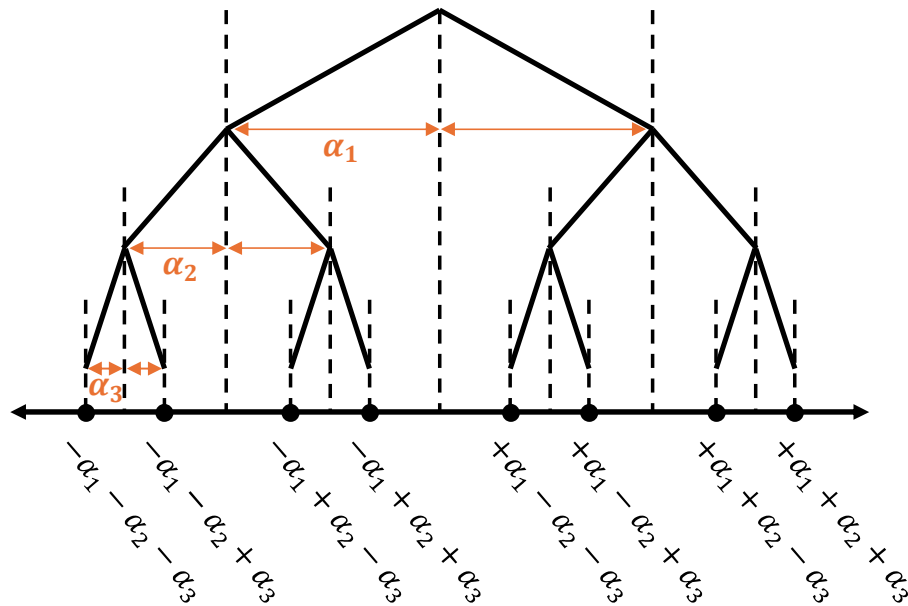
$$\dots$$



# Quantization Basics – Scheme

- **Non-uniform Quantization**

Binary-Code Quantization



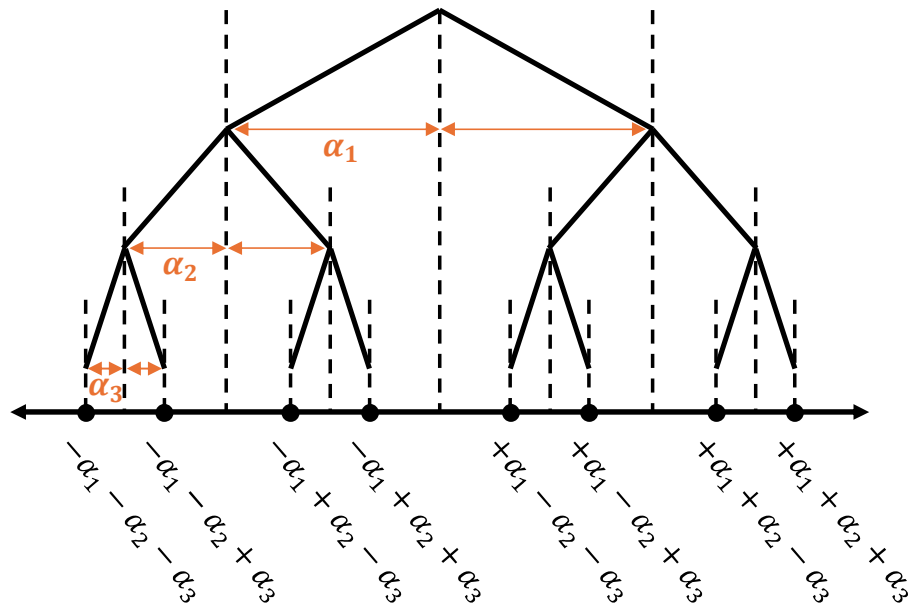
**Very hard to optimize this method**

1. Reconstruction is expensive
  - a lot of elementwise operations
2. Need special optimization for integer MM
  - BiQGEMM
  - LUT-GEMM

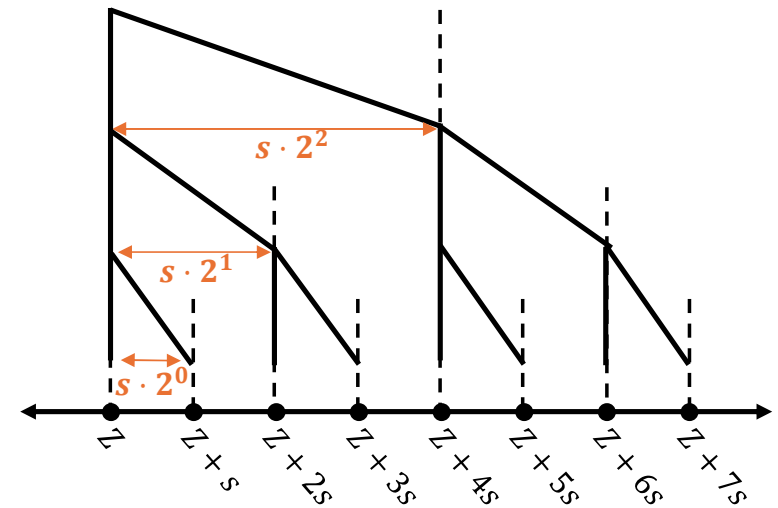
# Quantization Basics – Scheme

- Non-uniform Quantization

Binary-Code Quantization



Uniform Quantization



# Quantization Basics – Type

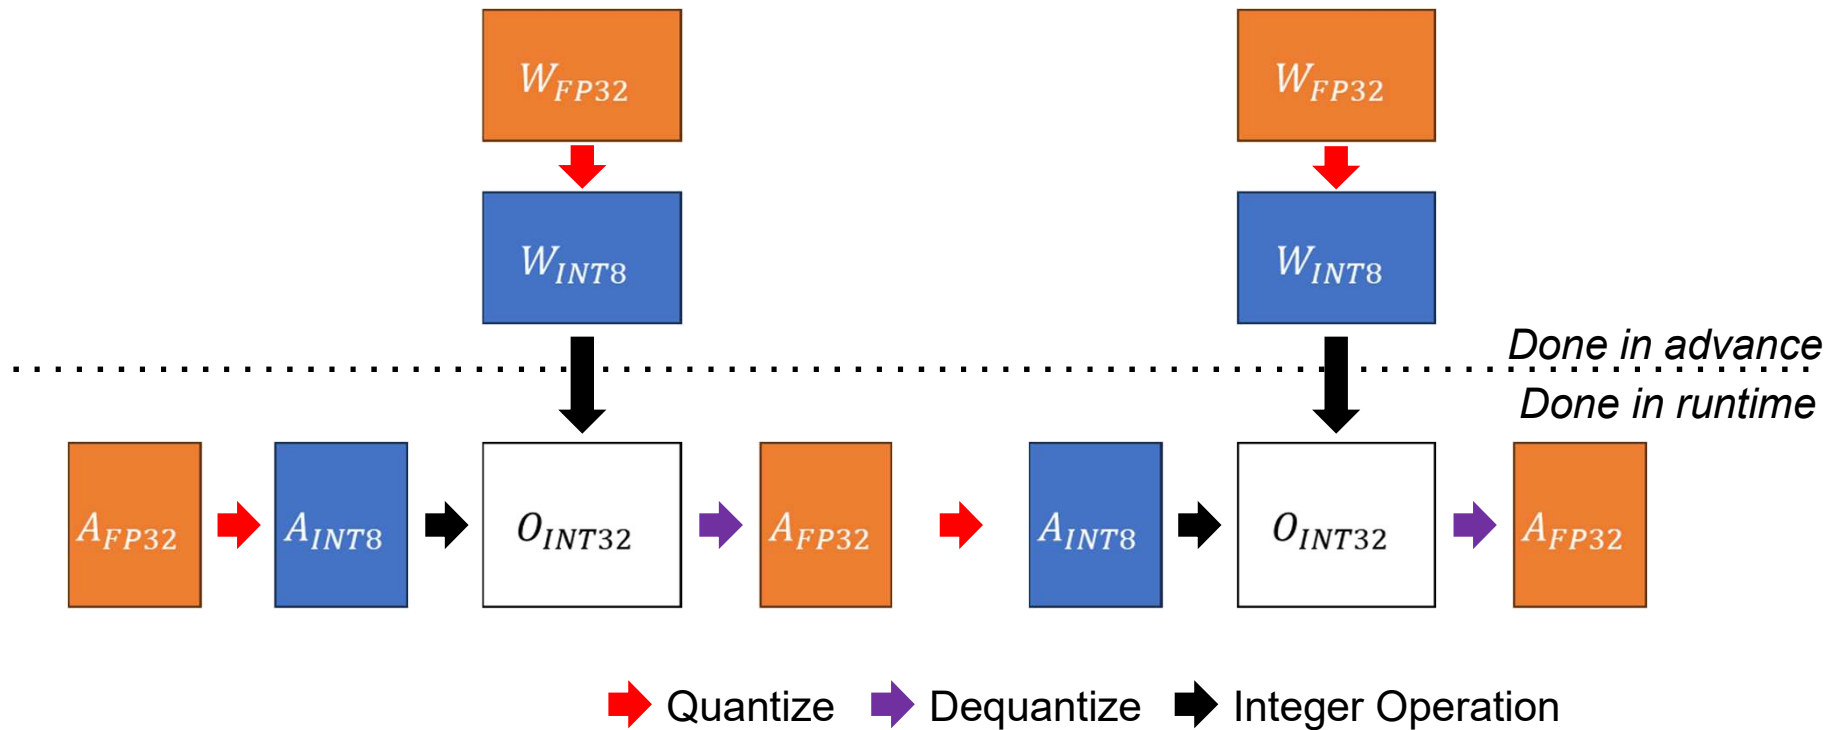
- Dynamic Quantization



➡ Quantize ➡ Dequantize ➡ Integer Operation

# Quantization Basics – Type

- Dynamic Quantization



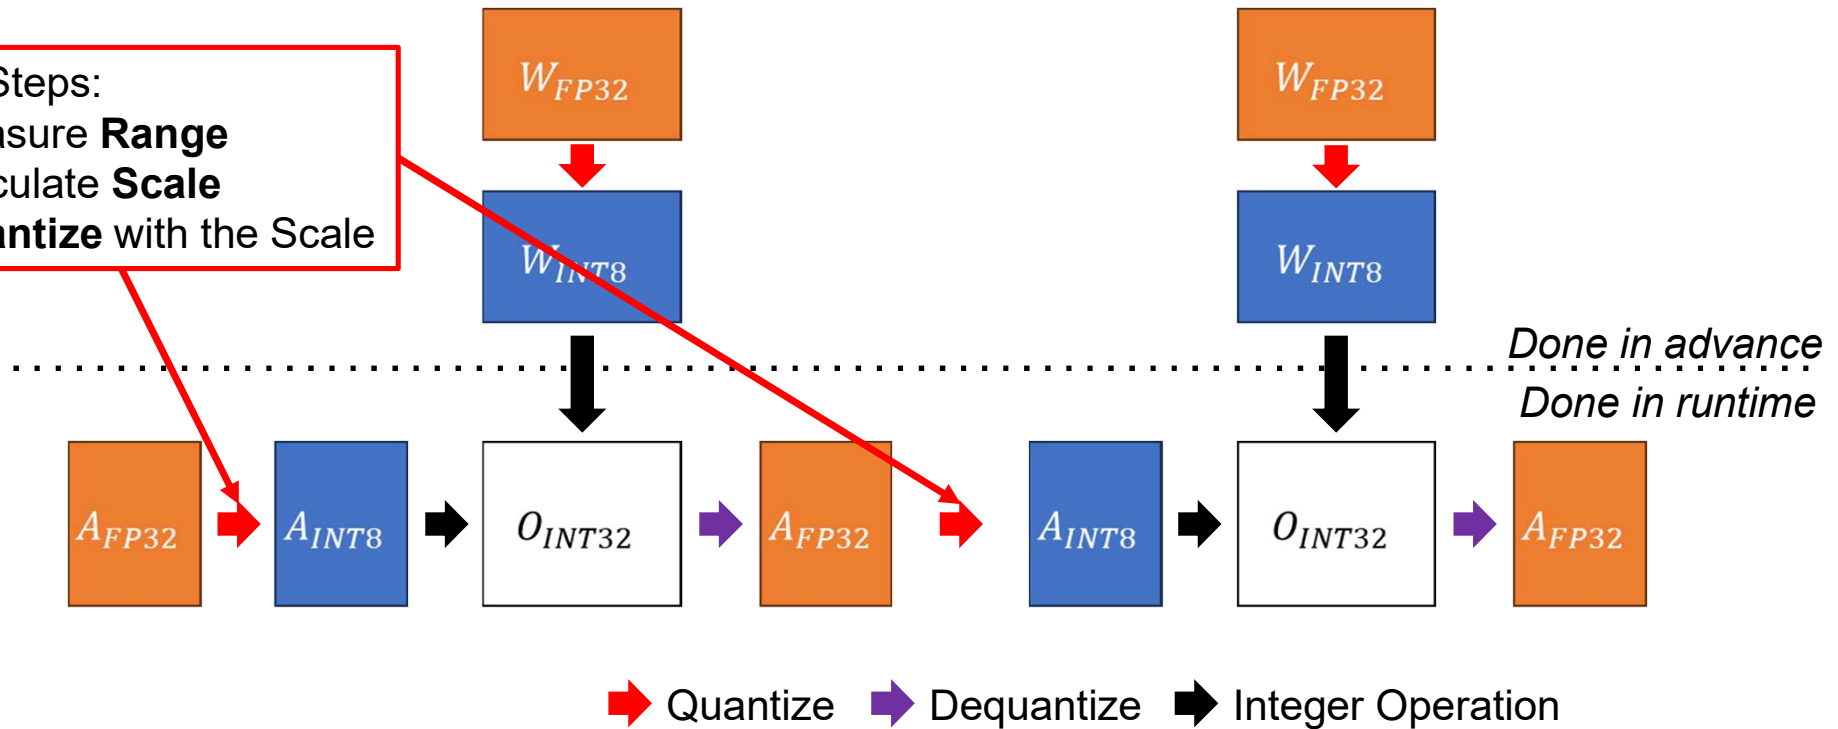


# Quantization Basics – Type

## • Dynamic Quantization

Three Steps:

1. Measure **Range**
2. Calculate **Scale**
3. **Quantize** with the Scale



# Quantization Basics – Type

- **Static Quantization**

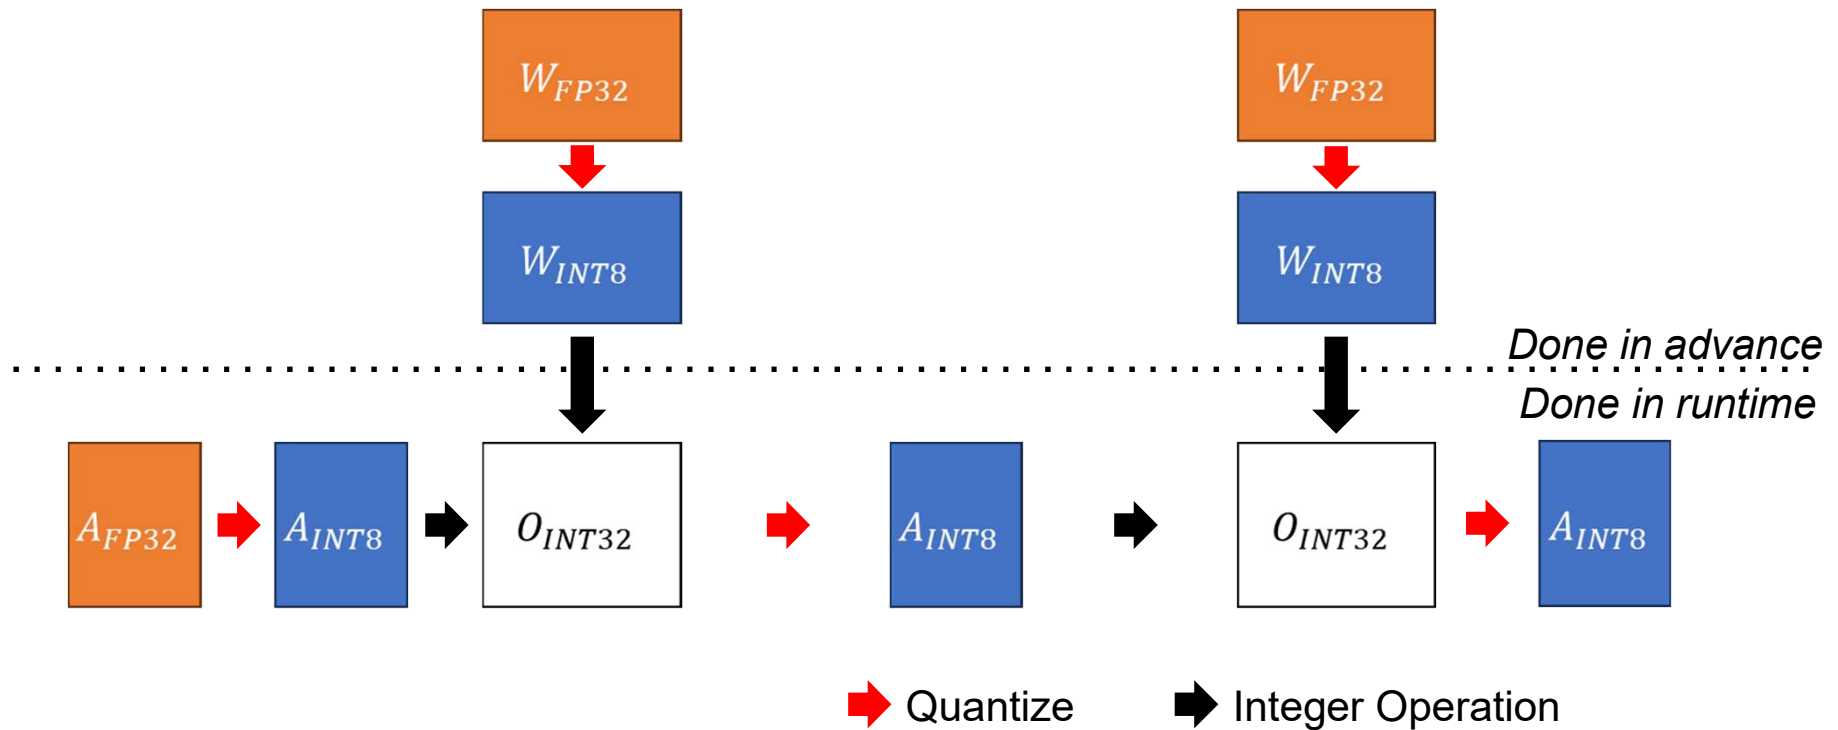


➡ Quantize

➡ Integer Operation

# Quantization Basics – Type

- Static Quantization

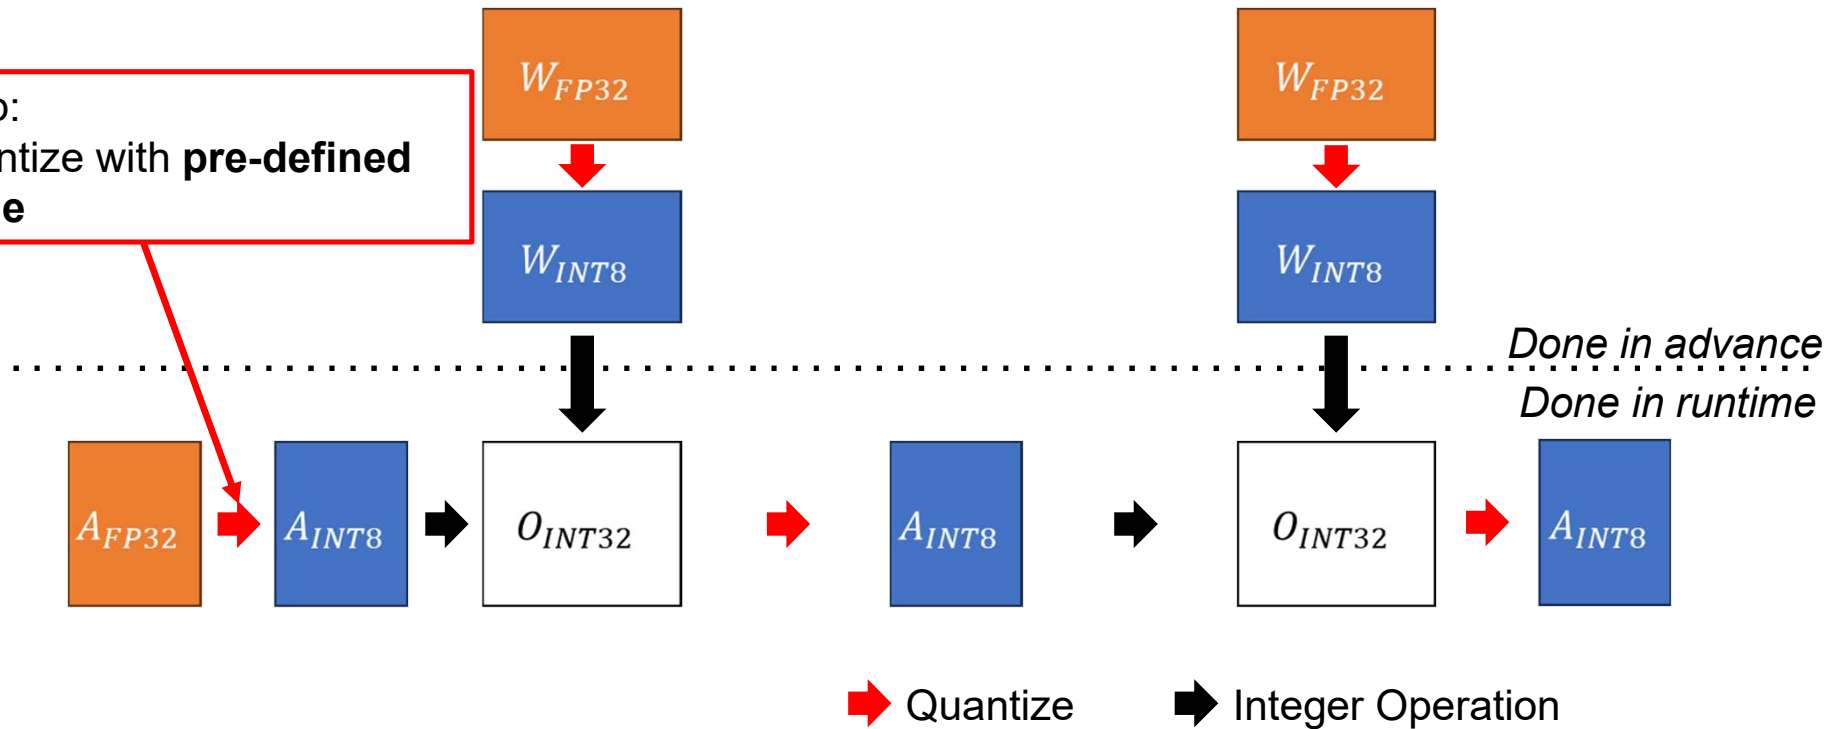


# Quantization Basics – Type

- **Static Quantization**

One Step:

1. Quantize with **pre-defined Scale**



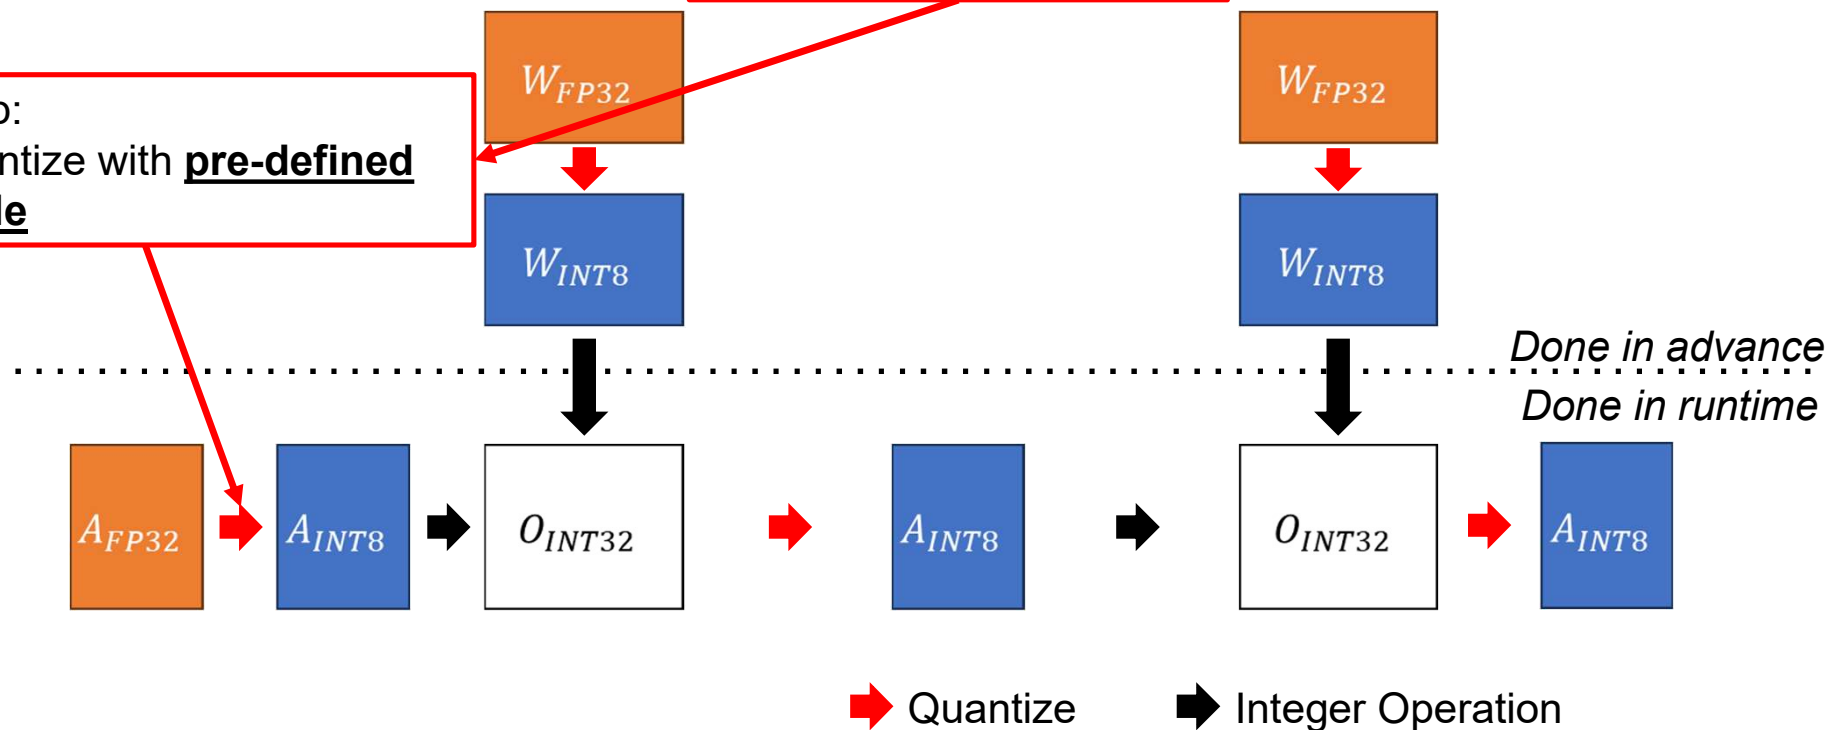
# Quantization Basics – Type

- **Static Quantization**

One Step:

1. Quantize with **pre-defined Scale**

Observe statistics using Calibration Dataset!



# Quantization Basics – Type

- **Static Quantization**
  - Post Training Quantization (PTQ)
  - Quantization Aware Training (QAT)

# Quantization Basics – Type

- **Static Quantization**

- Post Training Quantization (PTQ)
  - Quantize model **after training**
  - Calibrate range statistics with **only a small subset of data**
- Quantization Aware Training (QAT)

# Quantization Basics – Type

- **Static Quantization**

- Post Training Quantization (PTQ)
- Quantization Aware Training (QAT)
  - Quantize **with training/fine-tuning**
  - **Computation is done in FP**, but **fake quantize** to simulate integer operation



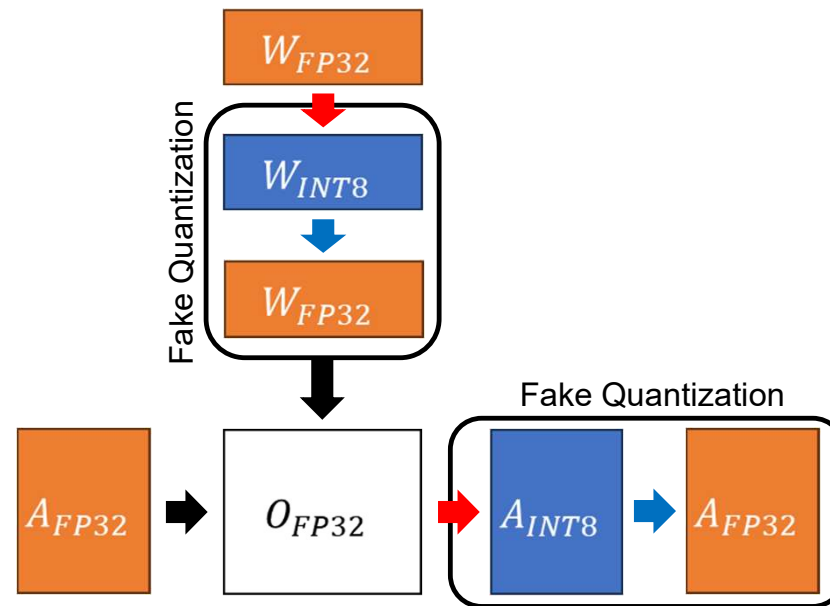
# Quantization Basics – Type

- **Static Quantization**

- Post Training Quantization (PTQ)
- Quantization Aware Training (QAT)
  - Quantize **with training/fine-tuning**
  - **Computation is done in FP**, but **fake quantize** to simulate integer operation
  - However, we can't naively implement the fake quantization due to existence of operation (ex. round) that does not have a gradient function

# Quantization Basics – Type

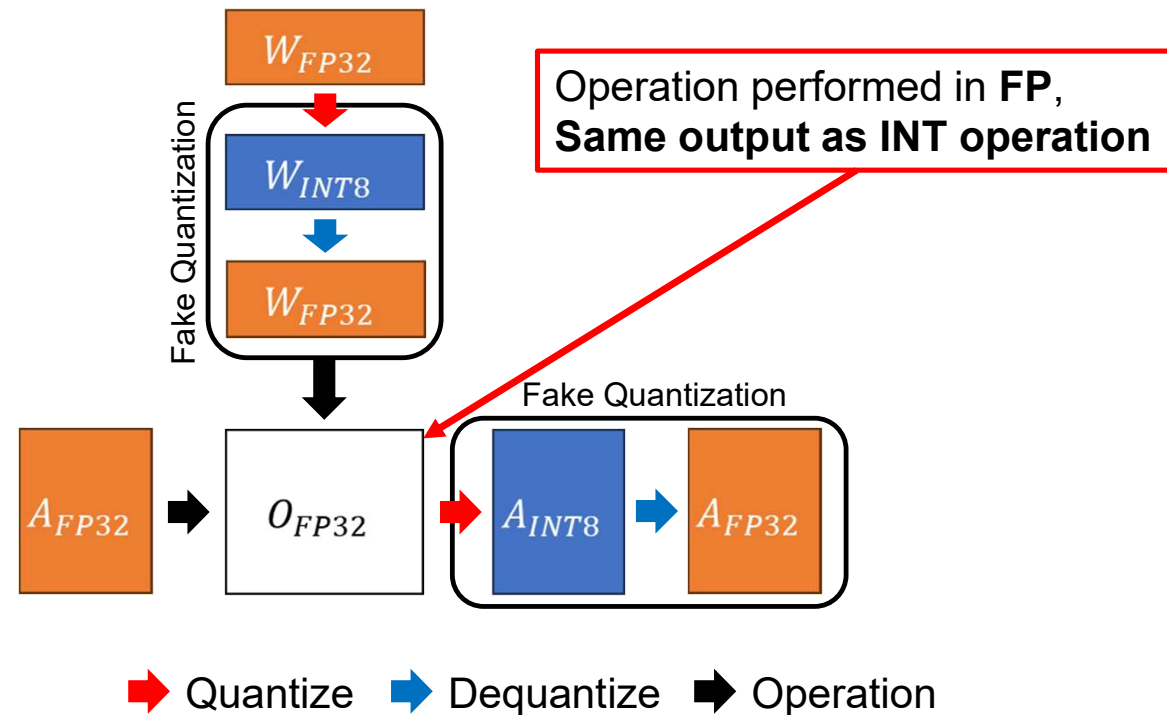
- **Quantization Aware Training : Fake Quantization**



➡ Quantize ➡ Dequantize ➡ Operation

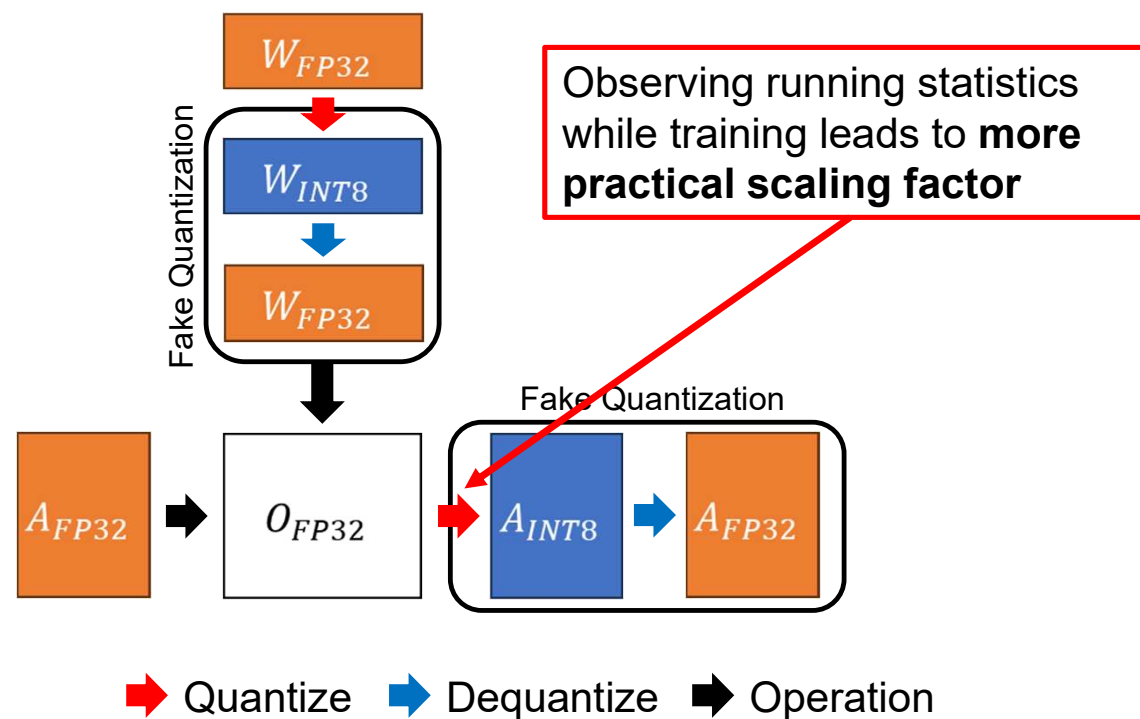
# Quantization Basics – Type

- Quantization Aware Training : Fake Quantization



# Quantization Basics – Type

- Quantization Aware Training : Fake Quantization



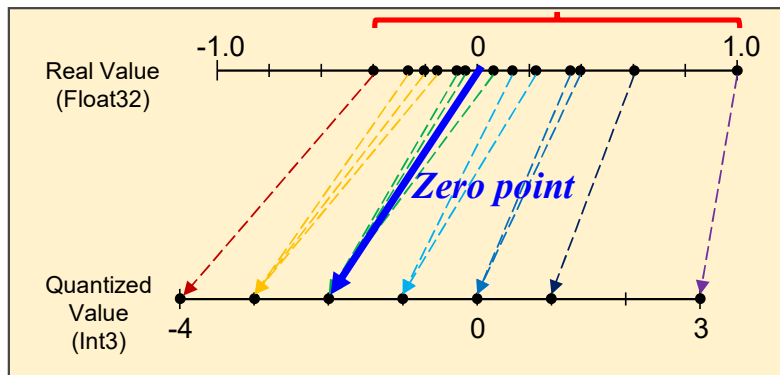
# Break

# Quantization Strategies

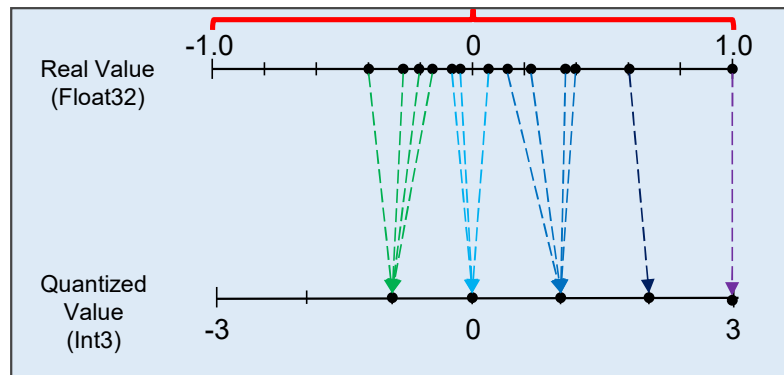
- Quantization Optimization Strategies
  - How can we **calculate** proper scaling factor?
  - What are the **common considerations** when using quantization?

# Quantization Strategies

- How can we **calculate** proper scaling factor?
  1. Simply use **Min/Max** value



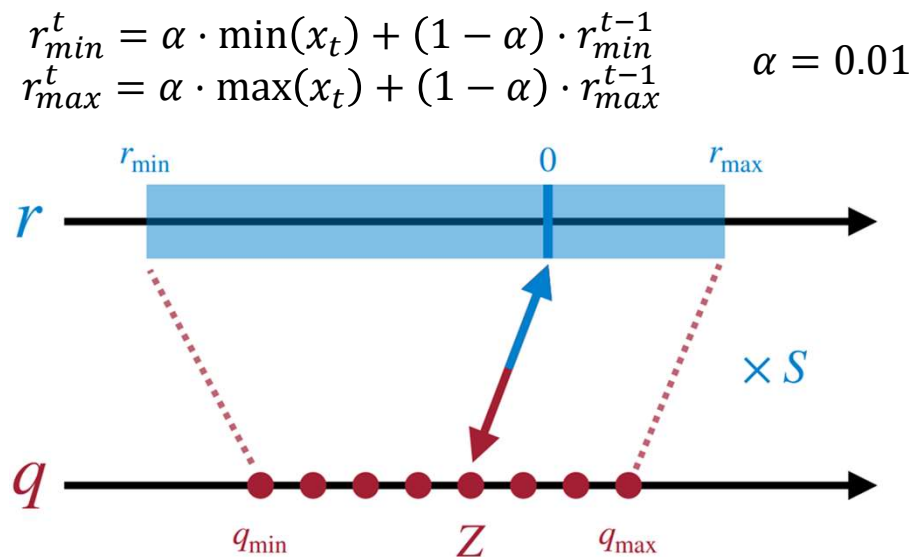
Asymmetric quantization



Symmetric quantization

# Quantization Strategies

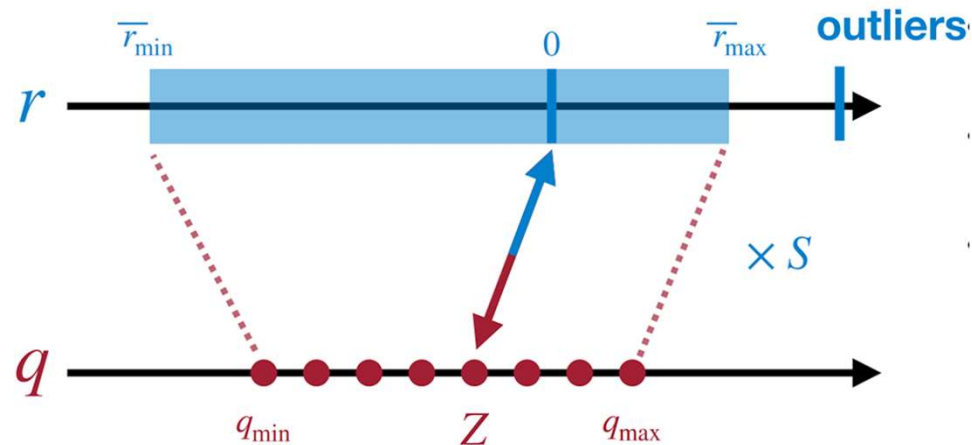
- How can we **calculate** proper scaling factor?
  2. Use **Exponential Moving Averages (EMA)**





# Quantization Strategies

- How can we **calculate** proper scaling factor?
  3. Remove **Outlier**
    - Observe statistics from **calibration samples** and **clamp the range** properly.

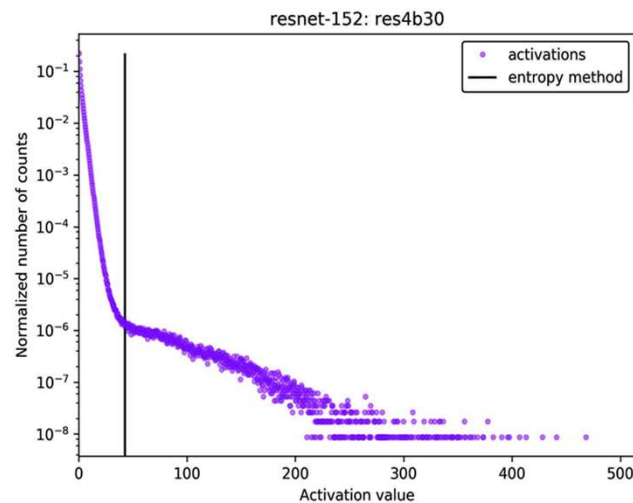


# Quantization Strategies

- How can we **calculate** proper scaling factor?

## 3-1. Entropy Calibration

- find proper clamping value which **minimize KL-divergence score**.

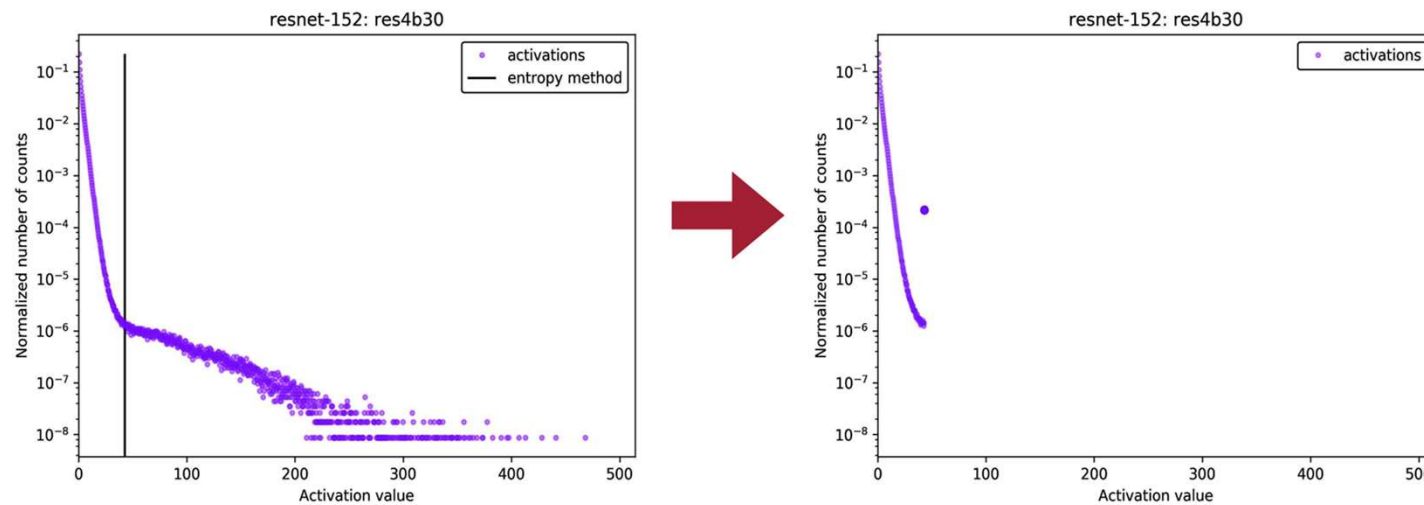


# Quantization Strategies

- How can we **calculate** proper scaling factor?

## 3-1. Entropy Calibration

- find proper clamping value which **minimize KL-divergence score**.

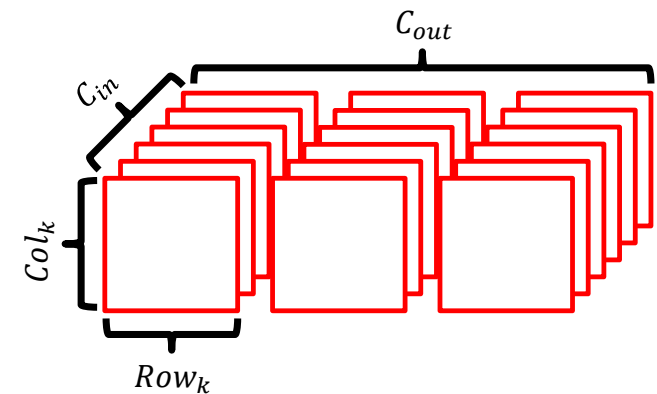


# Quantization Strategies

- How can we **calculate** proper scaling factor?

## 4. Quantization Granularity

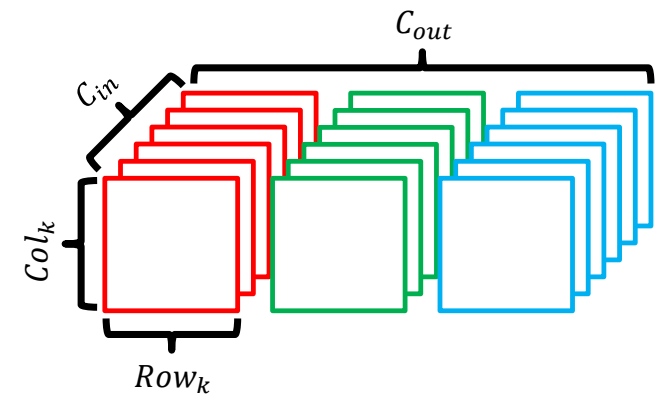
- **Layer-wise** Quantization
  - Calculate scaling factor **per-tensor**



**Convolution Weight filters**

# Quantization Strategies

- How can we **calculate** proper scaling factor?
  - **Quantization Granularity**
    - **Layer-wise** Quantization
      - Calculate scaling factor **per-tensor**
    - **Channel-wise** Quantization
      - Calculate scaling factor **per-channel**



**Convolution Weight filters**

# Quantization Strategies

- **Guideline for Quantization**

- Check the precision supported by hardware acceleration operations



Educators, Students, Makers

Commercial product developers

JETSON NANO 2GB  
5W | 10W

0.5 TFLOPS (FP16)

\$59

JETSON NANO  
5W | 10W

0.5 TFLOPS (FP16)

\$59

JETSON XAVIER NX  
10W | 15W

7 TFLOPS (FP16) | 21 TOPS (INT8)

\$399

JETSON AGX XAVIER  
10W | 15W | 20W

11 TFLOPS (FP16) | 32 TOPS (INT8)

\$899

JETSON AGX ORIN  
15W | 30W | 150W  
Up to 60W

270 TOPS (INT8)

Available Q1 2022

# Quantization Strategies

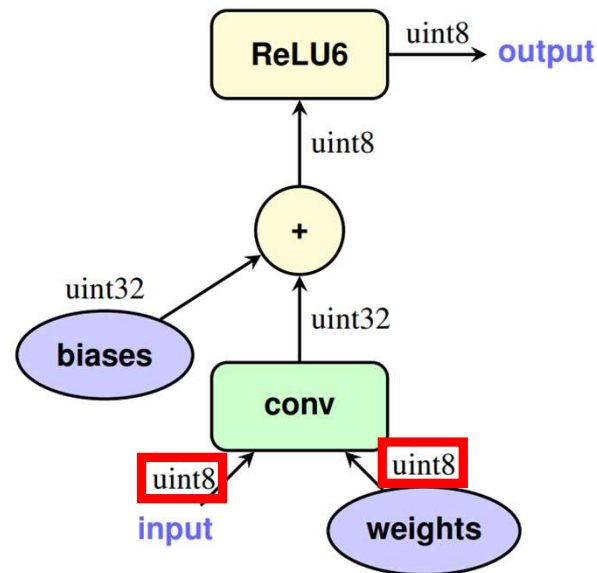
- **Guideline for Quantization**

- Quantizing all layers may not be optimal
  - It is necessary to consider the trade-off between **computational cost** and **accuracy loss** of the model.
  - Example : Compute amount only 0.1% of total inference; Accuracy drop significant (Depthwise-Convolution in MobileNetV3)

# Quantization Strategies

- **Guideline for Quantization**

- Input and weight must have the **same precision** for correct computation





# Quantization Strategies

- **Guideline for Quantization**

- Specific operations require special handling
  - **Non-Linear Activation** (GeLU, Softmax, etc.): FP16
  - **Normalization Layer**: Minimum 16-bit or FP16 or BN Folding
  - **Residual Connection**: Match scaling factor or use INT16, FP16

# Quantization Strategies

- **Guideline for Quantization**

- Challenges of Extremely Low-Bitwidth Quantization

- Changes in Output Distribution (Norm-layer statistics become inaccurate)
    - Necessity of Quantization Aware Training (QAT)

# Break